



Red Hat Reference Architecture Series

Deploying Oracle RAC Database 12c on Red Enterprise Linux 7

Best Practices

Roger Lopez, Principal Software Engineer

Version 1.1

March 2015





100 East Davie Street
Raleigh NC 27601 USA
Phone: +1 919 754 3700
Phone: 888 733 4281
Fax: +1 919 754 3701
PO Box 13588
Research Triangle Park NC 27709 USA

Linux is a registered trademark of Linus Torvalds. Red Hat, Red Hat Enterprise Linux and the Red Hat "Shadowman" logo are registered trademarks of Red Hat, Inc. in the United States and other countries.

Microsoft and Windows are U.S. registered trademarks of Microsoft Corporation.

UNIX is a registered trademark of The Open Group.

Intel, the Intel logo and Xeon are registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries.

All other trademarks referenced herein are the property of their respective owners.

© 2015 by Red Hat, Inc. This material may be distributed only subject to the terms and conditions set forth in the Open Publication License, V1.0 or later (the latest version is presently available at <http://www.opencontent.org/openpub/>).

The information contained herein is subject to change without notice. Red Hat, Inc. shall not be liable for technical or editorial errors or omissions contained herein.

Distribution of modified versions of this document is prohibited without the explicit permission of Red Hat Inc.

Distribution of this work or derivative of this work in any standard (paper) book form for commercial purposes is prohibited unless prior permission is obtained from Red Hat Inc.

The GPG fingerprint of the security@redhat.com key is:
CA 20 86 86 2B D6 9D FC 65 F6 EC C4 21 91 80 CD DB 42 A6 0E

Send feedback to refarch-feedback@redhat.com



Comments and Feedback

In the spirit of open source, we invite anyone to provide feedback and comments on any reference architectures. Although we review our papers internally, sometimes issues or typographical errors are encountered. Feedback allows us to not only improve the quality of the papers we produce, but allows the reader to provide their thoughts on potential improvements and topic expansion to the papers.

Feedback on the papers can be provided by emailing refarch-feedback@redhat.com. Please refer to the title within the email.

Staying In Touch

Join us on some of the popular social media sites where we keep our audience informed on new reference architectures as well as offer related information on things we find interesting.

Like us on Facebook:

<https://www.facebook.com/rhrefarch>

Follow us on Twitter:

<https://twitter.com/RedHatRefArch>

Plus us on Google+:

<https://plus.google.com/u/0/b/114152126783830728030/>



Table of Contents

1 Executive Summary.....	1
2 Reference Architecture Environment.....	2
2.1 Reference Architecture Overview.....	2
2.2 Network Topology.....	3
2.3 Hardware Details.....	4
2.4 File System Layout & Disk Space Details.....	5
2.5 Storage Layout.....	7
2.6 Swap Space.....	7
2.7 Security: Firewall Settings.....	8
2.8 Security: SELinux.....	9
3 Reference Architecture Configuration Details.....	10
3.1 Setting OS Hostname.....	10
3.2 Network Configuration.....	11
3.2.1 Configuring /etc/resolv.conf file.....	11
3.2.2 Public Network Configuration.....	11
3.2.3 Configure SCAN via DNS.....	14
3.2.4 Configure Virtual IP (VIP) via DNS.....	16
3.2.5 Private Network Configuration.....	17
3.3 OS Configuration.....	20
3.3.1 Using the Red Hat Subscription Manager.....	20
3.3.2 NTP Configuration.....	22
3.3.3 Oracle Database 12c Release 1 (12.1) Package Requirements.....	23
3.3.4 Configuring Security-Enhanced Linux (SELinux).....	25
3.3.5 Configuring Firewall Settings.....	25
3.3.6 Setting Virtual Memory.....	28
3.3.7 Setting Shared Memory (SHMMAX, SHMALL, SHMMNI).....	29
3.3.8 Setting Semaphores (SEMMSL, SEMMNI, SEMMNS).....	31
3.3.9 Ephemeral Network Ports.....	31
3.3.10 Optimizing Network Settings.....	32
3.3.11 Setting NOZEROCONF.....	32
3.3.12 Disabling the avahi-daemon service.....	32
3.3.13 Increasing synchronous I/O Requests.....	33
3.3.14 Increasing File Handles.....	33



3.3.15 Kernel Panic On OOPS Parameter.....	34
3.3.16 Reverse Path Filtering.....	34
3.3.17 User Accounts & Groups.....	35
3.3.18 Setting Shell Limits for the Grid and Oracle User.....	36
3.4 Storage Configuration.....	40
3.4.1 Setting up DM-Multipath.....	40
3.4.2 Partitioning Device Mapper Shared Disks.....	43
3.4.3 Configuring Oracle ASM Disks.....	43
3.4.3.1 Oracle ASMLib Alternative: Configuring udev Rules.....	43
3.4.4 Optimizing Database Storage using Automatic System Tuning.....	45
4 Oracle 12c R1 Configuration.....	49
4.1.1 Installing Oracle Grid Infrastructure (Required for ASM).....	49
4.1.2 Installing Oracle 12c R1 Database Software.....	57
4.1.3 Creating ASM Diskgroups via the ASM Configuration Assistant (ASMCA).....	64
4.1.4 Creating Pluggable Databases using Database Configuration Assistant (DBCA)....	67
4.1.5 Enabling HugePages.....	74
5 Logging into the Oracle Container Database 12c Release 1 (12.1).....	77
6 Post Installation Cleanup Tasks.....	80
6.1.1 Removal of Oracle Kernel Parameters from /etc/sysctl.conf.....	80
6.1.2 Removal of firewalld Trusted Source Address.....	81
7 Common Tasks when Managing Container Database (CDB) and Pluggable Databases (PDB).....	82
7.1 Connect to a CDB.....	82
7.2 Connect to a PDB.....	84
7.3 Managing a CDB.....	85
7.4 Managing a PDB.....	86
7.5 Location of Data files of PDBs & CDB.....	87
8 Conclusion.....	89
Appendix A: Revision History.....	90
Appendix B: Contributors.....	91
Appendix C: DM Multipath Configuration File.....	92



Appendix D: Firewall Configuration File.....	94
Appendix E: Huge Pages Script.....	95
Appendix F: Oracle Database Package Requirements Text File.....	96
Appendix G: Kernel Parameters.....	97
Appendix H: 98-oracle.conf.....	98
Appendix I: Limits Configuration File (99-grid-oracle-limits.conf).....	99
Appendix J: 99-oracle-asmdevices.rules.....	100
Appendix K: Sample Kickstart File.....	101
Appendix L: Custom Oracle Tuned Profile (script.sh & tuned.conf).....	102
Appendix M: Configuration Files.....	105
Appendix N: Troubleshooting ORA-* Errors.....	106
Appendix O: References.....	110



1 Executive Summary

IT organizations face challenges of optimizing Oracle database environments to keep up with the ever increasing workload demands and evolving security risks. This reference architecture provides a step-by-step deployment procedure with the latest best practices to install and configure an Oracle Real Application Clusters (RAC) Database 12c Release 1 (12.1)¹ with Oracle Automatic Storage Management (ASM). It is suited for system, storage, and database administrators deploying Oracle RAC Database 12c Release 1 (12.1) on Red Hat Enterprise Linux 7. It is intended to provide a Red Hat | Oracle reference architecture that focuses on the following tasks:

- Deploying Oracle Grid Infrastructure 12c Release 1 (12.1.0.2.0)
- Deploying Oracle RAC Database Software 12c Release 1 (12.1.0.2.0)
- Deploying an Oracle RAC Database 12c Release 1 (12.1.0.2.0) with shared SAN disks
- Using Oracle ASM disks with *udev* rules
- Securing the Oracle RAC Database 12c environment with *SELinux*

¹ Oracle Database 12c – supported on version 12.1.0.2.0 or higher on Red Hat Enterprise Linux 7



2 Reference Architecture Environment

This section focuses on the components used during the deployment of Oracle RAC Database 12c with Oracle Automatic Storage Management (ASM) on Red Hat Enterprise Linux 7 x86_64 in this reference architecture.

2.1 Reference Architecture Overview

A pictorial representation of the environment used in this reference environment is shown in **Figure 2.1.1: Reference Architecture Overview**




Figure 2.1.1: Reference Architecture Overview



2.2 Network Topology

The network topology used in this reference environment consists of two public switches using link aggregation to connect the two switches together (*Public Switch A* and *Public Switch B*), creating a single logical switch. A similar link aggregation is done for private switches (*Private Switch A* and *Private Switch B*), creating a single logical switch. Ethernet device *ens10f0* on the server connects to *Public Switch A*, while Ethernet device *ens10f1* on the server connects to *Public Switch B*. Ethernet devices *ens10f0* and *ens10f1* are bonded together as a bond device, *bond0*, providing high availability for the public network traffic. Ethernet device *ens10f2* on each server connects to *Private Switch A*, while Ethernet device *ens10f3* on each server connects to *Private Switch B*. Ethernet devices *ens10f2* and *ens10f3* take advantage of Oracle's Highly Available Internet Protocol (HAIP) for Oracle's private interconnect. HAIP can load balance Ethernet traffic for up to four Ethernet devices. Due to the use of Oracle's HAIP, no bond device is created for private Ethernet devices *ens10f2* and *ens10f3* on each node within the Oracle RAC Database cluster. **Figure 2.2.1: Network Topology** shows the pictorial representation of the network topology.

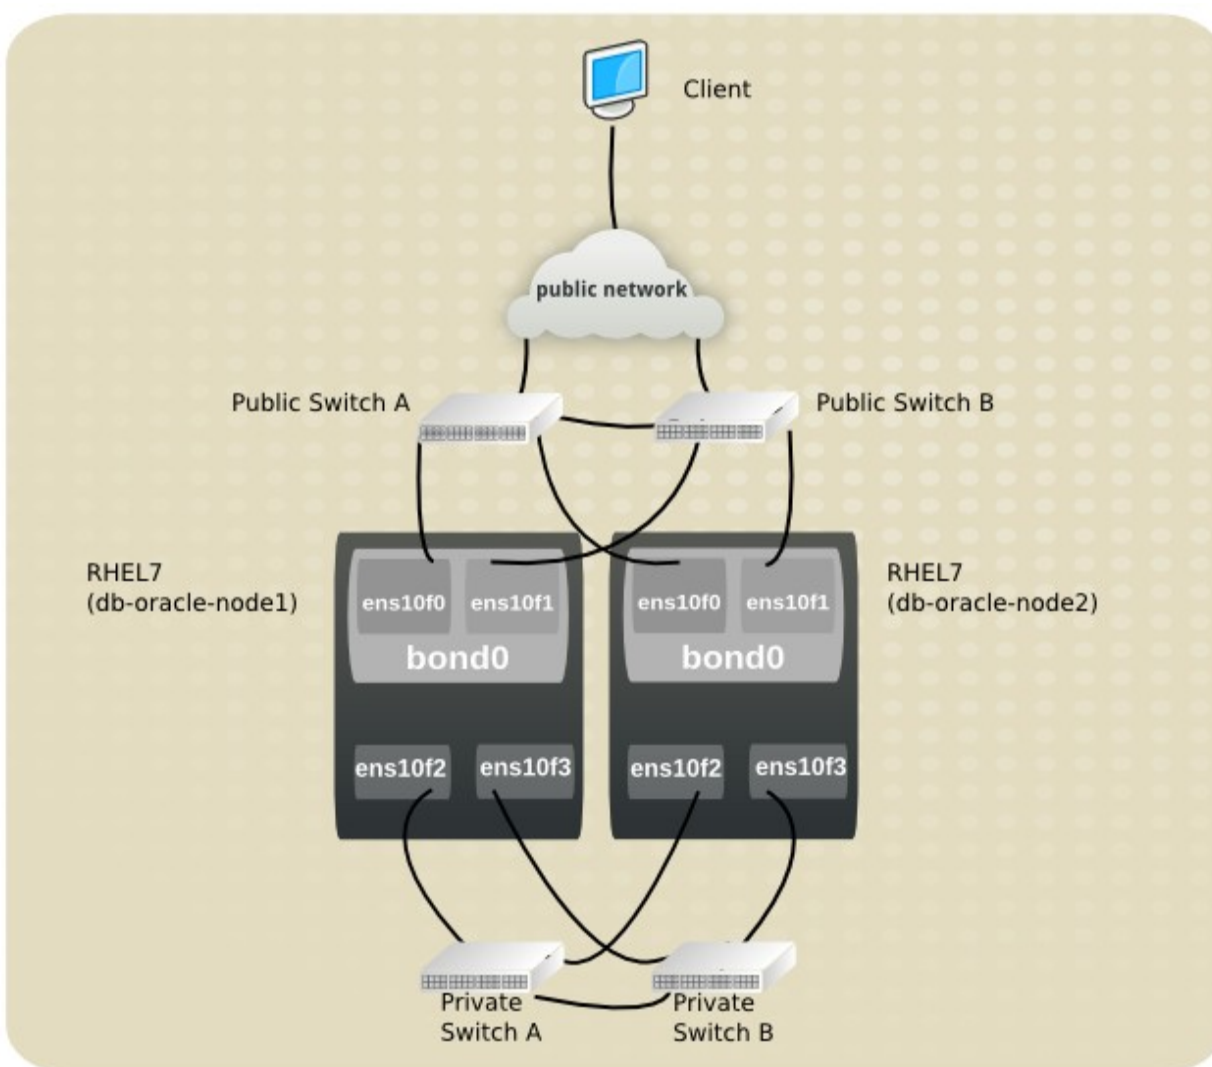


Figure 2.2.1: Network Topology



2.3 Hardware Details

The following are the hardware requirements to properly install Oracle RAC Database 12c on a x86_64 system:

- Minimum of 4 GB of RAM for the installation of both Oracle Grid Infrastructure and Oracle Database²
- The minimum of three Network Interface Cards (NIC) with the usage of direct attach storage or fibre channel storage; however, four NICs are recommended.
- Red Hat Enterprise Linux 7.x Server x86_64 with kernel 3.10.0-123 or higher
- Console access that supports 1024 x 768 resolution to ensure correct display of Oracle's Universal Installer (OUI).
- All nodes within the Oracle RAC Database environment require the same chip architecture. This reference architecture uses 64-bit processors on all nodes within the cluster.

Table 2.3.1: Server Details specifies the hardware for the server used within this reference environment. This hardware meets the minimum requirements for properly installing Oracle Database 12c Release 1 (12.1) on a x86_64 system.

Server Hardware	Specifications
Oracle RAC 12c Release 1 Server (<i>db-oracle-node1</i> , <i>db-oracle-node2</i>) [2 x HP ProLiant DL370 G6 Server]	Red Hat Enterprise Linux 7 kernel 3.10.0-123.el7.x86_64
	2 Socket, 8 Core (16 cores) Intel(R) Xeon(R) CPU W5580 @ 3.20 GHz
	48 GB of memory, DDR3 4096 MB @ 1333 MHz DIMMs
	2 x NetXen NX3031 1/10-Gigabit Network Interface Cards (NICs) for public network
	2 x NetXen NX3031 1/10-Gigabit Network Interface Cards (NICs) for private network
	1 x Qlogic ISP2532 8GB Fibre Channel Dual Port HBA ³

Table 2.3.1: Server Details

² Preparing your Cluster, Oracle Documentation - <https://docs.oracle.com/database/121/TDPRC/preparing.htm#TDPRC122>

³ Due to hardware limitations, this reference architecture uses only one dual port HBA. It is recommended to have at least two Fibre Channel Single Port HBAs for high availability.



Table 2.3.2: Switch Details specifies the Fibre Channel and Ethernet switches used within this reference environment.

Switch Hardware	
Fibre Channel	2 x Brocade Silkorm Fibre Switches
Public Network	1 x Juniper EX4200 Switch
Private Network (Private VLANs)	1 x Juniper EX4200 Switch

Table 2.3.2: Switch Details

Table 2.3.3: Storage Details specifies the storage used for storing Oracle data files within this reference environment.

Storage Hardware	Specification
HP StorageWorks MSA2324fc Dual Controller Array	24 x 146 GB 15K SAS Hard disks

Table 2.3.3: Storage Details

2.4 File System Layout & Disk Space Details

The following is the disk space requirements for properly installing Oracle Database 12c Release 1 (12.1) software for this reference environment.

Software	Disk Space
Oracle Grid Infrastructure Home (includes software files)	12 GB
Oracle Database Home Enterprise Edition (includes software files and data files)	12 GB
/tmp	1 GB

Table 2.4.1: Disk Space Requirements

NOTE: The actual amount of disk space consumed for Oracle Grid Infrastructure Home and Oracle Database Home Enterprise Edition may vary.



Table 2.4.2: File System Layout specifies the file system layout for the server used in this reference environment. The layout ensures the disk space requirements to properly install the Oracle Grid Infrastructure and Oracle Database software for Oracle RAC Database 12c Release 1 (12.1).

File System Layout	Disk Space Size
/	15 GB
/dev/shm	24 GB
/boot	248 MB
/home	8 GB
/tmp	4 GB
/u01	50 GB
/usr	5 GB
/var	8 GB

Table 2.4.2: File System Layout

Oracle RAC Database 12c Release 1 (12.1) recommends three volumes each of at least 4 GB in size to store the Oracle Cluster Registry (OCR), voting disks, and the Oracle Grid Infrastructure Management Repository (GIMR) within an Oracle ASM disk group with the use of normal redundancy. The OCR manages the Oracle Clusterware and Oracle RAC Database 12c Release 1 configuration information. The voting disk manages any information pertaining to the node membership. GIMR is a mandatory installation with the release of Oracle Grid Infrastructure 12c Release 1 (12.1.0.2). Via Oracle's documentation, it enables the Cluster Health Monitor, Oracle Database QoS Management, Rapid Home Provisioning, and provides a historical metric repository that simplifies viewing of past performance and diagnosis issues. While the size of the Oracle data files varies for each solution, the following are the Oracle data file sizes used for this reference environment.

Volume	Volume Size
Oracle Database Volume 1 (db1)	100 GB
Oracle Database Volume 2 (db2)	100 GB
Fast Recovery Area (fra)	200 GB
Oracle Redo Log Volume (redo)	10 GB
OCR & Voting Disks (ocrvote1)	4 GB
OCR & Voting Disks (ocrvote2)	4 GB
OCR & Voting Disks (ocrvote3)	4 GB

Table 2.4.3: Oracle OCR, Voting Disk, & Data File Sizes



2.5 Storage Layout

Table 2.5.1: Storage Disk Layout for Reference Architecture shows the storage disk layout for each volume.

Virtual Diskgroup Name	Volume Name	Volume Size	RAID Group Type	Hard Drive Count	Hot Spares Available	Size of Virtual Disk
vd01	db1	100 GB	Raid 10	8	1	586 GB
vd02	db2	100 GB	Raid 10	8	1	586 GB
vd03	fra	200 GB	Raid 5	5	0	586 GB
vd04	redo	10 GB	Raid 1	2	0	146 GB
vd04	ocrvote1	4 GB	Raid 1	2	0	146 GB
vd04	ocrvote2	4 GB	Raid 1	2	0	146 GB
vd04	ocrvote3	4 GB	Raid 1	2	0	146 GB
-	-	-	-	1	1	146 GB

Table 2.5.1: Storage Disk Layout for Reference Architecture

NOTE: The Hot Spare Available is a Global Hot Spare that can be applied to any virtual disk group in case of failure.

2.6 Swap Space

Swap space is determined by the amount of RAM found within the system. The following table displays the swap space recommendation. This reference environment allocates 16 GB of RAM for swap space.

RAM	Swap Space
2 GB up to 16 GB	Equal to the size of RAM
Greater than 16 GB	16 GB of RAM

Table 2.6.1: Recommended Swap Space

NOTE: When calculating swap space, ensure not to include RAM assigned for *hugepages*. More information on *hugepages* can be found in **Section 4.1.5 Enabling HugePages**



2.7 Security: Firewall Settings

Red Hat Enterprise Linux 7 introduces the dynamic firewall daemon, **firewalld**. **firewalld** provides a dynamically managed firewall with support for network/firewall zones to define the trust level of network connections or interfaces⁴. **firewalld** is the default firewall service in Red Hat Enterprise Linux 7, however, **iptables** service is still available. It is important to note that with the **iptables** service, every single change means flushing all the old rules and reading all of the new rules from the `/etc/sysconfig/iptables`. With **firewalld**, there is no re-creating of all the rules, only the differences are applied. Consequently, **firewalld** can change the setting during runtime without existing connections being lost⁵. For the purposes of this reference architecture, **firewalld** is used and is the preferred method of implementing firewall rules. This section focuses on providing the details required to run **firewall-cmd** successfully for an Oracle Database environment. **Table 2.7.1: Firewall Port Settings** lists the enabled ports in this reference environment.

Port	Services	Protocol	Description
-	ssh	TCP	Secure Shell (SSH)
1521	-	TCP	Oracle Transparent Network Substrate (TNS) Listener default port
5500	-	TCP	EM Express 12c default port

Table 2.7.1: Firewall Port Settings

⁴ Linux man pages - `man (1) firewalld`

⁵ 4.5.3 Comparison of firewalld to system-config-firewall and iptables - https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/Security_Guide/sec-Using_Firewalls.html



2.8 Security: SELinux

Starting with Oracle 11g Release 2 version 11.2.0.3, *SELinux* is supported for Oracle database environments. The system in this reference environment runs with *SELinux* enabled and set to *ENFORCING* mode. **Table 2.8.1: SELinux Packages** lists the required *SELinux* packages.

Package	Version
selinux-policy	3.12.1-153
selinux-policy-targeted	3.12.1-153

Table 2.8.1: SELinux Packages



3 Reference Architecture Configuration Details

This reference architecture focuses on the deployment of Oracle RAC Database 12c Release 1 (12.1) with Oracle Automatic Storage Management (ASM) on Red Hat Enterprise Linux 7 x86_64. The configuration is intended to provide a comprehensive Red Hat | Oracle solution. The key solution components covered within this reference architecture consists of:

- Red Hat Enterprise Linux 7
- Oracle Grid Infrastructure 12c Release 1 (12.1.0.2.0)⁶
- Oracle Database 12c Release 1 Software Installation (12.1.0.2.0)⁷
- Deploying an Oracle Database 12c Release 1 (12.1.0.2.0) with shared SAN disks
- Enabling Security-Enhanced Linux (*SELinux*)
- Configuring Device Mapper Multipathing
- Using *udev* Rules instead of Oracle ASMLib⁸

3.1 Setting OS Hostname

A unique host name is required for the installation of Oracle RAC Database 12c Release 1 (12.1) The host names within this reference environment is: *db-oracle-node1* and *db-oracle-node2*. To set the host name, please follow the instructions below.

Set the host name via the use of the **hostnamectl** command. An example of setting *db-oracle-node1* host name is shown below.

```
# hostnamectl set-hostname db-oracle-node1.cloud.lab.eng.bos.redhat.com
```

Verify the status:

```
# hostnamectl status
  Static hostname: db-oracle-node1.cloud.lab.eng.bos.redhat.com
        Icon name: computer
        Chassis: n/a
        Machine ID: eb00fab21b644a45ba9d2532554f7b49
        Boot ID: 503421cddcd54634ba926daca7e08ccb
  Operating System: Red Hat Enterprise Linux
        CPE OS Name: cpe:/o:redhat:enterprise_linux:7.0:GA:server
        Kernel: Linux 3.10.0-123.el7.x86_64
  Architecture: x86_64
```

⁶ Oracle Grid Infrastructure 12c (12.1.0.2.0) or higher supported version with Red Hat Enterprise Linux 7

⁷ Oracle Database 12c (12.1.0.2.0) or higher supported version with Red Hat Enterprise Linux 7

⁸ No Oracle ASMLib for Red Hat Enterprise Linux 7 is available at the time of this writing. My Oracle Support Article ID: 1089399.1



3.2 Network Configuration

The network configuration focuses on the proper setup of public and private network interfaces along with the DNS configuration for the Single Client Access Name (SCAN). The public bonded network interface provides an Oracle environment with high availability in case of a network interface failure. The High Availability Internet Protocol (HAIP) provides the private network interfaces with failover and load balancing across each private network interface. SCAN provides the Oracle RAC Database 12c Release 1 (12.1) environment a single name that can be used by any client trying to access an Oracle Database within the cluster.

3.2.1 Configuring `/etc/resolv.conf` file

The resolver is a set of routines in the C library that provides access to the Internet Domain Name System (DNS). The resolver configuration file contains information that is read by the resolver routines the first time they are invoked by a process. The file is designed to be human readable and contains a list of keywords with values that provide various types of resolver information.⁹ The `/etc/resolv.conf` file for this reference environment consists of two configuration options: `nameserver` and `search`. The search option is used to search for a host name that is part of a particular domain. The `nameserver` option is the IP address of the name server the systems (`db-oracle-node1`, `db-oracle-node2`) must query. If more than one `nameserver` is listed, the resolver library queries them in order. An example of the `/etc/resolv.conf` file used on each node of the reference environment is shown below.

```
cat /etc/resolv.conf
search cloud.lab.eng.bos.redhat.com
nameserver 10.19.143.247
nameserver 10.19.143.248
nameserver 10.19.255.2
```

3.2.2 Public Network Configuration

The public network configuration consists of two network interfaces bonded together to provide high availability. The example below shows how to bond physical interfaces `ens10f0` and `ens10f1` with a bond device labeled `bond0`. `NetworkManager` will be used to setup the bond device.

NOTE: `NetworkManager` is optional to setup bonding or Ethernet device interfaces, however, it is the method used within this reference environment.

Check the status of `NetworkManager`:

```
# systemctl status NetworkManager.service
NetworkManager.service - Network Manager
   Loaded: loaded (/usr/lib/systemd/system/NetworkManager.service; enabled)
   Active: active (running) since Tue 2014-10-28 11:28:22 CDT; 1 day 22h ago
   Main PID: 1989 (NetworkManager)
   CGroup: /system.slice/NetworkManager.service
           \u2514\u25001989 /usr/sbin/NetworkManager --no-daemon
```

⁹ Linux man pages – man resolv.conf



Enable *NetworkManager* to start upon boot time, if not already:

```
# systemctl enable NetworkManager.service
```

Startup the *NetworkManager* text user interface (TUI) **nmtui** to setup the bonded devices:

```
# nmtui
```

1. Select **Edit a connection**, and press **OK**.
2. Select **Add**, hit Enter.
3. Within the **New Connection** window, highlight **Bond**, and press **Create**.
4. Change **Profile Name** if desired, otherwise within *Device* type *bond0*
5. Select **Add**, and within the **New Connection** window, select **Ethernet**, and press **Create**.
6. Within the window, change **Profile name** if desired, and enter the appropriate network device, i.e. *ens10f0* and click **OK**.
7. Repeat step 6 for any additional network devices to add to the bond device.
8. Change *Mode* to **Active-Backup** and select one of the Ethernet devices as the **Primary**, such as *ens10f0*
9. Change **IPv4 Configuration** from **<Automatic>** to **<Manual>** and press **Show**.
10. Add **Addresses**, **Gateway**, **DNS servers**, **Search domains** as required.
11. Ensure **Automatically connect** and **Available to all users** is checked and press **OK**.
12. Exit the **nmtui** utility by pressing **Quit**.

Figure 3.2.2.1: Bond Device using Network Manager In the following page shows the setup of the bond device.

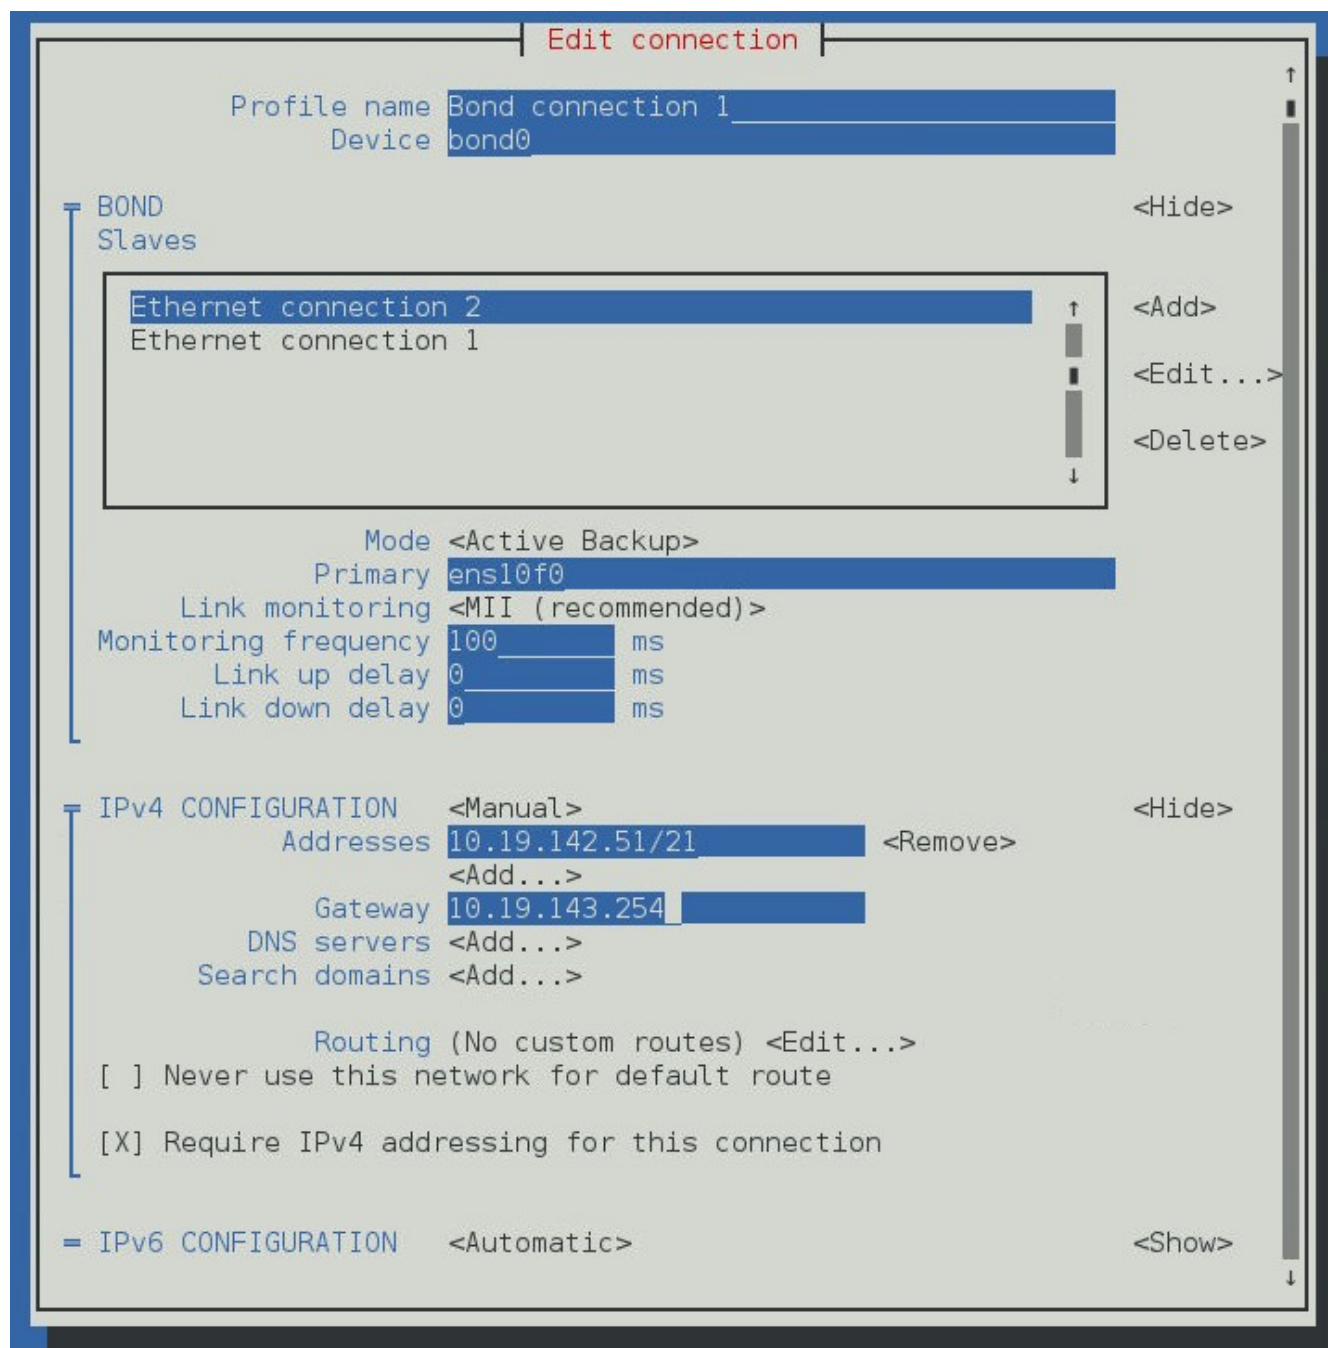


Figure 3.2.2.1: Bond Device using Network Manager

Restart *NetworkManager* service:

```
# systemctl restart NetworkManager.service
```



Once the bond0 device is configured on each node, use the **ping** command to verify connectivity as follows:

On node one labeled *db-oracle-node1*,

```
# ping db-oracle-node2
PING db-oracle-node2.cloud.lab.eng.bos.redhat.com (10.19.142.52) 56(84)
bytes of data.
64 bytes from db-oracle-node2.cloud.lab.eng.bos.redhat.com (10.19.142.52):
icmp_seq=1 ttl=64 time=0.150 ms
```

On node two labeled *db-oracle-node2*,

```
# ping db-oracle-node1
PING db-oracle-node1.cloud.lab.eng.bos.redhat.com (10.19.142.51) 56(84)
bytes of data.
64 bytes from db-oracle-node1.cloud.lab.eng.bos.redhat.com (10.19.142.51):
icmp_seq=1 ttl=64 time=0.151 ms
```

NOTE: Please ensure a DNS entry that resolves to the appropriate *hostname*. This reference architecture resolves the following IP address to the following host names:

IP	Hostname
10.19.142.51	db-oracle-node1
10.19.142.52	db-oracle-node2

Table 3.2.2.1: Public IP & Hostnames

3.2.3 Configure SCAN via DNS

SCAN provides a single name in which a client server can use to connect to a particular Oracle database. The main benefit of SCAN is the ability to keep a client connection string the same even if changes within the Oracle RAC Database environment occur, such as adding or removing of nodes within the cluster. The reason this works is because every client connection sends a request to the SCAN Listener, which then reroutes the traffic to an available VIP Listener within the Oracle RAC cluster to establish a database connection. The setup of SCAN requires the creation of a single name, no longer than 15 characters in length not including the domain suffix, resolving to three IP addresses using a round-robin algorithm from the DNS server. SCAN must reside in the same subnet as the public network within the Oracle RAC Database cluster and be resolvable without the domain suffix. Within the reference environment, the domain is *cloud.lab.eng.bos.redhat.com* and SCAN name is *db-oracle-scan*.

An example DNS entry for the SCAN is as follows:

```
db-oracle-scan    IN A  10.19.142.53
                  IN A  10.19.142.54
                  IN A  10.19.142.55
```



An example of the DNS entry for the SCAN to enable reverse lookups is as follows:

```
53    IN    PTR    db-oracle-scan.cloud.lab.eng.bos.redhat.com
54    IN    PTR    db-oracle-scan.cloud.lab.eng.bos.redhat.com
55    IN    PTR    db-oracle-scan.cloud.lab.eng.bos.redhat.com
```

On each node within the Oracle RAC cluster, verify the SCAN configuration within the DNS server is setup properly using the **nslookup** and **host** command as follows:

```
# nslookup db-oracle-scan
Server:      10.19.143.247
Address:     10.19.143.247#53

Name: db-oracle-scan.cloud.lab.eng.bos.redhat.com
Address: 10.19.142.55
Name: db-oracle-scan.cloud.lab.eng.bos.redhat.com
Address: 10.19.142.53
Name: db-oracle-scan.cloud.lab.eng.bos.redhat.com
Address: 10.19.142.54
```

```
# host db-oracle-scan
db-oracle-scan.cloud.lab.eng.bos.redhat.com has address 10.19.142.53
db-oracle-scan.cloud.lab.eng.bos.redhat.com has address 10.19.142.54
db-oracle-scan.cloud.lab.eng.bos.redhat.com has address 10.19.142.55
```

On each node within the Oracle RAC cluster, verify the SCAN configuration reverse lookup is setup properly using the **nslookup** and **host** command as follows:

```
# nslookup 10.19.142.53
Server:      10.19.143.247
Address:     10.19.143.247#53
53.142.16.10.in-addr.arpa  name = db-oracle-
scan.cloud.lab.eng.bos.redhat.com.
```

Repeat the above step for the reverse lookup on the remaining IP addresses (10.19.142.54, 10.19.142.55) used for the SCAN.

NOTE: The reference environment resolves the following IP address to the following host name:

IP	Hostname
10.19.142.53	db-oracle-scan
10.19.142.54	
10.19.142.55	

Table 3.2.3.1: SCAN IP & Hostname

For more information on SCAN, please refer to Oracle's documentation¹⁰.

¹⁰ Oracle Single Client Access Name (SCAN) -

<http://www.oracle.com/technetwork/products/clustering/overview/scan-129069.pdf>



3.2.4 Configure Virtual IP (VIP) via DNS

The virtual IP is an IP address assigned to each node within an Oracle RAC Database environment with the IP address residing in the public subnet. During the installation of the Oracle Grid Infrastructure, each VIP Listener registers with every SCAN Listener. The reason is because when a client sends a request, the SCAN Listener routes the incoming traffic to one of the VIP Listeners within the Oracle RAC Database cluster. If a client connection string uses the VIP to talk directly to the VIP Listener (as done in prior versions), every time changes to the Oracle RAC Database environment are made, such as adding or removing nodes within the cluster, the client connection string would require updating. Due to this, Oracle recommends always using the SCAN for the client connection string.

An example DNS entry for the VIPs is as follows:

```
db-oracle-node1-vip    IN    A    10.19.142.56
db-oracle-node2-vip    IN    A    10.19.142.57
```

On each node within the Oracle RAC cluster, verify the VIP address for *db-oracle-node1-vip* and *db-oracle-node2-vip* within the DNS server is setup properly using the **nslookup** and **host** command. An example of checking *db-oracle-node1-vip* can be seen below.

```
# nslookup db-oracle-node1-vip
Server:          10.19.143.247
Address:         10.19.143.247#53

Name: db-oracle-node1-vip.cloud.lab.eng.bos.redhat.com
Address: 10.19.142.56
```

```
# host db-oracle-node1-vip
db-oracle-node1-vip.cloud.lab.eng.bos.redhat.com has address 10.19.142.56
```

An example of the DNS entry for the SCAN to enable reverse lookups is as follows:

```
56    IN    PTR    db-oracle-node1-vip.cloud.lab.eng.bos.redhat.com
57    IN    PTR    db-oracle-node2-vip.cloud.lab.eng.bos.redhat.com
```



On each node within the Oracle RAC Database cluster, verify the VIP address reverse lookup for both VIP addresses (10.19.142.56 and 10.19.142.57) is setup properly using the **nslookup** and **host** command. An example is shown using VIP address 10.19.142.56 below.

```
# nslookup 10.19.142.56
Server:      10.19.143.247
Address:     10.19.143.247#53

56.142.16.10.in-addr.arpa  name = db-oracle-node1-
vip.cloud.lab.eng.bos.redhat.com.
```

```
# host 10.19.142.56
56.142.16.10.in-addr.arpa domain name pointer db-oracle-node1-
vip.cloud.lab.eng.bos.redhat.com.
```

NOTE: The VIP address should provide a *Destination Host Unreachable* response if an attempt to **ping** the VIP or VIP host name is attempted.

This reference environment resolves the following Virtual IP addresses to the following host names:

	Hostname
10.19.142.56	db-oracle-node1-vip
10.19.142.57	db-oracle-node2-vip

Table 3.2.4.1: Virtual IP & Hostnames

3.2.5 Private Network Configuration

The private network configuration consists of two network interfaces *ens10f2* and *ens10f3*. The private network is used to provide interconnect communication between all the nodes in the cluster. This is accomplished via Oracle's Redundant Interconnect, also known as Highly Available Internet Protocol (HAIP), that allows the Oracle Grid Infrastructure to activate and load balance traffic on up to four Ethernet devices for private interconnect communication. The example below shows how to setup physical interfaces *ens10f2* and *ens10f3* to be used with HAIP.

On each node, as the root user, startup the *NetworkManager* text user interface (TUI) **nmtui** to setup the bonded devices:

- ```
nmtui
```
1. Select **Edit a connection**, and press **OK**.
  2. Highlight **ens10f2**, and select **Edit...**
  3. Within the **Edit Connection** window, press **show** in **Ethernet** and set **MTU** to 9000 bytes, change **IPv4 Configuration** from **<Automatic>** to **<Manual>** and press **Show**.
  4. Add **Addresses**, **Gateway**, **DNS servers**, **Search domains** as required.





5. Ensure **Automatically connect** and **Available to all users** is checked and press **OK**.
6. Exit the **nmtui** utility by pressing **Quit**.
7. Repeat the above steps for **ens10f3** and any other private network devices.

The screenshot shows the 'Edit connection' window in Network Manager. The profile name is 'ens10f2' and the device is '00:25:B3:A9:B0:02 (ens10f2)'. The 'ETHERNET' section is expanded, showing 'Cloned MAC address' and 'MTU 9000 bytes'. The 'IPv4 CONFIGURATION' section is also expanded, showing 'Manual' mode with the address '192.11.142.52/24', a gateway, and DNS servers. The 'IPv6 CONFIGURATION' section is collapsed, showing 'Automatic' mode. At the bottom, the 'Automatically connect' and 'Available to all users' checkboxes are both checked. The window has '<Cancel>' and '<OK>' buttons at the bottom right.

```

Edit connection

Profile name ens10f2
Device 00:25:B3:A9:B0:02 (ens10f2)

ETHERNET <Hide>
 Cloned MAC address
 MTU 9000 bytes

IPv4 CONFIGURATION <Manual> <Hide>
 Addresses 192.11.142.52/24 <Remove>
 <Add...>
 Gateway
 DNS servers <Add...>
 Search domains <Add...>
 Routing (No custom routes) <Edit...>
 [] Never use this network for default route
 [X] Require IPv4 addressing for this connection

= IPv6 CONFIGURATION <Automatic> <Show>
[X] Automatically connect
[X] Available to all users

<Cancel> <OK>

```

**Figure 3.2.5.1: Setting Private Network using Network Manager**

**NOTE:** The MTU size is set to 9000 for the enablement of jumbo frames. Ensure jumbo frames are enabled on the private Ethernet switches.





After the setup of Ethernet devices *ens10f2* and *ens10f3* on each node, restart the *NetworkManager* service:

```
systemctl restart NetworkManager.service
```

Ensure all private Ethernet interfaces are set to different subnets on each node. If different subnets are not used and connectivity is lost, this can cause a node reboot within the cluster. For the reference environment, subnets 192.11.142.0/24 and 192.12.142.0/24 are used on each node within the Oracle RAC Database 12c cluster.

Verify connectivity on each node using the **ping** command.

On node one labeled *db-oracle-node1*,

```
ping 192.11.142.52
```

```
ping 192.12.142.52
```

On node two labeled *db-oracle-node2*,

```
ping 192.11.142.51
```

```
ping 192.12.142.51
```

| IP            | Ethernet Interface | Host            |
|---------------|--------------------|-----------------|
| 192.11.142.51 | ens10f2            | db-oracle-node1 |
| 192.12.142.51 | ens10f3            | db-oracle-node1 |
| 192.11.142.52 | ens10f2            | db-oracle-node2 |
| 192.12.142.52 | ens10f3            | db-oracle-node2 |

**Table 3.2.5.1: Private IP, Ethernet Interfaces, & Host**



## 3.3 OS Configuration

### 3.3.1 Using the Red Hat Subscription Manager

The **subscription-manager** command is used to register systems to the Red Hat Network (RHN) and to manage the subscription entitlements for the systems. The **--help** option can be specified on the command line to query the command for the available options. If the **--help** option is issued along with a command directive, then options available for the specific command directive are listed.

To use Red Hat Subscription Management for providing packages to a system, the system must first register with the service. To register a system, use the **subscription-manager** command and pass the register command directive. If the **--username** and **--password** options are specified, then the command will not prompt for the RHN Network authentication credentials.

An example of registering a system using **subscription-manager** is shown below.

```
subscription-manager register --username [User] --password '[Password]'
The system has been registered with id: abcd1234-ab12-ab12-ab12-481ba8187f60
```

After a system is registered, it must be attached to an entitlement pool. For the purposes of this reference environment, the **Red Hat Enterprise Linux Server** is the pool chosen. To identify and subscribe to the **Red Hat Enterprise Linux Server** entitlement pool, the following command directives are required.

```
subscription-manager list --available | grep -A8 "Red Hat Enterprise Linux
Server"
Subscription Name: Red Hat Enterprise Linux Server, Standard (8 sockets)
(Unlimited guests)
Provides: Red Hat Beta
 Oracle Java (for RHEL Server)
 Red Hat Enterprise Linux Server
 Red Hat Software Collections Beta (for RHEL Server)
SKU: RH0186633
Contract: 10541483
Pool ID: 8a85f98148a8c1a00148a961a0701b1b
Available: 47
Suggested: 1
Service Level: STANDARD
Service Type: L1-L3
```

```
subscription-manager attach --pool 8a85f98148a8c1a00148a961a0701b1b
Successfully attached a subscription for: Red Hat Enterprise Linux Server,
Standard (8 sockets) (Unlimited guests)
```



The Red Hat Enterprise Linux supplementary repository is part of subscribing to the Red Hat Enterprise Linux Server entitlement pool, however, it is disabled by default. To enable the supplementary repository, use the **subscription-manager** command.

**NOTE:** The following step is required in order to install the **compat-libstdc++-33** package that is required for a successful Oracle 12c Database install on Red Hat Enterprise Linux 7. The package is only available in the **rhel-7-server-optional-rpms** repository.

```
subscription-manager repos --enable=rhel-7-server-optional-rpms
Repo 'rhel-7-server-optional-rpms' is enabled for this system.
```

For more information on the use of Red Hat Subscription Manager, please visit the Red Hat Subscription management documentation<sup>11</sup>.

---

11 Red Hat Subscription Management - [https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Subscription\\_Management/1/html-single/RHSM/](https://access.redhat.com/documentation/en-US/Red_Hat_Subscription_Management/1/html-single/RHSM/)



### 3.3.2 NTP Configuration

The **ntpd** program is an operating system daemon which sets and maintains the system time, synchronizing with Internet standard time servers<sup>12</sup>. The **ntpd** program operates by exchanging messages with one or more configured servers at designated poll intervals.

To configure the **ntpd** daemon, follow the instructions below.

On all nodes within the Oracle RAC Database cluster,

1. If not installed, install **ntp** via **yum** as follows:

```
yum install ntp
```

2. Edit the `/etc/ntp.conf` file with a text editor such as **vi**.

```
vi /etc/ntp.conf
```

3. Locate the following public server pool section, and modify to include the appropriate NTP servers. For the purposes of this reference environment, only one NTP server is used, but three is recommended. The **iburst** option is added to speed up the time in which it takes to properly sync with the NTP servers.

```
Use public servers from the pool.ntp.org project.
Please consider joining the pool (http://www.pool.ntp.org/join.html).
server 10.5.26.10 iburst
```

4. Save all the changes within the `/etc/ntp.conf` file
5. The following `-x` option within the `/etc/sysconfig/ntpd` file needs to be added to prevent the time synched by the NTP daemon to be adjusted backward.

```
Command line options for ntpd
OPTIONS="-x -g"
```

6. Start the **ntpd** daemon via the command:

```
systemctl start ntpd.service
```

7. Ensure that the **ntpd** daemon is started when the system is booted.

```
systemctl enable ntpd.service
ln -s '/usr/lib/systemd/system/ntpd.service' '/etc/systemd/system/multi-
user.target.wants/ntpd.service'
```

---

<sup>12</sup> ntpd – Network Time Protocol (NTP) daemon man page – man ntpd (8)



8. Verify the **ntpd** daemon status.

```
systemctl status ntpd.service
ntpd.service - Network Time Service
Loaded: loaded (/usr/lib/systemd/system/ntpd.service; enabled)
Active: active (running) since Tue 2015-01-27 10:51:34 CST; 2s ago
Process: 7743 ExecStart=/usr/sbin/ntpd -u ntp:ntp $OPTIONS (code=exited,
status=0/SUCCESS)
Main PID: 7744 (ntpd)
CGroup: /system.slice/ntpd.service
 \u2514\u25007744 /usr/sbin/ntpd -u ntp:ntp -x -g
```

### 3.3.3 Oracle Database 12c Release 1 (12.1) Package Requirements

A specific set of packages is required to properly deploy Oracle RAC Database 12c Release 1 (12.1) on Red Hat Enterprise Linux 7 (x86\_64). The number of installed packages required varies depending on whether a default or minimal installation of Red Hat Enterprise Linux 7 (x86\_64) is performed. For the purposes of this reference environment, a minimal Red Hat Enterprise Linux 7 installation is performed to reduce the number of installed packages. A sample kickstart file as been provided within **Appendix M Configuration Files**. Red Hat Enterprise Linux 7 installation requires the following group packages:

| Required Group Packages |
|-------------------------|
| @Base                   |
| @Core                   |

**Table 3.3.3.1: Group Packages**



Oracle Grid Infrastructure 12.1 and Oracle Database 12.1 required x86\_64 RPM packages<sup>13</sup>:

| Required Packages   |                |
|---------------------|----------------|
| binutils            | libX11         |
| compat-libcap1      | libXau         |
| compat-libstdc++-33 | libaio         |
| gcc                 | libaio-devel   |
| gcc-c++             | libdmx         |
| glibc-devel         | glibc          |
| ksh                 | make           |
| libgcc              | sysstat        |
| libstdc++           | xorg-x11-utils |
| libstdc++-devel     | xorg-x11-xauth |
| libXext             | libXv          |
| libXtst             | libXi          |
| libxcb              | libXt          |
| libXmu              | libXxf86misc   |
| libXxf86dga         | LibXxf86vm     |
| nfs-utils           |                |

**Table 3.3.3.2: Required Packages**

After the installation of Red Hat Enterprise Linux 7 is completed, create a file, *req-rpm.txt*, that contains the name of each RPM package listed above on a separate line. For simplicity, this *req-rpm.txt* file is included in **Appendix F Oracle Database Package Requirements Text File**.

Use the **yum** package manager to install the packages and any of their dependencies with the following command:

```
yum install `awk '{print $1}' ./req-rpm.txt`
```

A minimum installation of Red Hat Enterprise Linux 7 does not install the *X Window System* server package, but only the required *X11* client libraries. In order to run the Oracle Universal Installer (OUI), a system with the *X Window System* server package installed is required. Using a system with *X Window System* installed, **SSH** into the Oracle Database server with the *-Y* option to ensure trusted *X11* forwarding is set. The command is as follows:

```
ssh -Y db-oracle-node1
```

Alternatively, if a system with the *X Window System* server package is unavailable, install the *X Window System* server package directly on the Oracle Database Server.

<sup>13</sup> Packages for Red Hat Enterprise Linux 7 -  
[https://docs.oracle.com/database/121/LADBI/pre\\_install.htm#LADBI7534](https://docs.oracle.com/database/121/LADBI/pre_install.htm#LADBI7534)



```
yum groupinstall "X Window System"
```

### 3.3.4 Configuring Security-Enhanced Linux (SELinux)

*SELinux* is an implementation of a mandatory access control (MAC) mechanism developed by the National Security Agency (NSA). The purpose of *SELinux* is to apply rules on files and processes based on defined policies. When policies are appropriately defined, a system running *SELinux* enhances application security by determining if an action from a particular process should be granted thus protecting against vulnerabilities within a system. The implementation of Red Hat Enterprise Linux 7 enables *SELinux* by default and appropriately sets it to the default setting of *ENFORCING*. It is highly recommended that *SELinux* be kept in *ENFORCING* mode when running Oracle RAC Database 12c Release 1 (12.1).

Verify that *SELinux* is running and set to *ENFORCING*:

On each node within the Oracle RAC cluster, as the root user,

```
sestatus
SELinux status: enabled
SELinuxfs mount: /sys/fs/selinux
SELinux root directory: /etc/selinux
Loaded policy name: targeted
Current mode: enforcing
Mode from config file: enforcing
Policy MLS status: enabled
Policy deny_unknown status: allowed
Max kernel policy version: 28
```

If the system is running in *PERMISSIVE* or *DISABLED* mode, modify the */etc/selinux/config* file and set *SELinux* to enforcing as shown below.

```
SELINUX=enforcing
```

The modification of the */etc/selinux/config* file takes effect after a reboot. To change the setting of *SELinux* immediately without a reboot, run the following command:

```
setenforce 1
```

For more information on Security-Enhanced Linux, please visit the [Red Hat Enterprise Linux 7 Security-Enhanced Linux User Guide](#)

### 3.3.5 Configuring Firewall Settings

Firewall access and restrictions play a critical role in securing your Oracle RAC Database 12c Release 1 (12.1) environment. It is not uncommon for corporations to be running hardware based firewalls to protect their corporate networks. Due to this, enabling firewall may not be required. However, this reference environment demonstrates how to successfully implement firewall settings for an Oracle RAC Database environment. The firewall rules described below only apply to the public network. Oracle recommends the private network should not have any firewall rules as this can cause issues with the installation of Oracle RAC Database, as well



as, disruption with the Oracle RAC Database private interconnect. It is highly recommended that the private network be isolated and communicate only between nodes locally. Red Hat Enterprise Linux 7 introduces the use of **firewalld**, a dynamic firewall daemon, instead of the traditional **iptables service**. **firewalld** works by assigning network zones to assign a level of trust to a network and its associated connections and interfaces<sup>14</sup>. The key difference and advantage of **firewalld** over **iptables service** is that it does not require flushing of old firewall rules to apply the new firewall rules. **firewalld** changes the settings during runtime without losing existing connections<sup>14</sup>. With the implementation of **firewalld**, the **iptables service** configuration file `/etc/sysconfig/iptables` does not exist. It is recommended that the firewall settings be configured to permit access to the Oracle Database network ports only from authorized database or database-management clients. For example, in order to allow access to a specific database client with an IP address of 10.19.142.54 and to make requests to the database server via SQL\*Net using Oracle's TNS (Transparent Network Substrate) Listener (default port of 1521), the following permanent firewall rule within the public zone must be added to the **firewalld** configuration.

On all nodes within the Oracle RAC cluster unless otherwise specified,

```
firewall-cmd --permanent --zone=public --add-rich-rule="rule family="ipv4"
source address="10.19.142.54" port protocol="tcp" port="1521" accept"
success
```

Likewise, if a particular database client with an IP address of 10.19.142.54 required access to the web-based EM Express that uses the default port of 5500, the following firewall rich rule must be added using the **firewall-cmd** command.

```
firewall-cmd --permanent --zone=public --add-rich-rule="rule family="ipv4"
source address="10.19.142.54" port protocol="tcp" port="5500" accept"
success
```

Ensure the firewall allows all traffic from the private network by accepting all traffic (trusted) from the private Ethernet interfaces *ens10f2* and *ens10f3* from all nodes within the Oracle RAC cluster. It is highly recommended that the private network be isolated and communicate only between nodes locally.

```
firewall-cmd --permanent --zone=trusted --change-interface=ens10f2
success

firewall-cmd --permanent --zone=trusted --change-interface=ens10f3
success
```

The following rules are added to satisfy the Oracle Installer's prerequisites. Once Oracle installation is complete, the following rules can be removed. Steps for removal shown upon the completion of the Oracle RAC Database installation in **Section 6.1.2 Removal of firewalld Trusted Source Address**

---

<sup>14</sup> Red Hat Enterprise Linux 7 – Using Firewalls: [https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/7/html/Security\\_Guide/sec-Using\\_Firewalls.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/Security_Guide/sec-Using_Firewalls.html)





On node one of the Oracle RAC cluster, add the public source IP of all remaining nodes. This reference environment only adds the public IP of node two of the Oracle RAC cluster as it is a two-node Oracle RAC Database environment.

```
firewall-cmd --permanent --zone=trusted --add-source=10.19.142.52/21
```

On node two of the Oracle RAC cluster, add the public source IP of all remaining nodes. This reference environment only adds the public IP of node one of the Oracle RAC cluster as it is a two-node Oracle RAC Database environment.

```
firewall-cmd --permanent --zone=trusted --add-source=10.19.142.51/21
```

Once the rules have been added, run the following command to activate:

```
systemctl restart firewalld.service
```

To verify the port 1521 & 5500 has been added and database client with IP address of 10.19.142.54 has been properly added, run the following command:

```
firewall-cmd --zone=public --list-all
public (default, active)
 interfaces: bond0 ens10f0 ens10f1
 sources:
 services: dhcpv6-client ssh
 ports:
 masquerade: no
 forward-ports:
 icmp-blocks:
 rich rules:
 rule family="ipv4" source address="10.19.142.54" port port="1521"
 protocol="tcp" accept
 rule family="ipv4" source address="10.19.142.54" port port="5500"
 protocol="tcp" accept
```

To verify the firewall rules being applied to the *trusted* zone for the private Ethernet interfaces, and temporarily for the source public IP run the following command:

Example of *db-oracle-node1*

```
firewall-cmd --zone=trusted --list-all
trusted (active)
 interfaces: ens10f2 ens10f3
 sources: 10.19.142.52/21
 services:
 ports:
 masquerade: no
 forward-ports:
 icmp-blocks:
 rich rules:
```

**NOTE:** A full listing of all the firewall settings for this reference architecture can be found at **Appendix D Firewall Configuration File.**



For more information of Red Hat Enterprise Linux 7 **firewalld**:

[https://access.redhat.com/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/7/html/Security\\_Guide/sec-Using\\_Firewalls.html](https://access.redhat.com/documentation/en-US/Red_Hat_Enterprise_Linux/7/html/Security_Guide/sec-Using_Firewalls.html)

### 3.3.6 Setting Virtual Memory

Tuning virtual memory requires the modification of five kernel parameters that affect the rate at which virtual memory is used within Oracle RAC databases. **It is important to note the recommended values listed are to be used as a starting point when setting virtual memory.**

A brief description<sup>15</sup> and recommended settings for the virtual memory parameters, as well as, the definition of dirty data are described below.

*SWAPPINESS*<sup>15</sup> - Starting with Red Hat Enterprise Linux 6.4 and above, the definition of swappiness has changed. Swappiness is defined as a value from 0 to 100 which controls the degree to which the system favors anonymous memory or the page cache. A high value improves file-system performance, while aggressively swapping less active processes out of memory. A low value avoids swapping processes out of memory, which usually decreases latency, at the cost of I/O performance. The default value is 60. Warning: Since Red Hat Enterprise Linux 6.4, setting swappiness to 0 will even more aggressively avoid swapping out, which increases the risk of out-of-memory (OOM) killing under strong memory and I/O pressure. To achieve the same behavior of swappiness as previous versions of Red Hat Enterprise Linux 6.4 in which the recommendation was to set swappiness to 0, set swappiness to the value of 1. The recommendation of swappiness for Red Hat Enterprise Linux 6.4 or higher running Oracle databases is now the value of 1.

*DIRTY\_DATA* – Dirty data is data that has been modified and held in the page cache for performance benefits. Once the data is flushed to disk, the data is clean.

*DIRTY\_RATIO*<sup>15</sup> – Contains, as a percentage of total system memory, the number of pages at which a process which is generating disk writes will itself start writing out dirty data. The default value is 20. The recommended value is 80. The reasoning behind increasing the value from the standard Oracle 15 recommendation to 80 is because dirty ratio defines the maximum percentage of total memory that can be filled with dirty pages before user processes are forced to write dirty buffers themselves during their time slice instead of being allowed to do more writes. All processes are blocked for writes when this occurs due to synchronous I/O, not just the processes that filled the write buffers. This can cause what is perceived as unfair behavior where a single process can hog all the I/O on a system. **As for all parameters in this reference architecture, there is no “one-size fits all” value and the recommendation should be only seen as a starting point.**

*DIRTY\_BACKGROUND\_RATIO*<sup>15</sup> – Contains, as a percentage of total system memory, the number of pages at which the background write back daemon will start writing out dirty data. The Oracle recommended value is 3.

**NOTE:** With the `dirty_background_ratio` set to 3 and `dirty_ratio` set to 80, the background write back daemon will start writing out the dirty data when it hits the 3% threshold

---

<sup>15</sup> RHEL7 Kernel Documentation (requires package `kernel-doc` to be installed) - `/usr/share/doc/kernel-doc-3.10.0/Documentation/sysctl/vm.txt`



asynchronously, however, none of that data is written synchronously until the `dirty_ratio` is 80% full which is what causes for all processes to be blocked for writes when this occurs.

`DIRTY_EXPIRE_CENTISECS`<sup>15</sup> - Defines when dirty in-memory data is old enough to be eligible for writeout. The default value is 3000, expressed in hundredths of a second. The Oracle recommended value is 500.

`DIRTY_WRITEBACK_CENTISECS`<sup>15</sup> - Defines the interval of when writes of dirty in-memory data are written out to disk. The default value is 500, expressed in hundredths of a second. The Oracle recommended value is 100.

Prior to making any changes to the `/etc/sysctl.conf`, ensure to create a backup as follows:

```
cp /etc/sysctl.conf /etc/sysctl.conf.bkup
```

```
vm.swappiness = 1
vm.dirty_background_ratio = 3
vm.dirty_ratio = 80
vm.dirty_expire_centisecs = 500
vm.dirty_writeback_centisecs = 100
```

For the changes to take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

**NOTE:** A full listing of all the kernel parameters modified within the `/etc/sysctl.conf` file can be found at **Appendix G Kernel Parameters**.

### 3.3.7 Setting Shared Memory (SHMMAX, SHMALL, SHMMNI)

Shared memory allows processes to communicate with each other by placing regions of memory into memory segments. In the case of Oracle, shared memory segments are used by the System Global Area (SGA) to store incoming data and control information. The size of Oracle's SGA impacts the amount of shared memory pages and shared memory segments to be set within a system. By default, Red Hat Enterprise Linux 7 provides a large amount of shared memory pages and segments. However, the appropriate allocation for a system depends on the size of the SGA within an Oracle RAC database instance.

In order to allocate the appropriate amount of shared memory pages and shared memory segments for a system running an Oracle database, the kernel parameters `SHMMAX`, `SHMALL`, and `SHMMNI` must be set.

`SHMMAX` – is the maximum size in bytes of a single shared memory segment

`SHMALL` – is the maximum total amount of shared memory pages

`SHMMNI` – is the maximum total amount of shared memory segments

A default installation of Red Hat Enterprise Linux 7 x86\_64 provides a maximum size of a single shared memory segment, `SHMMAX`, to 4294967295 bytes, equivalent to 4 GB -1 byte. This value is important since it regulates the largest possible size of one single Oracle SGA



shared memory segment. If the Oracle SGA is larger than the value specified by *SHMMAX* (default 4 GB-1 byte), then Oracle will be required to create multiple smaller shared memory segments to completely fit Oracle's SGA. This can cause a significant performance penalty, especially in NUMA environments. In an optimal NUMA configuration, a single shared memory segment for Oracle's SGA is created on each NUMA node. If *SHMMAX* is not properly sized and creates multiple shared memory segments, *SHMMAX* limitations may keep the system from evenly distributing the shared memory segments across each NUMA node. To determine a proper *SHMMAX* value, the Oracle Database parameter **SGA\_MAX\_SIZE** needs to be analyzed. If intending to set **SGA\_MAX\_SIZE** to a value larger than 4 GB, set *SHMMAX* to a size in bytes larger than the **SGA\_MAX\_SIZE**. If in doubt on how to properly set *SHMMAX*, a value of 4398046511104 (4 TB) can be used. This value is purposely set higher than the architectural memory limits to ensure that any Oracle SGA value set within an Oracle RAC database instance may fit in one single shared memory segment.

After calculating *SHMMAX*, the next step is to determine the maximum amount of shared memory pages (*SHMALL*) in a system by capturing system's page size in bytes. The following command can be used to obtain the system page size.

```
getconf PAGE_SIZE
4096
```

A default installation of Red Hat Enterprise Linux 7 x86\_64 provides a *SHMALL* value of 268435456, the equivalent of 1 TB in system pages. This is determined by the following formula:

$$SHMALL \text{ IN BYTES } * PAGE\_SIZE$$

From the example above, 268435456 bytes \* 4096 bytes = 1099511627776 bytes = 1 TB in system pages

To ensure an adequate amount of memory pages are allocated to a single Oracle SGA, it is recommended that the value of *SHMALL* be set to the following formula:

$$SHMMAX \text{ IN BYTES } / PAGE\_SIZE$$

For example, on a system with the *SHMMAX* set to 4398046511104, the value of *SHMALL* is calculated as follows:

```
echo "4398046511104 / 4096" | bc
1073741824
```

*SHMMNI* is the maximum total amount of shared memory segments. A default installation of Red Hat Enterprise Linux 7 x86\_64 provides a *SHMMNI* default value of 4096. By optimizing the *SHMMAX* value with one shared memory segment per Oracle SGA, this parameter reflects the maximum number of Oracle and ASM instances that can be started on a system. Oracle recommends the value of *SHMMNI* to be left at the default value of 4096.

Snippet of the */etc/sysctl.conf* file:

```
kernel.shmmax = 4398046511104
kernel.shmall = 1073741824
kernel.shmmni = 4096
```



**NOTE:** If the current value found within */etc/sysctl.conf* for any parameter is higher than the value calculated for *SHMMAX* and *SHMALL* on any nodes within the Oracle RAC cluster, do not change the value found within the */etc/sysctl.conf* file.

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

**NOTE:** A full listing of all the kernel parameters modified within the */etc/sysctl.conf* file can be found at **Appendix G Kernel Parameters**.

### 3.3.8 Setting Semaphores (SEMMSL, SEMMNI, SEMMNS)

Red Hat Enterprise Linux 7 provides semaphores for synchronization of information between processes. The kernel parameter *sem* is composed of four parameters:

*SEMMSL* – is defined as the maximum number of semaphores per semaphore set

*SEMMNI* – is defined as the maximum number of semaphore sets for the entire system

*SEMMNS* – is defined as the total number of semaphores for the entire system

**NOTE:** *SEMMNS* is calculated by *SEMMSL* \* *SEMMNI*

*SEMOPM* – is defined as the total number of semaphore operations performed per semop system call.

The following line is required within the */etc/sysctl.conf* file to provide sufficient semaphores for Oracle:

```
kernel.sem = 250 32000 100 128
```

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

### 3.3.9 Ephemeral Network Ports

Oracle recommends that the ephemeral default port range be set starting at 9000 to 65500. This ensures that all well known ports used by Oracle and other applications are avoided. To set the ephemeral port range, modify the */etc/sysctl.conf* file on each node of the Oracle RAC Database cluster and add the following line:

```
net.ipv4.ip_local_port_range = 9000 65500
```

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```



### 3.3.10 Optimizing Network Settings

Optimizing the network settings for the default and maximum buffers for the application sockets in Oracle is done by setting static sizes to *RMEM* and *WMEM*. The *RMEM* parameter represents the receive buffer size, while the *WMEM* represents the send buffer size. The recommended values by Oracle are configured within the */etc/sysctl.conf* file on each node of the Oracle RAC cluster.

```
net.core.rmem_default = 262144
net.core.rmem_max = 4194304
net.core.wmem_default = 262144
net.core.wmem_max = 1048576
```

In order to make the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

### 3.3.11 Setting NOZEROCONF

On each node within the Oracle RAC Database cluster, set the value of *NOZEROCONF* to *yes* within the */etc/sysconfig/network* file. Setting *NOZEROCONF* ensures that the route 169.254.0.0/16 is not added to the routing table.

```
NOZEROCONF=yes
```

### 3.3.12 Disabling the avahi-daemon service

From the Red Hat Customer Portal article: <https://access.redhat.com/solutions/25463>,

The Avahi website defines Avahi as: 'a system which facilitates service discovery on a local network. This helps to plug the laptop or computer into a network and instantly be able to view other people who you can chat with, find printers to print to, or find files being shared...' Avahi-daemon (on by default on Red Hat Enterprise Linux) can interfere with Oracle RAC's multicast heartbeat causing the application-layer interface to assume it has been disconnected on a node and reboot the node. It is not recommended to remove the package due to its many dependencies. The **avahi** libraries are being used by many packages on a system.

On each node within the Oracle RAC Database cluster, stop and disable the avahi services run the following commands:

```
systemctl stop avahi-dnssconfd
systemctl stop avahi-daemon
Warning: Stopping avahi-daemon, but it can still be activated by:
avahi-daemon.socket
```

To keep the avahi services off persistently across reboots, on each node run the following:

```
systemctl disable avahi-dnssconfd
systemctl disable avahi-daemon
rm '/etc/systemd/system/dbus-org.freedesktop.Avahi.service'
rm '/etc/systemd/system/multi-user.target.wants/avahi-daemon.service'
rm '/etc/systemd/system/sockets.target.wants/avahi-daemon.socket'
```



### 3.3.13 Increasing synchronous I/O Requests

The kernel parameter *FS.AIO-MAX-NR* sets the maximum number of current asynchronous I/O requests. Oracle recommends setting the value to 1048576. In order to add *FS.AIO-MAX-NR* to 1048576, modify the */etc/sysctl.conf* file on each node of the Oracle RAC cluster as follows:

```
fs.aio-max-nr = 1048576
```

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

### 3.3.14 Increasing File Handles

The kernel parameter *FS.FILE-MAX* sets the maximum number of open file handles assigned to the Red Hat Enterprise Linux 7 operating system. Oracle recommends that for each Oracle RAC database instance found within a system, allocate  $512 \times PROCESSES$  in addition to the open file handles already assigned to the Red Hat Enterprise Linux 7 operating system. *PROCESSES* within a database instance refers to the maximum number of processes that can be concurrently connected to the Oracle RAC database by the *oracle* user. The default value for *PROCESSES* is 300 for Oracle RAC Database 12c Release 1 (12.1). To properly calculate the *FS.FILE-MAX* for a system, first identify the current *FS.FILE-MAX* allocated to the system via the following command:

```
cat /proc/sys/fs/file-max
4897258
```

Next, add all the *PROCESSES* together from each Oracle RAC database instance found within the system and multiply by 512 as seen in the following command.

```
echo "512 * 300" | bc
153600
```

**NOTE:** To determine the current *PROCESSES* value, log into each Oracle RAC database instance and run the following command below. Since no Oracle RAC database has yet been created within this reference environment, the default value of 300 *PROCESSES* is used.

```
SQL> show parameter processes;
```

| NAME      | TYPE    | VALUE |
|-----------|---------|-------|
| processes | integer | 300   |

Finally, add the current *FS.FILE-MAX* value with the new value found from multiplying  $512 \times PROCESSES$  to attain the new *FS.FILE-MAX* value.

```
echo "4897258 + 153600" | bc
5050858
```





While the value of the `FS.FILE-MAX` parameter varies upon every environment, this reference environment sets the value at 6815744. Oracle recommends a value no smaller than 6815744. Due to the calculation in the above example equating to 5050858, the minimum Oracle recommended value is used. In order to add `FS.FILE-MAX` to 6815744, modify the `/etc/sysctl.conf` file on each node of the Oracle RAC cluster as follows:

```
fs.file-max = 6815744
```

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
```

**NOTE:** It is recommended to revisit the `FS.FILE-MAX` value if the `PROCESSES` value is increased for any Oracle RAC databases created.

**NOTE:** A full listing of all the kernel parameters modified within the `/etc/sysctl.conf` file can be found at **Appendix G Kernel Parameters**.

### 3.3.15 Kernel Panic On OOPS Parameter

The kernel parameter `kernel.panic_on_oops` controls the kernel's behavior when an oops or bug is encountered. By default the value is set to 1, however, the oracle installer requires it to be set within the `/etc/sysctl.conf` file. Add the following parameter to the `/etc/sysctl.conf` file as follows:

```
kernel.panic_on_oops = 1
```

In order for the changes take effect immediately, run the following command:

```
sysctl -p /etc/sysctl.conf
```

**NOTE:** A full listing of all the kernel parameters modified within the `/etc/sysctl.conf` file can be found at **Appendix G Kernel Parameters**.

### 3.3.16 Reverse Path Filtering

Red Hat Enterprise Linux 7 defaults to the use of Strict Reverse Path filtering. The reason strict mode is the default is to prevent IP spoofing from Distributed Denial-of-service (DDoS) attacks. However, having strict mode enabled on the private interconnect of an Oracle RAC database cluster may cause disruption of interconnect communication. It is recommended to set the `RP_FILTER` from strict mode to loose mode. Loosening the security on the private Ethernet interfaces should not be of concern as best practices recommend for an isolated private network that can only communicate between nodes specifically for Oracle's private interconnect.

To satisfy the Oracle Installer prerequisite, add the following modifications to the `/etc/sysctl.conf` on each node of the Oracle RAC cluster as follows:

```
net.ipv4.conf.ens10f2.rp_filter = 2
net.ipv4.conf.ens10f3.rp_filter = 2
```





In addition to the above, please include the following modification within the *98-oracle.conf* file located within the */etc/sysctl.d/* directory. The *98-oracle.conf* file can be found within

## Appendix H 98-oracle.conf

In order for the changes take effect immediately, run the following command on each node of the Oracle RAC cluster:

```
sysctl -p /etc/sysctl.conf
[Output Appreciated ...]
net.ipv4.conf.ens10f2.rp_filter = 2
net.ipv4.conf.ens10f3.rp_filter = 2

sysctl -p /etc/sysctl.d/98-oracle.conf
net.ipv4.conf.ens10f2.rp_filter = 2
net.ipv4.conf.ens10f3.rp_filter = 2
```

**NOTE:** Once the installation of Oracle is complete, the entries found within the */etc/sysctl.conf* file are removed and the values are only stored in the *98-oracle.conf* file within */etc/sysctl.d* directory. See **Section 6.1.1 Removal of Oracle Kernel Parameters from */etc/sysctl.conf*** for more details.

## 3.3.17 User Accounts & Groups

Prior to the installation of Oracle RAC Database 12c Release 1 (12.1), Oracle recommends the creation of a *grid* user for the Oracle Grid Infrastructure and an *oracle* user for the Oracle RAC Database software installed on the system. For the purposes of this reference environment, the Oracle Grid Infrastructure software owner is the user *grid* and the Oracle Database software owner is the user *oracle*. Each user is designated different groups to handle specific roles based on the software installed. However, the creation of separate users requires that both the *oracle* user and the *grid* user have a common primary group, the Oracle central inventory group (*OINSTALL*).

The following are the recommended system groups created for the installation of the Oracle RAC Database and part of the *oracle* user.

OSDBA group (*DBA*) – determines OS user accounts with DBA privileges

OSOPER group (*OPER*) – an optional group created to assign limited DBA privileges (*SYSOPER* privilege) to particular OS user accounts

With the release of Oracle RAC Database 12c (12.1), Oracle introduces three additional database groups that are optional, but recommended. The recommended following groups provide the ability to separate administrative privileges for common database operations and limit privileged access of the OSDBA group (*DBA*).

OSBACKUPDBA group (*BACKUPDBA*) – an optional group created to assign limited administrative privileges (*SYSBACKUP* privilege) to a user for database backup and recovery

OSDGDBA group (*DGDBA*) – an optional group created to assign limited administrative



privileges (*SYSDG* privilege) to a user for administering and monitoring Oracle Data Guard

OSKMDBA group (*KMDBA*) – an optional group created to assign limited administrative privileges (*SYSKM* privilege) to a user for encryption key management when using Oracle Wallet Manager

The following are the recommended system groups created for the installation of the Oracle Grid Infrastructure and part of the `grid` user

OSDBA group (*ASMDBA*) – provides administrative access to Oracle ASM instances

OSASM group (*ASMADMIN*) – provides administrative access for storage files via the *SYSASM* privilege

OSOPER group (*ASMOPER*) – an optional group created to assign limited DBA privileges with regards to ASM to particular OS user accounts

As the root user on each node, create the following user accounts, groups, and group assignments using a consistent UID and GID assignments across your organization:

```
groupadd --gid 54321 oinstall
groupadd --gid 54322 dba
groupadd --gid 54323 asmdba
groupadd --gid 54324 asmoper
groupadd --gid 54325 asmadmin
groupadd --gid 54326 oper
groupadd --gid 54327 backupdba
groupadd --gid 54328 dgdba
groupadd --gid 54329 kmdba

useradd --uid 54321 --gid oinstall --groups dba,oper,asmdba,asmoper,\
> backupdba,dgdba,kmdba oracle
passwd oracle

useradd --uid 54322 --gid oinstall --groups dba,asmadmin,asmdba,asmoper grid
passwd grid
```

Verify the `oracle` and `grid` user on each Oracle RAC database cluster node correctly displays the appropriate primary and supplementary groups via the commands:

```
id oracle
uid=54321(oracle) gid=54321(oinstall)
groups=54321(oinstall),54322(dba),54323(asmdba),54324(asmoper),54326(oper),
54327(backupdba),54328(dgdba),54329(kmdba)
id grid
uid=54322(grid) gid=54321(oinstall)
groups=54321(oinstall),54322(dba),54323(asmdba),54324(asmoper),54325(asmadmin)
```

### 3.3.18 Setting Shell Limits for the Grid and Oracle User

Oracle recommends the following settings for the soft and hard limits for the number of open file descriptors (`nfile`), number of processes (`nproc`), and size of the stack segment (`stack`) allowed by each user respectively. The purpose of setting these limits is to prevent a system wide crash that could be caused if an application, such as Oracle, were allowed to exhaust all of the OS resources under an extremely heavy workload.



Create a file labeled *99-grid-oracle-limits.conf* within */etc/security/limits.d/* as follows:

```
touch /etc/security/limits.d/99-grid-oracle-limits.conf
```

**NOTE:** The reason that the */etc/security/limits.conf* file is not directly modified is due to the order in which limit files are read in the system. After reading the */etc/security/limits.conf* file, files within the */etc/security/limits.d/* directory are read. If two files contain the same entry, the entry read last takes precedence. For more information visit Red Hat Article: “What order are the limit files in the limits.d directory read in?”<sup>16</sup>

Within the */etc/security/limits.d/99-grid-oracle-limits.conf* file, add the following soft and hard limits for the *oracle* and *grid* user:

```
oracle soft nproc 16384
oracle hard nproc 16384
oracle soft nofile 1024
oracle hard nofile 65536
oracle soft stack 10240
oracle hard stack 32768
grid soft nproc 16384
grid hard nproc 16384
grid soft nofile 1024
grid hard nofile 65536
grid soft stack 10240
grid hard stack 32768
```

Due to Bug 15971421<sup>17</sup>, the soft limit of *nproc* is not adjusted at runtime by the Oracle database. Due to this, if the *nproc* limit is reached, the Oracle database may become unstable and not be able to fork additional processes. A high enough value for the maximum number of concurrent threads for the given workload must be set, or use the hard limit value of 16384 as done above if in doubt.

**NOTE:** Modifications made to the *99-grid-oracle-limits.conf* file take effect immediately. However, please ensure that any previously logged in *oracle* or *grid* user sessions (if any) are logged out and logged back in for the changes to take effect.

As the root user, on each node of the Oracle RAC Database cluster, create a shell script labeled *oracle-grid.sh* within */etc/profile.d/* to create the *ulimits* for the *oracle* and *grid* user. The contents of the *oracle-grid.sh* script:

```
#Setting the appropriate ulimits for oracle and grid user
if [$USER = "oracle"]; then
 if [$SHELL = "/bin/ksh"]; then
 ulimit -u 16384
 ulimit -n 65536
 else
 ulimit -u 16384 -n 65536
 fi
fi
```

16 What order are the limits files in the limits.d directory read in? -

<https://access.redhat.com/site/solutions/199993>

17 Oracle Documentation 12.1 -

[http://docs.oracle.com/cd/E16655\\_01/readmes.121/e17908/toc.htm#BABICCB](http://docs.oracle.com/cd/E16655_01/readmes.121/e17908/toc.htm#BABICCB)



```
fi
if [$USER = "grid"]; then
 if [$SHELL = "/bin/ksh"]; then
 ulimit -u 16384
 ulimit -n 65536
 else
 ulimit -u 16384 -n 65536
 fi
fi
```

**NOTE:** While the *ulimit* values can be set directly within the */etc/profile* file, it is recommended to create a custom shell script within */etc/profile.d* instead. The *oracle-grid.sh* script can be downloaded from the **Appendix M Configuration Files**



As oracle and grid user, verify the *ULIMIT* values by running the following command:

```
ulimit -a
core file size (blocks, -c) 0
data seg size (kbytes, -d) unlimited
scheduling priority (-e) 0
file size (blocks, -f) unlimited
pending signals (-i) 385878
max locked memory (kbytes, -l) 14854144
max memory size (kbytes, -m) unlimited
open files (-n) 65536
pipe size (512 bytes, -p) 8
POSIX message queues (bytes, -q) 819200
real-time priority (-r) 0
stack size (kbytes, -s) 10240
cpu time (seconds, -t) unlimited
max user processes (-u) 16384
virtual memory (kbytes, -v) unlimited
file locks (-x) unlimited
```



## 3.4 Storage Configuration

The following storage configuration section describes the best practices for setting up device mapper multipath, the use of *udev* rules for ASM disk management, and the use of the *tuned* package for optimal performance.

### 3.4.1 Setting up DM-Multipath

*Device mapper multipath* provides the ability to aggregate multiple I/O paths to a newly created device mapper mapping to achieve high availability, I/O load balancing, and persistent naming. The following procedures provide the best practices to installing and configuring *device mapper multipath* devices.

**NOTE:** Ensure Oracle RAC database volumes are accessible via the operating system on all nodes within the Oracle RAC Database cluster prior to continuing with the section below.

**The following instructions are required on each node within the Oracle RAC Database 12c cluster.**

1. As the root user, install the *device-mapper-multipath* package using the **yum** package manager.

```
yum install device-mapper-multipath
```

2. Copy the *multipath.conf* file found within */usr/share/doc/device-mapper-multipath-0.4.9/* to */etc/*

```
cp /usr/share/doc/device-mapper-multipath-0.4.9/multipath.conf /etc/
```

3. Capture the scsi id of the local disk(s) on the system.

```
/usr/lib/udev/scsi_id --whitelisted --replace-whitespace --
device=/dev/sda
3600508b1001030353434363646301200
```

4. Uncomment and modify the blacklist section within the */etc/multipath.conf* file to include the *scsi id* of the local disk on the system. Once complete, save the changes made to the *multipath.conf* file.

```
blacklist {
 wwid 3600508b1001030353434363646301200
 devnode "^(ram|raw|loop|fd|md|dm-|sr|scd|st)[0-9]*"
 devnode "^hd[a-z]"
}
```

5. Start the *multipath* daemon.

```
systemctl start multipathd.service
```

6. Enable the *multipath* daemon to ensure it is started upon boot time.

```
systemctl enable multipathd.service
```



- Identify the *dm- device*, *size*, and *WWID* of each *device mapper* volume for Oracle data disks and recovery disks. In this example, volume *mpathb* is identified via the following command:

```
multipath -ll
```

```

multipath alias name world wide identifier (WWID) dm- device
 ↓ ↓ ↓
size → mpathb (3600c0ff000d7e7a89e85ac5101000000) dm-10 HP,MSA2324fc
 |
 | size=186G features='1 queue_if_no_path' hwhandler='0' wp=rw
 | +- policy='round-robin 0' prio=130 status=active
 | | - 3:0:0:3 sdd 8:48 active ready running
 | | - 3:0:1:3 sdh 8:112 active ready running
 | | - 4:0:0:3 sdt 65:48 active ready running
 | | - 4:0:1:3 sdx 65:112 active ready running
 | +- policy='round-robin 0' prio=10 status=enabled
 | | - 3:0:2:3 sdl 8:176 active ready running
 | | - 3:0:3:3 sdp 8:240 active ready running
 | | - 4:0:2:3 sdab 65:176 active ready running
 | | - 4:0:3:3 sdaf 65:240 active ready running

```

**Figure 3.4.1.1: Multipath Device (mpathb)**

**Figure 3.4.1.1: Multipath Device (mpathb)** properly identifies the current multipath alias name, size, WWID, and dm device. This information is required for the application of a custom alias to each volume as shown in step 9.

**NOTE:** Due to the size of some of the disks being the same, compare the *WWID* found on the system with the *WWID* located on the storage array to ensure the multipath alias name assigned corresponds with the name of the volume from the storage array.

- Uncomment the defaults section found within the */etc/multipath.conf* file.

```

defaults {
 udev_dir /dev
 polling_interval 10
 path_selector "round-robin 0"
 path_grouping_policy multibus
 getuid_callout "/lib/udev/scsi_id --whitelisted"
 --device="/dev/%n"
 prio alua
 path_checker readsector0
 rr_min_io 100
 max_fds 8192
 rr_weight priorities
 failback immediate
 no_path_retry fail
 user_friendly_names yes
}

```



**NOTE:** The standard options can be customized to better fit your storage array capabilities. Check with your storage vendor for details.

- Uncomment the multipath section found within the `/etc/multipath.conf` file and create an alias for each *device mapper* volume in order to enable persistent naming of those volumes. Once complete, save the changes made to the `multipath.conf` file. The output should resemble the example below. For reference, refer to the Oracle data volumes created for this reference environment displayed in **Table 2.4.3: Oracle OCR, Voting Disk, & Data File Sizes**.

```

multipaths {
 multipath {
 wwid 3600c0ff000d7e7a899d8515101000000
 alias db1
 }
 multipath {
 wwid 3600c0ff000dabfe5a7d8515101000000
 alias db2
 }
 multipath {
 wwid 3600c0ff000d7e7a8dbd8515101000000
 alias fra
 }
 multipath {
 wwid 3600c0ff000dabfe5f4d8515101000000
 alias redo
 }
}

```

- Restart the *device mapper multipath* daemon.

```
systemctl restart multipathd.service
```

- Verify the *device mapper* paths and aliases are displayed properly. Below is an example of one *device mapper* device labeled *fra*.

```

multipath -ll
fra (3600c0ff000d7e7a89e85ac5101000000) dm-10 HP,MSA2324fc
size=186G features='1 queue_if_no_path' hwhandler='0' wp=rw
|-+- policy='round-robin 0' prio=130 status=active
| |- 3:0:0:3 sdd 8:48 active ready running
| |- 3:0:1:3 sdh 8:112 active ready running
| |- 4:0:0:3 sdt 65:48 active ready running
| `-- 4:0:1:3 sdx 65:112 active ready running
`-+- policy='round-robin 0' prio=10 status=enabled
| |- 3:0:2:3 sdl 8:176 active ready running
| |- 3:0:3:3 sdp 8:240 active ready running
| |- 4:0:2:3 sdab 65:176 active ready running
| `-- 4:0:3:3 sdaf 65:240 active ready running

```

- Repeat all above steps for all nodes within the Oracle RAC Database cluster.





## 3.4.2 Partitioning Device Mapper Shared Disks

On the first node of the Oracle RAC Database cluster, create a partition for each device mapper volume (*ocrvote1*, *ocrvote2*, *ocrvote3*, *db1*, *db2*, *fra*, *redo*) using **parted** as displayed below for device *db1*.

```
parted /dev/mapper/db1 mklabel gpt mkpart primary "1 -1"
Information: You may need to update /etc/fstab.
```

Once the partition is created, a newly created device mapper device is created as *db1p1*.

```
ls -l /dev/mapper/db1p1
lrwxrwxrwx. 1 root root 8 Apr 16 15:15 /dev/mapper/db1p1 -> ../dm-11
```

**NOTE:** A newly created partition alias name ending in a **number** i.e. *db1* requires the alias name followed by *p1* such as *db1p1* seen above. If *p1* is missing, please run the following **kpartx** command to add the partition mappings to the device mapper disks.

```
kpartx -a /dev/mapper/db1
```

**NOTE:** If a newly created partition alias name ends in a **letter** i.e. *fra*, the alias name will be followed by just the partition number, i.e. *fra1*.

Once the partitions are created, on all the nodes in the Oracle RAC Database cluster, run the following **kpartx** command on each device mapper volume in order to update the partition table mapping.

**NOTE:** If the following **kpartx** command does not add the *p1* suffix to the partitions ending in a **number**, reboot the system.

## 3.4.3 Configuring Oracle ASM Disks

The configuration of Oracle ASM requires the use of either *udev* rules or Oracle's *ASMLib*. Oracle's *ASMLib* is an optional utility used to manage and assist users with Oracle ASM devices and is not required for proper operation of Oracle ASM disks. Moreover, Oracle *ASMLib* does not have any impact on Oracle Database performance and requires a kernel module labeled **kmod-oracleasm** and proprietary user space utilities that need to be built for specific kernel versions to properly function. *udev* rules represent an alternative to Oracle's *ASMLib* and does not require additional kernel modules or a specific kernel version, thus keeping an overall smaller footprint on the Linux system and allowing more flexibility when upgrading kernel versions. This reference architecture only documents Red Hat's native *udev* rules as at the time of this writing Oracle does not provide an **oracleasm-lib** package for Red Hat Enterprise Linux 7.

### 3.4.3.1 Oracle ASMLib Alternative: Configuring udev Rules

The configuration of Oracle ASM requires the use of either *udev* rules or Oracle's *ASMLib* (currently not available on Red Hat Enterprise Linux 7). This section, focuses on the best practices of using Red Hat's native *udev* rules to setup the appropriate permissions for each



device mapper disk.

1. On the first node of the Oracle RAC cluster as the root user, identify the *Device Mapper Universally Unique Identifier (DM\_UUID)* for each *device mapper* volume. The example below shows the *DM\_UUID* for the partitions of the volumes labeled *ocrvote1p1,ocrvote2p1,ocrvote3p1,db1p1,db2p1,fra1*, and *redo1*.

```
for i in ocrvote1p1 ocrvote2p1 ocrvote3p1 db1p1 db2p1 fra1 redo1; do
printf "%s %s\n" "$i" "$(udevadm info --query=all --name=/dev/mapper/$i |
grep -i dm_uuid)"; done
ocrvote1p1 E: DM_UUID=part1-mpath-3600c0ff000dabfe596a0f65101000000
ocrvote2p1 E: DM_UUID=part1-mpath-3600c0ff000dabfe5a2a0f65101000000
ocrvote3p1 E: DM_UUID=part1-mpath-3600c0ff000dabfe5b4a0f65101000000
db1p1 E: DM_UUID=part1-mpath-3600c0ff000d7e7a899d8515101000000
db2p1 E: DM_UUID=part1-mpath-3600c0ff000dabfe5a7d8515101000000
fra1 E: DM_UUID=part1-mpath-3600c0ff000d7e7a8dbd8515101000000
redo1 E: DM_UUID=part1-mpath-3600c0ff000dabfe5f4d8515101000000
```

2. Create a file labeled *99-oracle-asmdevices.rules* within */etc/udev/rules.d/*
3. Within *99-oracle-asmdevices.rules* file, create rules for each device similar to the example below:

```
KERNEL=="dm- *", ENV{DM_UUID}=="part1-mpath-
3600c0ff000dabfe5f4d8515101000000", OWNER="grid", GROUP="asmadmin", MODE="06
60"
```

To understand the rule above, it can be read as follows:

If any *dm- device (dm-\*)* matches the *DM\_UUID* of *part1-mpath-3600c0ff000dabfe5f4d8515101000000*, assign that *dm- device* to be owned by the *grid* user and part of the *ASMADMIN* group with the permission mode set to *0660*. The *0660* value provides read and write permissions to the user *grid* and owner *asmadmin*.

4. Save the file labeled *99-oracle-asmdevices.rules*
5. Copy the *99-oracle-asmdevices.rules* file to each node within the Oracle RAC Database cluster using the **scp** command and enter the appropriate password credentials for the other nodes. The example below shows how to copy the file to node two of the Oracle RAC Database 12c cluster.

```
scp /etc/udev/rules.d/99-oracle-asmdevices.rules db-oracle-
node2:/etc/udev/rules.d/
root@db-oracle-node2's password:
99-oracle-asmdevices.rules 100% 834 0.8KB/s 00:00
```

6. On each node within the Oracle RAC Database cluster, locate the *dm-* device for each Oracle related partition. An example of how to find the *dm-* device for each partition is to run the following command:

```
for i in db1p1 db2p1 fra1 redo1 ocrvote1p1 ocrvote2p1 ocrvote3p1; do
printf "%s %s\n" "$i" "$(ls -ll /dev/mapper/$i)"; done
db1p1 lrwxrwxrwx. 1 root root 8 May 20 20:39 /dev/mapper/db1p1 -> ../dm-11
db2p1 lrwxrwxrwx. 1 root root 8 May 20 20:39 /dev/mapper/db2p1 -> ../dm-12
```



```
fra1 lrwxrwxrwx. 1 root root 8 May 20 20:39 /dev/mapper/fra1 -> ../dm-13
redo1 lrwxrwxrwx. 1 root root 8 May 20 20:39 /dev/mapper/redo1 -> ../dm-14
ocrvote1p1 lrwxrwxrwx. 1 root root 8 Jan 28 12:11 /dev/mapper/ocrvote1p1
-> ../dm-18
ocrvote2p1 lrwxrwxrwx. 1 root root 8 Jan 28 12:11 /dev/mapper/ocrvote2p1
-> ../dm-19
ocrvote3p1 lrwxrwxrwx. 1 root root 8 Jan 28 12:11 /dev/mapper/ocrvote3p1
-> ../dm-20
```

7. On each node within the Oracle RAC Database cluster, apply and test the rules created within the *99-oracle-asmdevices.rules* by running a **udevadm test** on each device. The example below demonstrates a **udevadm test** on *dm-11*

```
udevadm test /sys/block/dm-11
[... Output Abbreviated ...]
udevadm_test: DM_NAME=db1p1
udevadm_test: DM_UUID=part1-mpath-3600c0ff000d7e7a86485ac5101000000
udevadm_test: DM_SUSPENDED=0
udevadm_test: DEVLINKS=/dev/mapper/db1p1 /dev/disk/by-id/dm-name-db1p1
/dev/disk/by-id/dm-uuid-part1-mpath-3600c0ff000d7e7a86485ac5101000000
/dev/block/253:11
udevadm_test: ID_FS_TYPE=oracleasm
```

8. Confirm each device has the desired permissions. Example of *db1p1* -> *dm-11*, with owner set to *grid* and group set to *asmadmin*.

```
ls -lh /dev/dm-11
brw-rw----. 1 grid asmadmin 253, 11 Jun 6 20:59 /dev/dm-11
```

**NOTE:** If the desired permissions are not visible, please reboot the particular node from the Oracle RAC Database cluster.

**NOTE:** For simplicity, this *99-oracle-asmdevices.rules* file is included in **Appendix J 99-oracle-asmdevices.rules**

### 3.4.4 Optimizing Database Storage using Automatic System Tuning

The **tuned** package in Red Hat Enterprise Linux 7 is recommended for automatically tuning the system for common workloads via the use of profiles. Each profile is tailored for different workload scenarios such as: throughput performance, balanced, & high network throughput.

It is recommended to create a custom profile (*oracle*) with similar settings as the throughput-performance profile but that disables Transparent HugePages (THP) and sets the kernel parameters for Oracle RAC databases workload environments. For more information on why THP is disabled, see **Section 4.1.5 Enabling HugePages**. The table **Table 3.4.4.1: Profile Comparison** provides details between the balanced (default) profile, throughput-performance profile, and the custom profile labeled *oracle*.



| Tuned Parameters                    | balanced     | throughput-performance | oracle                          |
|-------------------------------------|--------------|------------------------|---------------------------------|
| I/O Elevator                        | deadline     | deadline               | deadline                        |
| CPU governor                        | OnDemand     | performance            | performance                     |
| kernel.sched_min_granularity_ns     | auto-scaling | 10ms                   | 10ms                            |
| kernel.sched_wake_up_granularity_ns | 3ms          | 15ms                   | 15ms                            |
| Disk read-ahead                     | 128 KB       | 4096 KB                | 4096 KB                         |
| vm.dirty_ratio                      | 20%          | 40%                    | 80% <sup>18</sup>               |
| File-system barrier                 | on           | on                     | on                              |
| Transparent HugePages               | on           | on                     | Off                             |
| vm.dirty_background_ratio           | 10%          | 10%                    | 3% <sup>18</sup>                |
| vm.swappiness                       | 60%          | 10%                    | 1% <sup>18</sup>                |
| energy_perf_bias                    | Normal       | Performance            | Performance                     |
| min_perf_pct (intel_pstate_only)    | Auto-scaling | Auto-scaling           | Auto-scaling                    |
| tcp_rmem_default                    | Auto-scaling | Auto-scaling           | 262144 <sup>18</sup>            |
| tcp_wmem_default                    | Auto-scaling | Auto-scaling           | 262144 <sup>18</sup>            |
| udp_mem (pages)                     | Auto-scaling | Auto-scaling           | Auto-scaling <sup>18</sup>      |
| vm.dirty_expire_centisecs           | -            | -                      | 500 <sup>18</sup>               |
| vm.dirty_writeback_centisecs        | -            | -                      | 100 <sup>18</sup>               |
| Kernel.shmmax                       | -            | -                      | 4398046511104 <sup>18</sup>     |
| Kernel.shmall                       | -            | -                      | 1073741824 <sup>18</sup>        |
| Kernel.sem                          | -            | -                      | 250 32000 100 128 <sup>18</sup> |
| fs.file-max                         | -            | -                      | 6815744 <sup>18</sup>           |
| fs.aio-max-nr                       | -            | -                      | 1048576 <sup>18</sup>           |
| ip_local_port_range                 | -            | -                      | 9000 65500 <sup>18</sup>        |
| tcp_rmem_max                        | -            | -                      | 4194304 <sup>19</sup>           |
| tcp_wmem_max                        | -            | -                      | 1048576 <sup>18</sup>           |

18 The value explicitly set within the `/etc/sysctl.conf` file has precedence over values set by tuned. Oracle values set in `/etc/sysctl.conf` are done to satisfy prerequisites set by Oracle's installer and will be removed post install and reside only within the oracle tuned profile.



**Table 3.4.4.1: Profile Comparison**

The following procedures provide the steps that are required to install, enable, and select the **oracle** profile.

On each node within the Oracle RAC Database cluster, as the root user,

1. Install the **tuned** package via the yum package manager.

```
yum install tuned
```

2. Enable **tuned** to ensure it is started upon boot time.

```
systemctl enable tuned.service
```

3. Start the **tuned** service

```
systemctl start tuned.service
```

4. Download the accompanying **tar.gz** file for this reference architecture and copy the oracle folder to the `/usr/lib/tuned` directory. For simplicity, **Appendix L Custom Oracle Tuned Profile (script.sh & tuned.conf)** provides the two files `script.sh` and `tuned.conf` that reside within the oracle profile directory.

```
cd /path/to/oracle/profile
cp -r oracle /usr/lib/tuned/
```

5. Activate the newly added oracle profile:

```
tuned-adm profile oracle
```

6. Verify that THP is now disabled via the following command:

```
cat /sys/kernel/mm/transparent_hugepage/enabled
always madvise [never]
```

7. Disable *transparent* huge pages persistently across reboots by adding to the kernel boot command line within the `/etc/default/grub` and add within the **GRUB\_CMDLINE\_LINUX** the following:

```
cat /etc/default/grub
GRUB_TIMEOUT=5
GRUB_DISTRIBUTOR="$(sed 's, release .*$,,g' /etc/system-release)"
GRUB_DEFAULT=saved
GRUB_DISABLE_SUBMENU=true
GRUB_TERMINAL_OUTPUT="console"
GRUB_CMDLINE_LINUX="rd.lvm.lv=myvg/swap rd.lvm.lv=myvg/usr
vconsole.font=latarcyrheb-sun16 rd.lvm.lv=myvg/root crashkernel=auto
vconsole.keymap=us rhgb quiet transparent_hugepage=never"
GRUB_DISABLE_RECOVERY="true"
```



8. For the grub changes to take effect, run the following command:

```
grub2-mkconfig -o /boot/grub2/grub.cfg
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-3.10.0-123.el7.x86_64
Found initrd image: /boot/initramfs-3.10.0-123.el7.x86_64.img
Found linux image: /boot/vmlinuz-0-rescue-
41c535c189b842eea5a8c20cbd9bfff26
Found initrd image: /boot/initramfs-0-rescue-
41c535c189b842eea5a8c20cbd9bfff26.img
done
```

For simplicity, the *oracle* profile consisting of two files, *script.sh* & *tuned.conf* is included in **Appendix M Configuration Files**

**NOTE:** If at any point in time a revert to the original settings are required with persistence across reboots, the following commands can be run on each node within the Oracle RAC Database cluster:

```
systemctl stop tuned.service
systemctl disable tuned.service
```

**NOTE:** Even if reverting to the original settings, it is recommended to keep *transparent* huge pages disabled within the */etc/default/grub* file.



# 4 Oracle 12c R1 Configuration

## 4.1.1 Installing Oracle Grid Infrastructure (Required for ASM)

The installation of the Oracle Grid Infrastructure for Oracle RAC Database 12c Release 1 (12.1) is required for the use of Oracle ASM. Prior to the installation of the Oracle Grid Infrastructure, ensure that the prerequisites from the following sections have been met:

- **Section 2 Reference Architecture Environment**
- **Section 3 Reference Architecture Configuration Details**

**NOTE:** The reference environment uses the `/u01/app/grid` as the grid base. The owner is set to `grid` and the group is set to `oinstall`. Run the following commands to create the `grid` base directory and set the appropriate permissions:

On each node within the Oracle RAC environment as the `root` user, create the following directory structure and set the proper permissions.

```
mkdir --parents /u01/app/grid
chown --recursive grid.oinstall /u01/
```

The following steps are intended **only** for node one of the Oracle RAC Database environment unless otherwise specified.

1. Download the two Oracle Grid Infrastructure software files<sup>19</sup> from the Oracle Software Delivery Cloud site<sup>20</sup>.
2. As the `grid` user, create a temporary directory within `/u01/app/grid/grid-software` to store the Oracle Grid Software **zip** file, move the Oracle Grid Software **zip** file to the `/u01/app/grid/grid-software` location, ensure the Oracle Grid Software **zip** has the proper permissions and unpack its contents.

As the `grid` user,

```
mkdir /u01/app/grid/grid-software
mv V46096-01_1of2.zip /u01/app/grid/grid-software
mv V46096-01_2of2.zip /u01/app/grid/grid-software
```

As the `root` user,

```
chown grid.oinstall V46096-01_1of2.zip
chown grid.oinstall V46096-01_2of2.zip
```

As the `grid` user,

```
cd /u01/app/grid/grid-software
unzip V46096-01_1of2.zip
unzip V46096-01_2of2.zip
```

<sup>19</sup> Oracle Database 12c Release 1 Grid Infrastructure (12.1.0.2.0) - [V46096-01\\_1of2.zip](http://edelivery.oracle.com) and [V46096-01\\_2of2.zip](http://edelivery.oracle.com) from <http://edelivery.oracle.com>



3. As the grid user, start the OUI via the command:

```
/u01/app/grid/grid-software/grid/runInstaller
```

**NOTE:** Ensure to **SSH** with the **-Y** option as the grid user from the client server. Otherwise, the following error will occur.

```
/u01/app/grid/grid-software/grid/runInstaller
Starting Oracle Universal Installer...

Checking Temp space: must be greater than 415 MB. Actual 3689 MB
Passed
Checking swap space: must be greater than 150 MB. Actual 16415 MB
Passed
Checking monitor: must be configured to display at least 256 colors
>>> Could not execute auto check for display colors using command
/usr/bin/xdpyinfo. Check if the DISPLAY variable is set. Failed <<<<

Some requirement checks failed. You must fulfill these requirements
before continuing with the installation,

Continue? (y/n) [n] n
```

1. Within the **Installation Option** window, select **Install and Configure Oracle Grid Infrastructure for a Cluster** and click **Next**.
2. Within the **Cluster Type** window, select between **Configure a Standard cluster** or **Configure a Flex cluster** and click **Next**. This reference architecture focuses on using a standard cluster.
3. Within the **Installation Type** window, select **Advanced Installation** and click **Next**.
4. Within the **Product Languages** window, select the appropriate language, and click **Next**.
5. Within the **Grid Plug and Play Information** window, provide the appropriate credentials for the SCAN and click **Next**. This reference architecture provides the following SCAN credentials:
  - Cluster Name: *db-ora-cluster*
  - SCAN Name: *db-oracle-scan.cloud.lab.eng.bos.redhat.com*
  - SCAN Port: 1521
  - Uncheck **Configure GNS**

**NOTE:** The SCAN Name is the name registered with the DNS Server as seen in **Section 3.2.3 Configure SCAN via DNS**





Oracle Grid Infrastructure 12c Release 1 Installer – Step 5 of 18

### Grid Plug and Play Information

Single Client Access Name (SCAN) allows clients to use one name in connection strings to connect to the cluster as a whole. Client connect requests to the SCAN name can be handled by any cluster node.

Cluster Name:

SCAN Name:

SCAN Port:

☐ **Configure GNS**

☒ Configure nodes Virtual IPs as assigned by the Dynamic Networks

☒ Create a new GNS

GNS VIP Address:

GNS Sub Domain:

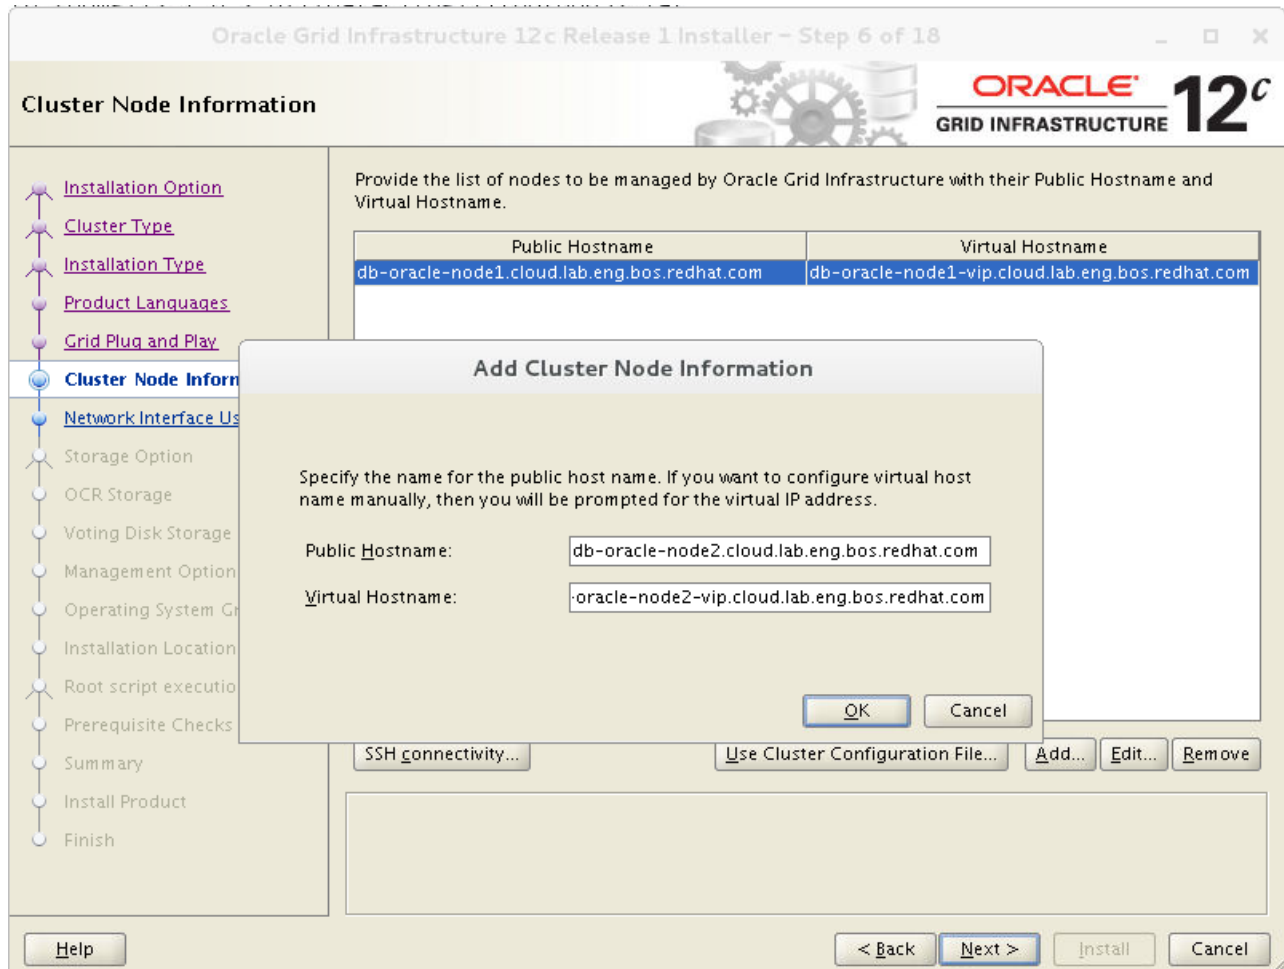
☐ Use Shared GNS

GNS Client Data:

Figure 4.1.1.1: Oracle Plug and Play Window



6. Within the same **Cluster Node Information** window, click the **Add** button to add each node within the Oracle RAC Database cluster and click OK. Each node within the Oracle RAC cluster requires the public hostname and VIP information as seen in **Figure 4.1.1.2: Cluster Node Information Window, Add button**



**Figure 4.1.1.2: Cluster Node Information Window, Add button**

7. Within the same **Cluster Node Information** window, select the **SSH Connectivity** button to set the passwordless SSH connectivity by entering the **OS Password:** credentials for the grid user and clicking **Setup**. Once a dialog box returns with the 'Successfully established passwordless SSH connectivity between the selected nodes', click **OK** and click **Next** to continue to the next window.
8. Within the **Network Interface Usage** window, select the Interface Name, *bond0*, to be set as the Interface Type *Public* and the Interface Name, *ens10f2* and *ens10f3*, to be set as the Interface Type *Private*. Any other Interfaces should be set to *Do Not Use*. Select **Next** to continue.
9. Within the **Storage Option** window, this reference architecture selects **Use Standard ASM for storage** and click **Next**.
10. Within the **Create ASM Disk Group** window, provide the following:



- a Disk Group Name, i.e. *OCRVOTE*
- Redundancy Level
  - *EXTERNAL* – redundancy provided by the storage system RAID, and not by Oracle ASM.
  - *NORMAL* – provides two-way mirroring by Oracle ASM, thus provided two copies of every data extent.
  - *HIGH* – provides three-way mirroring by Oracle ASM thus enduring the loss of two ASM disks within different failure groups.

The following **Table 4.1.1 Table 4.1.1.1: ASM Disk Group Window** provides the details for the reference environment.

| ASM Disk Group Window                                                    |                                                                                                 |                                    |
|--------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|------------------------------------|
| Disk Group Name                                                          | <i>OCRVOTE</i>                                                                                  |                                    |
| Redundancy Level                                                         | Normal                                                                                          |                                    |
| Device Mapper multipath and udev rules, Disks Assigned to ASM Disk Group | <i>/dev/mapper/ocrvote1p1</i><br><i>/dev/mapper/ocrvote2p1</i><br><i>/dev/mapper/ocrvote3p1</i> | Header Status:<br><i>Candidate</i> |

**Table 4.1.1.1: ASM Disk Group Window**



- Allocation Unit (AU) Size set to 4 MB<sup>20</sup>
  - A 4 MB AU Size is used to decrease the amount of extents Oracle needs to manage. With less extents to manage, CPU utilization and memory consumption is reduced thus improving performance. The AU Size varies depending on the type of Oracle workload, I/O size per transaction, and overall diskgroup size. There is no “best size” for AU size, but a good starting point is a 4 MB. Please ensure to visit Oracle's documentation<sup>21</sup> for more information.

To display the appropriate candidate disks, click on the **Change Discovery Path** button and enter as the **Disk Discovery Path** one of the following as appropriate:

- For Device Mapper devices, type: `/dev/mapper/*`




Figure 4.1.1.3: Create ASM Disk Group Window

<sup>20</sup> Oracle ASM Extents - [http://docs.oracle.com/cd/E16655\\_01/server.121/e17612/asmcon.htm#OSTMG94063](http://docs.oracle.com/cd/E16655_01/server.121/e17612/asmcon.htm#OSTMG94063)



**NOTE:** There is a significant size increase for the OCR and Voting Disk Diskgroup. Starting with Oracle Database 12.1.0.2.0, Oracle has included the installation of the Grid Infrastructure Management Repository<sup>21</sup>. For a two node Oracle RAC as shown within this reference architecture, a minimum of 11059 MB of space is required when using Normal Redundancy. If using External Redundancy, only 5528 MB required.

11. Click **Next** once complete within the **Create ASM Disk Group** window
12. Within the **ASM Password** window, specify the password for the `SYS` and `ASMSNMP` user accounts.
13. Within the **Failure Isolation Support** window, select whether to use or not use the Intelligent Platform Management Interface (IPMI). This reference architecture selects the **Do not use Intelligent Platform Management Interface (IPMI)** radio button and clicks **Next**.
14. Within the **Management Options** window, **Register with Enterprise Manager (EM) Cloud Control** is unchecked, click **Next**
15. Within the **Operating System Groups** window, select the appropriate OS groups. The values as created and assigned within this reference environment are as follows:
  - Oracle ASM Administrator Group – `ASMADMIN`
  - Oracle ASM DBA Group – `ASMDBA`
  - Oracle ASM Operator Group – `ASMOPER`
11. Within the **Installation Location** window, specify the appropriate Oracle base and software locations. The values set by this reference environment are as follows:
  - `ORACLE BASE` - `/u01/app/grid`
  - `SOFTWARE LOCATION` - `/u01/app/grid/product/12.1.0/grid`
12. Within the **Create Inventory** window, specify the inventory directory. The values set by this reference environment are as follows:
  - Inventory Directory - `/u01/app/oraInventory`
13. Within the **Root script execution configuration** window, select the check box labeled **Automatically run configuration scripts** and enter the root user credentials. The step specifying the root user credentials in order to run specific configuration scripts automatically at the end of the installation is optional. For the purposes of this reference environment, the root credentials are given in order to speed up the Oracle Grid Infrastructure installation process.
14. Within the **Prerequisite Checks** window, review the status and ensure there are no errors prior to continuing the installation. For failures with a status set to Fixable, select the **Fix & Check Again** button. The execution of the **Fix & Check Again** button provides a `runfixup.sh` script provided by the OUI. If selected **Automatically run configuration scripts** from the previous step, the Oracle OUI uses the root credentials and runs the

---

21 Oracle Documentation New Features -  
<https://docs.oracle.com/database/121/CWLIN/whatsnew.htm#CHDEJFCG>



*fixup.sh* script automatically. Otherwise, as root, on each node within the Oracle RAC cluster, run the */tmp/CVU\_12.1.0.2.0\_grid/runfixup.sh* and click on the **Check Again** button once the *runfixup.sh* has finished. For errors that can be ignored, select the Ignore All check box and click **Next**.

The following check errors are common and extra details are provided below.

- Task resolv.conf Integrity – This task checks consistency of the file */etc/resolv.conf* across all nodes. Manually verify that all nodes across the Oracle RAC Database cluster provide the appropriate response with the use of **nslookup**. If the appropriate response is achieved, this error can be safely ignored.
- Network Time Protocol (NTP) – This task verifies cluster time synchronization on clusters. Manually verify that **ntp** is running on all nodes within the Oracle RAC cluster. Please review **Section 3.3.2 NTP Configuration** for more details on properly setting up NTP. If NTP is properly running and configured, this error can be safely ignored.
- Device Checks for ASM – This task verifies that specified devices meet the requirements for Oracle ASM. Once the prerequisite of **cvuqdisk** rpm is installed, this warning will disappear.
- */dev/shm* mounted as temporary file system – Mounting of */dev/shm* is required when using Automatic Memory Management (AMM). Due to this reference environment, taking advantage of HugePages, */dev/shm* is not used. This warning can be safely ignored.

15. Within the **Summary** window, review all the information provided, and select **Install** to start the installation.

16. During the installation process, within the **Oracle Grid Infrastructure** pop up window, select **yes** to allow the installer to run as the root user to execute the configuration scripts.

**NOTE:** Oracle Cluster Verification Utility will provide a **Failed** status due to NTP task error. Click **Next** within the **Install Product** window, and select **Yes** within the **Oracle Grid Infrastructure 12c Release 1 Installer** popup window.

17. Within the **Finish** window, verify the installation was successful and click **Close**.

18. Verify the installation was successful using the **crsctl** command as follows:

As the grid user:

```
$GRID_HOME/bin/crsctl check crs
CRS-4638: Oracle High Availability Services is online
CRS-4537: Cluster Ready Services is online
CRS-4529: Cluster Synchronization Services is online
CRS-4533: Event Manager is online
```

**NOTE:** The Grid Home used in this reference environment is: */u01/app/12.1.0/grid*





## 4.1.2 Installing Oracle 12c R1 Database Software

Prior to the installation of the Oracle RAC Database 12c Release 1 (12.1), ensure the following prerequisites from the following sections have been met:

- **Section 2 Reference Architecture Environment**
- **Section 3 Reference Architecture Configuration Details**

**NOTE:** The reference environment uses the `/u01/app/oracle` as the Oracle base. The owner is set to `oracle` and the group is set to `OINSTALL`. Run the following commands to create the `oracle` base directory set the appropriate permissions:

On each node within the Oracle RAC environment, as the `root` user, create the following directory structure and set the proper permissions.

```
mkdir /u01/app/oracle
chown --recursive oracle.oinstall /u01/app/oracle
```

The following steps are intended for **only** node one of the Oracle RAC Database 12c environment unless otherwise specified.

1. Download the two Oracle Database software files<sup>22</sup> from the Oracle Software Delivery Cloud site.
2. As the `oracle` user, create a temporary directory within `/u01/app/oracle/oracle-software` to store the Oracle Database Software **zip** files, move the Oracle Database software **zip** files to the `/u01/app/oracle/oracle-software` location, ensure the Oracle Database software **zip** has the proper permissions and unpack its contents.

As the `oracle` user,

```
mkdir /u01/app/oracle/oracle-software
mv V46095-01_1of2.zip V46095-01_2of2.zip /u01/app/oracle/oracle-software/
```

As the `root` user,

```
chown oracle.oinstall /u01/app/oracle/oracle-software/V46095-01_1of2.zip
chown oracle.oinstall /u01/app/oracle/oracle-software/V46095-01_2of2.zip
```

As the `oracle` user,

```
cd /u01/app/oracle/oracle-software
unzip V46095-01_1of2.zip
unzip V46095-01_2of2.zip
```

3. As the `oracle` user, start the OUI via the command:

```
/u01/app/oracle/oracle-software/database/runInstaller
```

<sup>22</sup> Oracle Database 12c Release 1 (12.1.0.2.0) - [V46095-01\\_1of2.zip](http://edelivery.oracle.com) and [V46095-01\\_2of2.zip](http://edelivery.oracle.com) from <http://edelivery.oracle.com>



**NOTE:** Ensure to issue **SSH** with the **-Y** option as the **oracle** user from the client server. Otherwise, the following error will occur.

```
/u01/app/oracle/oracle-software/database/runInstaller
Starting Oracle Universal Installer...

Checking Temp space: must be greater than 500 MB. Actual 3461 MB
Passed
Checking swap space: must be greater than 150 MB. Actual 20479 MB
Passed
X11 connection rejected because of wrong authentication.
X11 connection rejected because of wrong authentication.
Checking monitor: must be configured to display at least 256 colors
>>> Could not execute auto check for display colors using command /
usr/bin/xdpyinfo. Check if the DISPLAY variable is set. Failed <<<<

Some requirement checks failed. You must fulfill these requirements
before continuing with the installation,

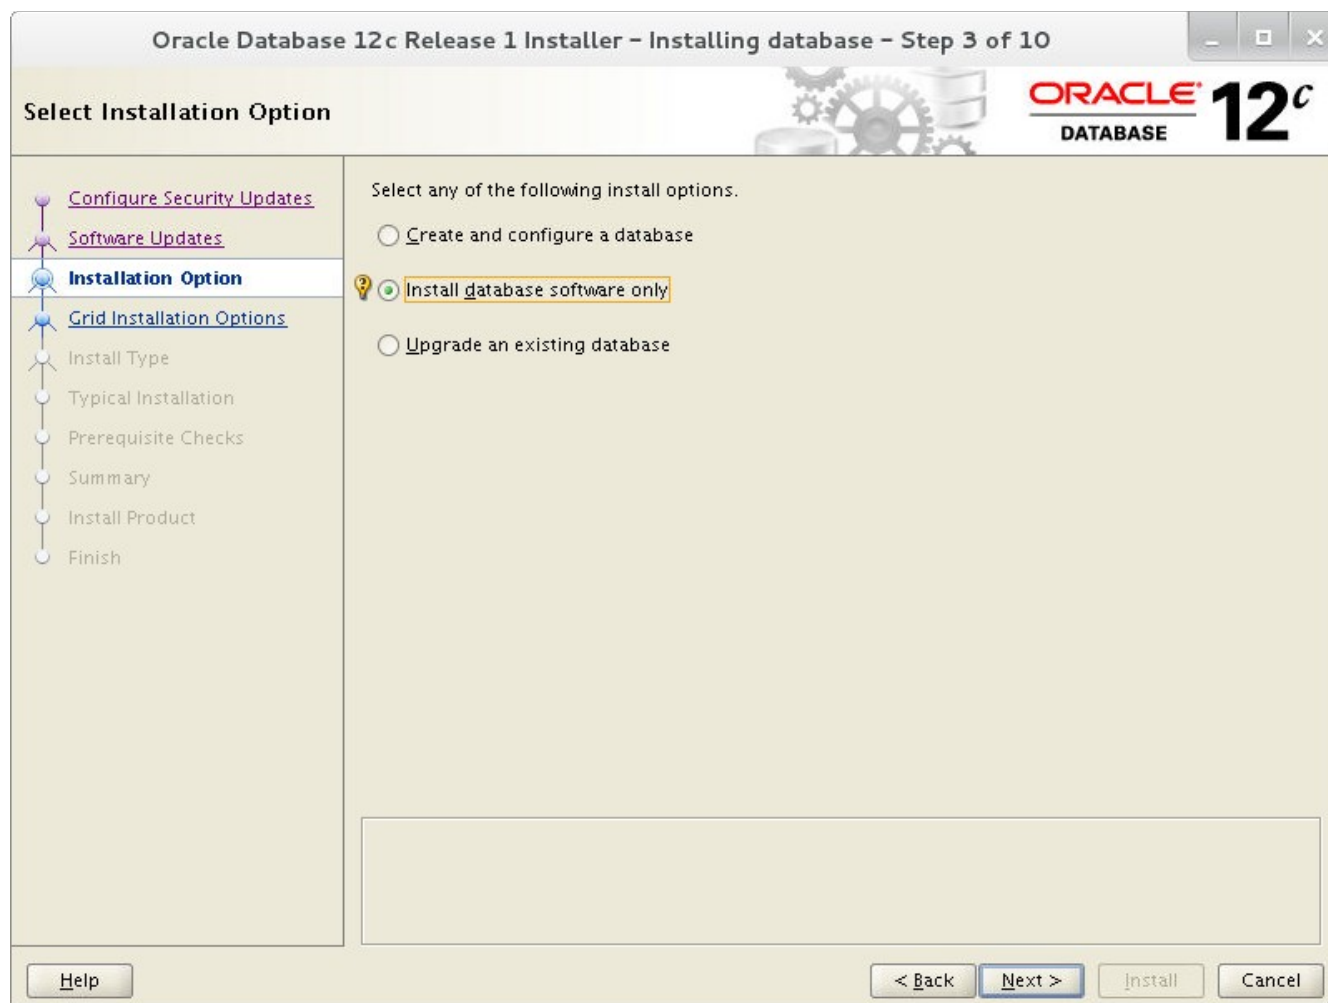
Continue? (y/n) [n] n
```

4. Within the **Configure Security Updates** window, provide the My Oracle Support email address for the latest security issues information. Otherwise uncheck the **I wish to receive security updates via My Oracle Support** and click **Next**.






5. Within the **Installation Option** window, select **Install database software only** and click **Next**.



**Figure 4.1.2.1: Installation Option Window**




6. Within the **Grid Installation Options** window, select **Oracle Real Application Clusters database installation** as the type of database installation being performed.



**Figure 4.1.2.2: Grid Installation Options Window**



7. Within the **Select List of Nodes** window, ensure all nodes for the Oracle RAC database cluster are checked and click on the **SSH Connectivity** button. Within the **OS Password:** dialog box enter the user `oracle`'s password and click **Setup**. Once a dialog box returns with 'Successfully established passwordless SSH connectivity between the selected nodes', click **OK** and **Next** to continue.



**Figure 4.1.2.3: Select List of Nodes Window**

8. Within the **Product Languages** window, select the appropriate language for the installation.
9. Within the **Database Edition** window, select the **Enterprise Edition** and click **Next**.
10. Within the **Installation Location** window, select the appropriate Oracle base and software location and click **Next**. For the purposes of this reference environment, the following values are set:
  - `ORACLE BASE` - `/u01/app/oracle`
  - `SOFTWARE LOCATION` - `/u01/app/oracle/product/12.1.0/dbhome_1`



11. Within the **Operating System Groups** window, select the appropriate OS groups and click **Next**. For the purposes of this reference environment, the following values are set as:

- Database Administrator Group – *DBA*
- Database Operator Group – *OPER*
- Database Backup and Recovery Group – *BACKUPDBA*
- Data Guard Administrative group – *DGDBA*
- Encryption Key Management Administrative group – *KMDBA*

12. Within the **Prerequisite Checks** window, review the status and ensure there are no errors prior to continuing the installation. For failures with a status set to Fixable, select the **Fix & Check Again** button. The execution of the **Fix & Check Again** button provides a *runfixup.sh* script provided by the OUI. As root, run the *runfixup.sh* and click on the **Check Again** button once the *runfixup.sh* has finished. For errors that can be ignored, select the Ignore All check box and click **Next**.

- The following check errors are common and extra details are provided below.
  - Task resolv.conf Integrity – This task checks consistency of the file */etc/resolv.conf* across all nodes. Manually verify that all nodes across the Oracle RAC Database cluster provide the appropriate response with the use of **nslookup**. If the appropriate response is achieved, this error can be safely ignored.
  - Clock Synchronization – This task checks to see if NTP daemon or service is running. If the service is not running, it provides this error and displays a PRVF-5415 error. However, the OUI contains a bug that does not see the NTP service is actually running and provides a not running status. Manually verify that all nodes across the Oracle RAC Database cluster are running the **ntpd** service. Please review **Section 3.3.2 NTP Configuration** for more details on properly setting up NTP. If NTP is properly running and configured, this error can be safely ignored.
  - Maximum locked memory check – This task checks if memlock is set within the */etc/security/limits.conf* file and is only a warning. Setting *memlock* allows the oracle user to lock a certain amount of memory from physical RAM that isn't swapped out. The value is expressed in kilobytes and is important from the Oracle perspective because it provides the oracle user permission to use huge pages. This warning can be safely ignored at the moment of installation as it is configured later during the setup of huge pages. More information on huge pages can be found in **Section 4.1.5 Enabling HugePages**
  - */dev/shm* mounted as temporary file system – Mounting of */dev/shm* is required when using Automatic Memory Management (AMM). Due to this reference environment, taking advantage of HugePages, */dev/shm* is not used. This warning can be safely ignored



13. Within the **Prerequisite Checks** window, select **yes** to the Oracle Database 12c Release 1 Installer popup window.
14. Within the **Summary** window, review all the information provided, and select **Install** to start the installation.
15. Once the installation completes, execute the scripts within the **Execute Configuration scripts** window. As the root user on each node, run the following:

```
/u01/app/oracle/product/12.1.0/dbhome_1/root.sh
```

```
Performing root user operation.
```

```
The following environment variables are set as:
```

```
ORACLE_OWNER= oracle
```

```
ORACLE_HOME= /u01/app/oracle/product/12.1.0/dbhome_1
```

```
Enter the full pathname of the local bin directory: [/usr/local/bin]:
```

```
The contents of "dbhome" have not changed. No need to overwrite.
```

```
The contents of "oraenv" have not changed. No need to overwrite.
```

```
The contents of "coraenv" have not changed. No need to overwrite.
```

```
Entries will be added to the /etc/oratab file as needed by
Database Configuration Assistant when a database is created
Finished running generic part of root script.
```

```
Now product-specific root actions will be performed.
```

**NOTE:** In the example above, `/u01/app/oracle/product/12.1.0/dbhome_1` is the Oracle home directory.

16. Click **OK** within the **Execute Configuration scripts** window.
17. Within the **Finish** window, verify the installation was successful and click **Close**.



### 4.1.3 Creating ASM Diskgroups via the ASM Configuration Assistant (ASMCA)

Prior to the creation of an Oracle RAC database, create the Fast Recovery Area (FRA), Redo Logs, and Oracle ASM diskgroups via Oracle's ASM Configuration Assistant (ASMCA).

The following steps should be done on node one of the Oracle RAC Database cluster environment.

1. **ssh** with the **-Y** option as the **grid** user is required prior to running **asmca**.
2. As the **grid** user, start **asmca** via the following command:

```
/u01/app/12.1.0/grid/bin/asmca
```

**NOTE:** In the example above, **/u01/app/12.1.0/grid** is the **grid** home directory.

3. Via the **asmca** application, select the **Disk Groups** tab and click **Create**.

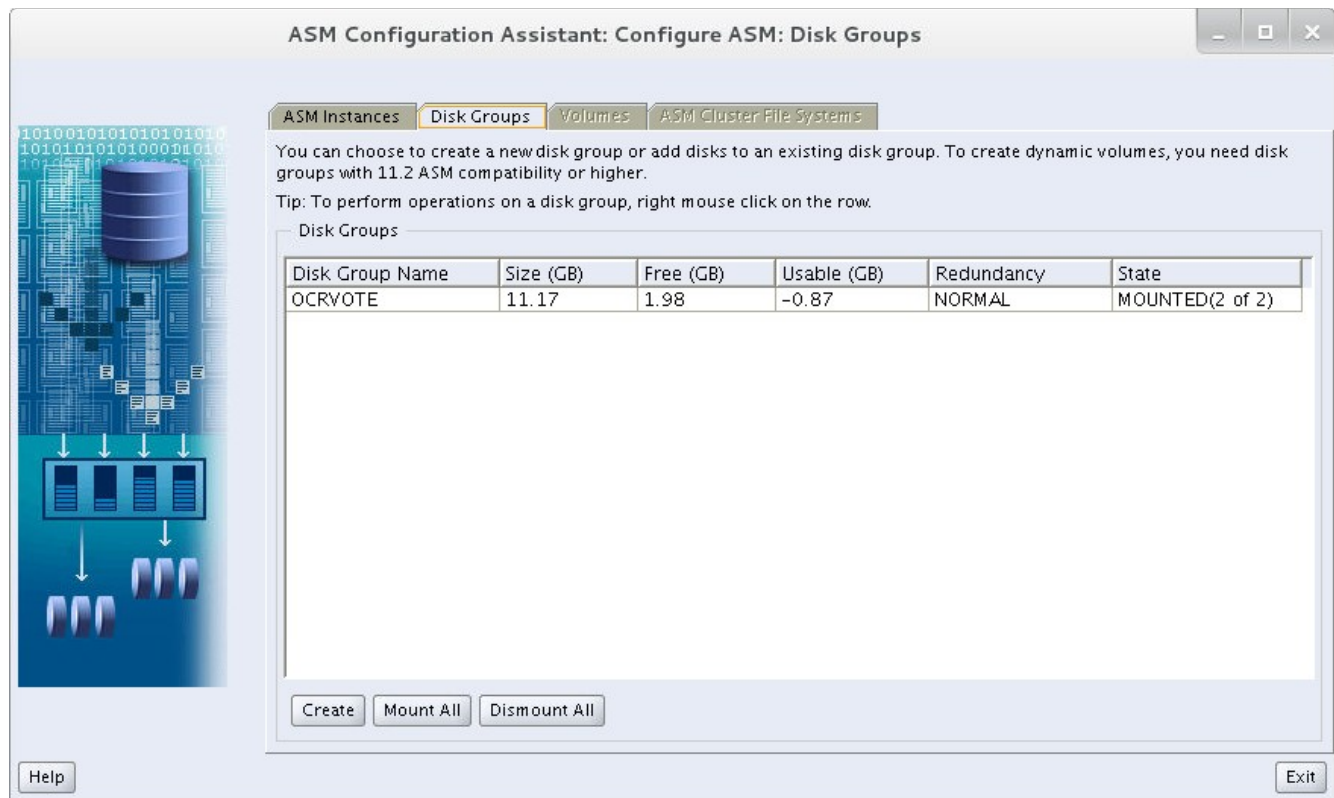


Figure 4.1.3.1: ASMCA Disk Groups Tab



4. Within the **Create Disk Group** window, provide the following:
- A name for the disk group, i.e. *FRADG*
  - Redundancy level for the disk group, i.e. *External Redundancy*
  - Selection of the disks to be added to the disk group, i.e. */dev/mapper/fra1*
  - Click the **Show Advanced Options** radio button and provide the appropriate AU Size. This reference environment uses an AU Size of 4 MB for disk groups all disks groups: *DATA*, *FRADG*, and *REDODG*.

| ASMCA Disk Group |                                                                                           |                                                |
|------------------|-------------------------------------------------------------------------------------------|------------------------------------------------|
| Diskgroup Name   | Disk Path                                                                                 | Header Status                                  |
| DATADG           | If using device mapper multipath:<br><i>/dev/mapper/db1p1</i><br><i>/dev/mapper/db2p1</i> | If using device mapper multipath:<br>Candidate |
| FRADG            | If using device mapper multipath:<br><i>/dev/mapper/fra1</i>                              | If using device mapper multipath:<br>Candidate |
| REDODG           | If using device mapper multipath:<br><i>/dev/mapper/redo1</i>                             | If using device mapper multipath:<br>Candidate |

**Table 4.1.3.1: ASMCA Create Disk Group**





### Create Disk Group

Disk Group Name

Redundancy

Redundancy is achieved by storing multiple copies of the data on different failure groups. Normal redundancy needs disks from at least two different failure groups, and high redundancy from at least three different failure groups.

☐ High ☐ Normal ☒ External (None)

Select Member Disks

☒ Show Eligible ☐ Show All

Quorum failure groups are used to store voting files in extended clusters and do not contain any user data. They require ASM compatibility of 11.2 or higher.

| <input type="checkbox"/>            | Disk Path         | Header Status | Disk Name | Size (MB) | Quorum                   |
|-------------------------------------|-------------------|---------------|-----------|-----------|--------------------------|
| <input checked="" type="checkbox"/> | /dev/mapper/db1p1 | CANDIDATE     |           | 95365     | <input type="checkbox"/> |
| <input checked="" type="checkbox"/> | /dev/mapper/db2p1 | CANDIDATE     |           | 95365     | <input type="checkbox"/> |
| <input type="checkbox"/>            | /dev/mapper/fra1  | CANDIDATE     |           | 190733    | <input type="checkbox"/> |
| <input type="checkbox"/>            | /dev/mapper/redo1 | CANDIDATE     |           | 9535      | <input type="checkbox"/> |

Note: If you do not see the disks which you believe are available, check the Disk Discovery Path and read/write permissions on the disks. The Disk Discovery Path limits set of disks considered for discovery.

Disk Discovery Path: /dev/mapper/\* Change Disk Discovery Path

Disk Group Attributes

An allocation unit (AU) is the fundamental unit in which contiguous disk space is allocated to ASM files. ASM file extent size is a multiple of AUs. The AU size cannot be modified later.

Allocation Unit Size (MB)

Specify minimum software versions for ASM, Database and ASM volumes that this disk group need to be compatible with.

ASM Compatibility

Database Compatibility

ADVM Compatibility

Refer Oracle Automatic Storage Management Administrator's Guide for more details on the Compatibility matrix.

Hide Advanced Options OK Cancel Help

**Figure 4.1.3.2: ASM Disk Group Creation Window**

**NOTE:** To display the appropriate eligible disks, click on the **Change Discovery Path** button and enter as the 'Disk Discovery Path' one of the following as appropriate:

- For Device Mapper devices, type: **/dev/mapper/\***

**NOTE:** A 4 MB AU Size is used to decrease the amount of extents Oracle needs to manage. With less extents to manage, CPU utilization and memory consumption is reduced thus improving performance. The AU Size varies depending on the type of Oracle workload, I/O size per transaction, and overall diskgroup size. There is no “best size” for AU size, but a good starting point is a 4 MB. Please ensure to visit Oracle's documentation<sup>21</sup> for more information.





Click the **OK** button once the steps above are complete.

5. Repeat steps 2, 3, and 4 to configure a disk group for the redo logs and Fast Recovery Area (FRA).

**NOTE:** Separation of redo logs into a separate Oracle ASM disk group is optional, but recommended.

6. Once all disk groups are created, click the **Exit** button from the main **ASM Configuration Assistant** window. Click **yes**, when asked to confirm quitting the application.

## 4.1.4 Creating Pluggable Databases using Database Configuration Assistant (DBCA)

With the introduction to Oracle Database 12c, Oracle introduced the Multitenant architecture. The Multitenant architecture provides the ability to consolidate multiple databases known as pluggable databases (PDBs) into a single container database (CDB). It provides advantages<sup>23</sup> which include easier management and monitoring of the physical database, fewer patches and upgrades, performance metrics consolidated into one CDB, and sizing one SGA instead of multiple SGAs. While using the Multitenant architecture is optional, this reference architecture focuses on describing the step-by-step procedure of taking advantage of it. When creating an Oracle RAC database, the recommended method is the usage of the **dbca** utility.

Prior to getting into the details of installing a container database (CDB) and deploying pluggable databases (PDB), an overview of the key concepts of the Multitenant Architecture is provided.

Container<sup>23</sup> – is a collection of schemas, objects, and related structures in a multitenant container database (CDB) that appears logically to an application as a separate database. Within a CDB, each container has a unique ID and name.

A CDB consists of two types of containers: the root container and all the pluggable databases that attach to a CDB.

Root container<sup>23</sup> – also called *the root*, is a collection of schemas, schema objects, and nonschema objects to which all PDBs belong. Every CDB has one and only one root container, which stores the system metadata required to manage PDBs (no user data is stored in the root container). All PDBs belong to the root. The name of the root container is `CDB$ROOT`.

PDB<sup>23</sup> – is a user-created set of schemas, objects, and related structures that appears logically to an application as a separate database. Every PDB is owned by `sys`, which is a common user in the CDB, regardless of which user created the CDB.

For more information on Oracle's Multitenant architecture, visit Oracle's documentation<sup>23</sup>.

---

<sup>23</sup> Purpose of a Multitenant Environment -

[http://docs.oracle.com/cd/E16655\\_01/server.121/e17636/cdb\\_intro.htm#ADMIN13509](http://docs.oracle.com/cd/E16655_01/server.121/e17636/cdb_intro.htm#ADMIN13509)



The following section describes the step-by-step procedure to create a container database (CDB) that holds two pluggable databases (PDB) thus taking advantage of Oracle's Multitenant architecture.

The following steps should be done on node one of the Oracle RAC Database cluster environment.

1. **ssh** with the -Y option as the `oracle` user prior to running **dbca**.
2. As the `oracle` user, run the **dbca** utility via the command:

```
/u01/app/oracle/product/12.1.0/dbhome_1/bin/dbca &
```

**NOTE:** In the example above, `/u01/app/oracle/product/12.1.0/dbhome_1` is the Oracle home directory.

3. Within the **Database Operations** window, select **Create a Database** radio button and click **Next**.
4. Within the **Creation Mode** window, select **Advanced Mode** radio button and click **Next**.
5. Within the **Database Template** window, select ensure **Database Type** is set to **Oracle Real Application Clusters (RAC) database** and **Custom Database** radio button and click **Next**.

**NOTE:** This reference environment sets **Configure Type** to **Admin-Managed**. If interested in Policy-Managed Configuration Type, more information can be found within Oracle's documentation – Using Server Pools with Oracle RAC Databases<sup>24</sup>.

---

24 Oracle Documentation – Using Server Pools with Oracle RAC -  
<https://docs.oracle.com/database/121/RILIN/srvpool.htm#RILIN1063>



6. Within the **Database Identification** window, set a global database name and SID prefix, i.e. *cdb*. Check the check box that reads **Create as Container Database**. Select the number of PDBs to install and provide a PDB Name Prefix, i.e. *pdb* and click **Next**. This reference environment creates two PDBs.

The screenshot shows the 'Database Identification' window in the Oracle Database Configuration Assistant. The window title is 'Database Configuration Assistant - Create Database - Step 4 of 15'. The Oracle logo and '12c' are visible in the top right. On the left is a navigation pane with steps: Database Operation, Creation Mode, Database Template, Database Identification (selected), Database Placement, Management Options, Database Credentials, Storage Locations, Database Options, Initialization Parameters, Creation Options, Prerequisite Checks, Summary, Progress Page, and Finish. The main area contains instructions: 'Provide the identifier information required to access the database uniquely. An Oracle database is uniquely identified by a Global Database Name, typically of the form "name.domain". A database is referenced by an Oracle instance on each cluster database node. Specify a prefix to be used to name the cluster database instances.' Below this are input fields: 'Global Database Name:' with 'cdb', and 'SID Prefix:' with 'cdb'. There are three radio button options: 'Create As Container Database' (checked), 'Create an Empty Container Database', and 'Create a Container Database with one or more PDBs'. Below the third option is a 'Number of PDBs:' spinner set to '2' and a 'PDB Name Prefix:' field with 'pdb'. At the bottom are buttons: 'Help', '< Back', 'Next >', 'Finish', and 'Cancel'.

**Figure 4.1.4.1: Database Identification Window**

7. Within the **Database Placement** window, move all available nodes within the Oracle RAC Database cluster to the **Selected** side via the arrow button and click **Next**.
8. Within the **Management Options** window, ensure the check box **Run Cluster Verification Utility (CV) Checks Periodically** and **Configure Enterprise Manager (EM) Database Express** is checked (default) and click **Next**.
9. Within the **Database Credentials** window, provide the administrative passwords for *SYS*, *SYSTEM*, and *PDBADMIN*, click **Next**.



10. Within the **Storage Locations** window, select the appropriate **Storage Type** and storage locations. For the purposes of this reference environment, the following selections were made:

- Database Files Storage Type – Automatic Storage Management (ASM)
  - Use Common Location for All Database Files selected
    - Database File Locations: +DATADG
    - Use Oracle-Managed Files (checked)
- Recovery Related Files Storage Type – Automatic Storage Management (ASM)
  - Storage Locations – Specify Fast Recovery Area
    - Fast Recovery Area: +FRADG
    - Fast Recovery Area Size: 190635 Megabytes<sup>25</sup> (max size created of volume)
    - Enable Archiving (checked)

Database Configuration Assistant – Create Database – Step 8 of 15

**Storage Locations**

Database files Storage Type: Automatic Storage Management (ASM) ▼

☐ Use Database File Locations from Template

☒ Use Common Location for All Database Files

File Location: +DATADG Browse...

☒ Use Oracle-Managed Files Multiplex Redo Logs and Control Files...

Choose the recovery options for the database.

Recovery files Storage Type: Automatic Storage Management (ASM) ▼

☒ Specify Fast Recovery Area

Fast Recovery Area: +FRADG Browse...

Fast Recovery Area Size: 190635 MB

☒ Enable Archiving Edit Archive Mode Parameters

File Location Variables...

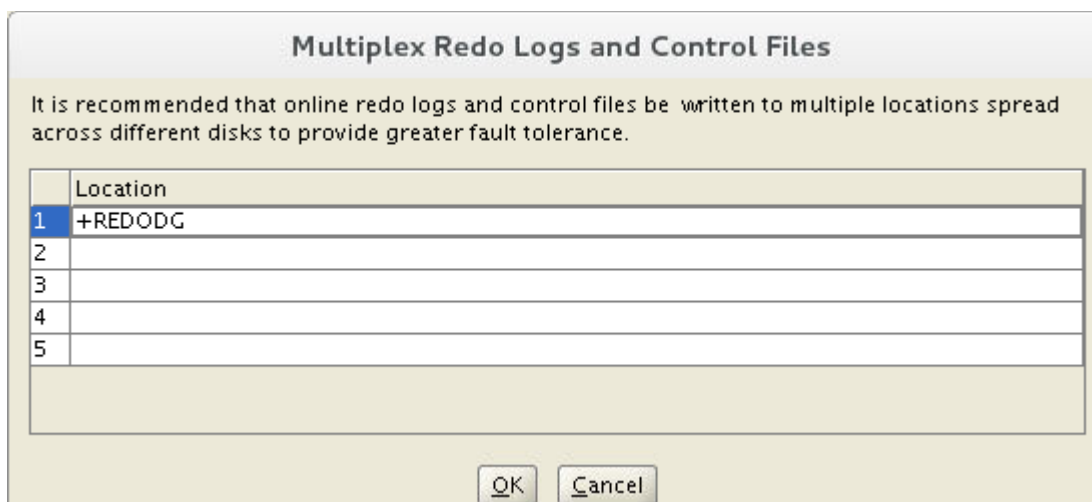
Help < Back Next > Finish Cancel

**Figure 4.1.4.2: Storage Locations Window**

<sup>25</sup> Doc ID [305648.1] – “What is Flash Recovery Area and how to configure it?” at <http://support.oracle.com>



11. Within the same **Storage Locations** window, select the **Multiplex Redo Logs and Control Files** button and within the **Multiplex Redo Logs and Control Files** window, add the Redo Logs diskgroup, i.e. +REDODG and click **OK**, then click **Next** in the **Storage Locations** window.



**Figure 4.1.4.3: Multiplex Redo Logs and Control Files Window**

12. Within the **Database Options** window, all the components should be greyed out. Click **Next**.
13. Within the **Database Vault & Label Security** window, this reference environment unchecks **Configure Database Vault** and **Configure Label Security**. These products provide a level of security with regards to access control decisions at the object level as done by Oracle Database Vault and access control decisions at the row level as done by Oracle Label Security. For more information on these two products, be sure to visit Oracle's Documentation about Oracle Database Vault<sup>26</sup> and Oracle Label Security<sup>27</sup>.

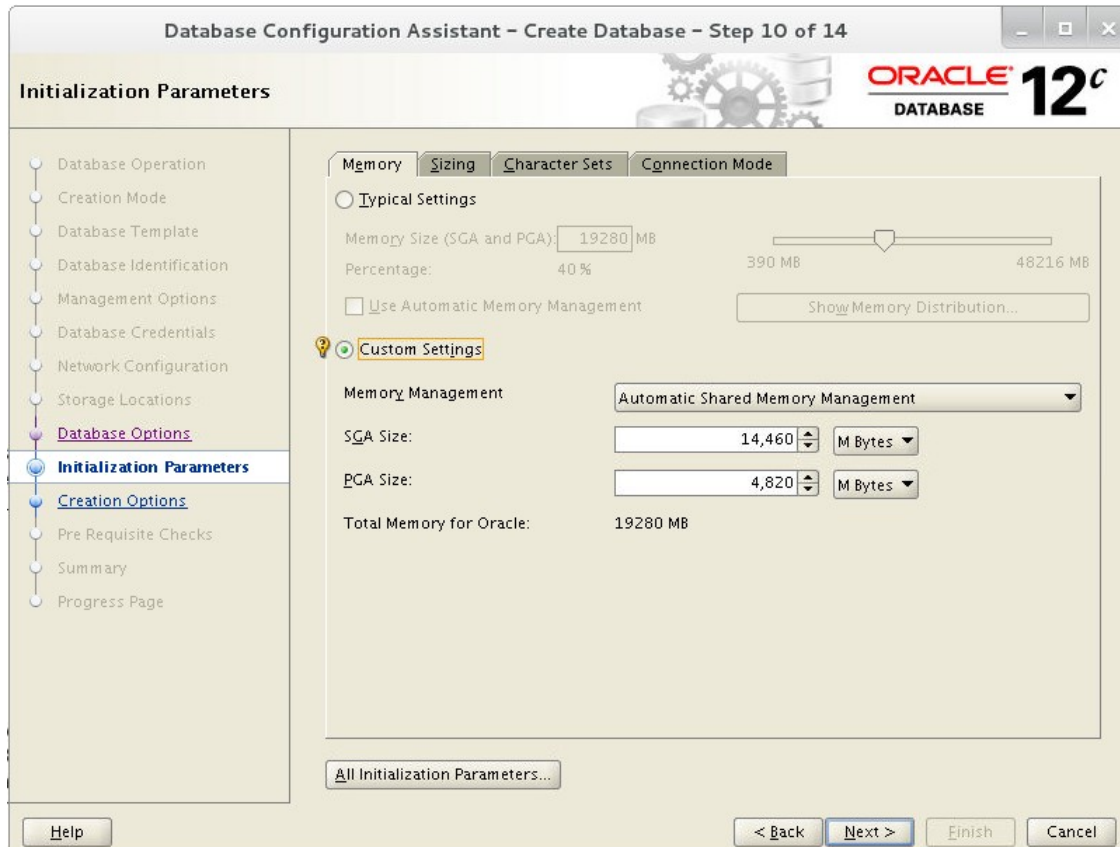
26 Oracle Documentation – Oracle Database Vault  
<http://docs.oracle.com/database/121/DVADM/dvintro.htm#DVADM001>

27 Oracle Documentation – Oracle Label Security  
<http://docs.oracle.com/database/121/OLSAG/intro.htm#OLSAG001>



14. Within the **Initialization Parameters** window under the **Memory** tab, select **Custom Settings** and enter the appropriate values for the SGA and PGA size. It is recommended that the Memory Management be set as **Automatic Shared Memory Management**. The values set for the reference environment with regards to SGA and PGA are the following, however, these values vary for every Oracle database environment.

- SGA – 14460 Megabytes
- PGA – 4820 Megabytes



**Figure 4.1.4.4: Initialization Parameters Window, Memory Tab**

**NOTE:** It is recommended to modify the values of the SGA and PGA based on the Oracle database workload requirements.

15. Within the **Sizing** tab of the **Initialization Parameters** window, appropriately set the block size and maximum number of user processes that can be simultaneously connected to the database. For the purpose of this reference environment, the defaults are used.





**Figure 4.1.4.5: Initialization Parameters Window, Sizing Tab**

**NOTE:** The **Character Sets** tab and **Connection Mode** tab within the **Initialization Parameters** window are not pictorially represented in this reference architecture, however, the default values are set.

16. Within the **Creation Options** window, ensure the **Create Database** box is checked and click **Next**.
17. Within the **Summary** window, review the **Create Database – Summary**, and click **Finish** to start the database creation.
18. Within the **Finish** window, click **Close**.



## 4.1.5 Enabling HugePages

Transparent Huge Pages (THP) are implemented within Red Hat Enterprise Linux 7 to improve memory management by removing many of the difficulties of manually managing huge pages by dynamically allocating huge pages as needed. Red Hat Enterprise Linux 7, by default, uses *transparent* huge pages also known as *anonymous* huge pages. Unlike static huge pages, no additional configuration is needed to use them. Huge pages can boost application performance by increasing the chance a program will have quick access to a memory page. Unlike traditional huge pages, *transparent* huge pages can be swapped out (as smaller 4kB pages) when virtual memory clean up is required.

Unfortunately, Oracle Databases do not take advantage of *transparent* huge pages for interprocess communication. In fact, My Oracle Support [ID 1557478.1]<sup>28</sup> states to disable THP due to unexpected performance issues or delays when THP is found to be enabled. To reap the benefit of huge pages for an Oracle database, it is required to allocate *static* huge pages and disable THP. Due to the complexity of properly configuring huge pages, it is recommended to copy the bash shell script found within **Appendix E Huge Pages Script** and run the script once the database is up and running. The reasoning behind allocating huge pages once the database is up and running is to provide a proper number of pages to handle the running shared memory segments. The steps are as follows:

On **node one** within the Oracle RAC Database environment,

1. Copy the bash script found within **Appendix E Huge Pages Script** and save it as *huge\_pages\_settings.sh*
2. As the root user, ensure the *huge\_pages\_settings.sh* is executable by running the following command:

```
chmod +x huge_pages_settings.sh
```

3. As the root user, execute the *huge\_pages\_settings.sh* script as follows:

```
./huge_pages_settings.sh
Recommended setting within the kernel boot command line: hugepages = <value>
Recommended setting within /etc/security/limits.d/99-grid-oracle-limits.conf:
oracle soft memlock <value>
Recommended setting within /etc/security/limits.d/99-grid-oracle-limits.conf:
oracle hard memlock <value>
```

---

28 ALERT: Disable Transparent HugePages on SLES11,RHEL6,OEL6 and UEK2 Kernels (DOC ID: 1557478.1)





On **each node** within the Oracle RAC Database cluster,

4. Add the number of hugepages provided by the *huge\_pages\_settings.sh* script that was run on Oracle RAC Database node one to the kernel boot command line within the */etc/default/grub* on each Oracle RAC Database node as follows:

```
vi /etc/default/grub
GRUB_TIMEOUT=5
GRUB_DISTRIBUTOR="$(sed 's, release .*$,g' /etc/system-release)"
GRUB_DEFAULT=saved
GRUB_DISABLE_SUBMENU=true
GRUB_TERMINAL_OUTPUT="console"
GRUB_CMDLINE_LINUX="rd.lvm.lv=myvg/swap rd.lvm.lv=myvg/usr
vconsole.font=latarcyrheb-sun16 rd.lvm.lv=myvg/root crashkernel=auto
vconsole.keymap=us rhgb quiet transparent_hugepage=never
hugepages=<value-provided-by-script>"
GRUB_DISABLE_RECOVERY="true"
```

**NOTE:** Allocating the number of huge pages within the kernel boot command line is the most reliable method due to memory not yet becoming fragmented.<sup>29</sup>

5. For the grub changes to take effect, run the following command:

```
grub2-mkconfig -o /boot/grub2/grub.cfg
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-3.10.0-123.el7.x86_64
Found initrd image: /boot/initramfs-3.10.0-123.el7.x86_64.img
Found linux image: /boot/vmlinuz-0-rescue-
41c535c189b842eea5a8c20cbd9bfff26
Found initrd image: /boot/initramfs-0-rescue-
41c535c189b842eea5a8c20cbd9bfff26.img
done
```

6. Oracle requires setting the soft and hard limits to *memlock*. Setting *memlock* allows the oracle user to lock a certain amount of memory from physical RAM that isn't swapped out. The value is expressed in kilobytes and is important from the Oracle perspective because it provides the oracle user permission to use huge pages. This value should be slightly larger than the largest SGA size of any of the Oracle Database instances installed in an Oracle environment. To set *memlock*, add within */etc/security/limits.d/99-grid-oracle-limits.conf* the following:

```
oracle soft memlock <value-provided-by-script>
oracle hard memlock <value-provided-by-script>
```

**NOTE:** This step satisfies the warning message provided during the Oracle Database Software Installation with regards to the Maximum locked memory check.

7. Reboot the system to ensure the huge pages setting takes effect properly.

<sup>29</sup> <https://www.kernel.org/doc/Documentation/vm/hugetlbpage.txt>



8. Verify the value provided by the *huge\_pages\_settings.sh* matches the total number of huge pages available on the system with the following command:

```
cat /proc/meminfo | grep -i hugepages_total
HugePages_Total: <value-provided-by-script>
```

9. Verify the current status of the *transparent* huge pages is set to *NEVER* via the following command:

```
cat /sys/kernel/mm/transparent_hugepage/enabled
always madvise [never]
```

**NOTE:** Starting with Oracle Database version 11.2.0.2, the initialization parameter “USE\_LARGE\_PAGES” was introduced to allocate huge pages on a per database use case. The default value for Oracle Database 12.1.0.2.0 is set to true. **Section 5 Logging into the Oracle Container Database 12c Release 1 (12.1)** shows how to set “USE\_LARGE\_PAGES” to the recommended value of **only** to ensure huge pages are always used upon Oracle database startup. For more information on the parameter and its value refer to My Oracle Support<sup>30</sup>.

**NOTE:** Huge pages is not compatible with Automatic Memory Management (AMM).

---

<sup>30</sup> USE\_LARGE\_PAGES To Enable HugePages [ID 1392497.1]



## 5 Logging into the Oracle Container Database 12c Release 1 (12.1)

This section focuses on ensuring once the Oracle Database 12c Release 1 (12.1) deployment is complete, the `oracle` user can successfully log into the Oracle container database (CDB), and ensure the Oracle database is using the allocated huge pages. The following steps provide the details.

The following example steps are done on node one of the Oracle RAC environment unless otherwise specified.

As the `oracle` user,

1. Set the environment variable for `ORACLE_HOME` with the location of your Oracle Database 12c Release 1 (12.1) home. This reference environment sets `ORACLE_HOME` to `/u01/app/oracle/product/12.1.0/dbhome_1`

```
export ORACLE_HOME=/u01/app/oracle/product/12.1.0/dbhome_1
echo $ORACLE_HOME
/u01/app/oracle/product/12.1.0/dbhome_1
```

**NOTE:** As a precaution, ensure not to include a trailing forward slash (/) when exporting the `ORACLE_HOME`.

2. Set the Oracle System ID (`ORACLE_SID`) used to identify the CDB database.

```
export ORACLE_SID=cdb1
echo $ORACLE_SID
cdb1
```

3. Invoke the **sqlplus** binary to log into the Oracle instance as a `sysdba`.

```
$ORACLE_HOME/bin/sqlplus / as sysdba;
```

```
SQL*Plus: Release 12.1.0.2.0 Production on Tue Nov 18 10:43:55 2014
```

```
Copyright (c) 1982, 2014, Oracle. All rights reserved.
```

```
Connected to:
```

```
Oracle Database 12c Enterprise Edition Release 12.1.0.2.0 - 64bit Production
With the Partitioning, Automatic Storage Management, OLAP, Advanced
Analytics and Real Application Testing options
```

4. Verify the current value of the Oracle parameter ***use\_large\_pages***

```
SQL> show parameter use_large_pages;
```

| NAME            | TYPE   | VALUE |
|-----------------|--------|-------|
| -----           | -----  | ----- |
| use_large_pages | string | TRUE  |



**CAUTION:** Ensure there is enough physical RAM on the nodes within the Oracle RAC Database cluster to place the entire SGA in large pages, otherwise the Oracle RAC database instance(s) will not start. If there is not enough physical RAM on the Oracle RAC Database cluster to place the entire SGA into large pages, leave the default setting of true and ignore the remaining steps in this section.

5. Set the value of the Oracle parameter ***use\_large\_pages*** to the value of ***only***.

```
SQL> alter system set use_large_pages=only scope=spfile sid='*';

System altered.
```

6. On all nodes in the Oracle RAC Database cluster, shutdown the Oracle database instance.

```
SQL> shutdown immediate;
Database closed.
Database dismounted.
ORACLE instance shut down.
```

7. On all nodes in the Oracle RAC Database cluster, restart the Oracle database instance.

```
SQL> startup;
ORACLE instance started.

Total System Global Area 2.0176E+10 bytes
Fixed Size 3478432 bytes
Variable Size 2952793184 bytes
Database Buffers 1.7180E+10 bytes
Redo Buffers 40247296 bytes
Database mounted.
Database opened.
```

**NOTE:** Ensure to set the appropriate ORACLE\_SID and ORACLE\_HOME within each node.

8. Verify the current value of the Oracle parameter ***use\_large\_pages*** is now set to ***only***.

```
SQL> show parameter use_large_pages;
```

| NAME            | TYPE   | VALUE |
|-----------------|--------|-------|
| -----           | -----  | ----- |
| use_large_pages | string | ONLY  |



9. Open the container database's alert log, named `alert_<name-of-cdb>.log`, located under the `$ORACLE_BASE/diag/rdbms/<name-of-cdb>/<name-of-cdb-instance#>/trace/` using a text editor, such as `vi`, and search for the following snippet to ensure that the System Global Area (SGA) is 100% in large pages.

`$ORACLE_BASE/diag/rdbms/<name-of-cdb>/<name-of-cdb-instance#>/trace/alert_<name-of-cdb>.log`

```

Tue Nov 18 10:45:41 2014
Dump of system resources acquired for SHARED GLOBAL AREA (SGA)

Tue Nov 18 10:45:41 2014
Per process system memlock (soft) limit = 19G
Tue Nov 18 10:45:41 2014
Expected per process system memlock (soft) limit to lock
SHARED GLOBAL AREA (SGA) into memory: 19G
Tue Nov 18 10:45:41 2014
Available system pagesizes:
4K, 2048K
Tue Nov 18 10:45:41 2014
Supported system pagesize(s):
Tue Nov 18 10:45:41 2014
 PAGESIZE AVAILABLE_PAGES EXPECTED_PAGES ALLOCATED_PAGES ERROR(s)
Tue Nov 18 10:45:41 2014
 2048K 9668 9665 9665 NONE
Tue Nov 18 10:45:41 2014
Reason for not supporting certain system pagesizes:
Tue Nov 18 10:45:41 2014
4K - Large pagesizes only
Tue Nov 18 10:45:41 2014

```

**NOTE:** This reference environment's SGA size is set to 19 GB, however, this value will vary depending on the value provided when creating an Oracle RAC database using `dbca`.



## 6 Post Installation Cleanup Tasks

This section describes the removal of certain tasks that were set in order to satisfy the prerequisites of the Oracle Universal Installer (OUI). While these prerequisites are required, the OUI does not handle them appropriately for a Red Hat Enterprise Linux 7 distribution.

### 6.1.1 Removal of Oracle Kernel Parameters from `/etc/sysctl.conf`

Changing kernel parameters within the `/etc/sysctl.conf` file directly can leave a system susceptible to having kernel parameter entries being overwritten, replaced, or removed all together if there were to be an update to the initscripts RPM. If not using the oracle tuned profile as described in this reference architecture, please ensure to have all the oracle related kernel parameters set within `98-oracle.conf` file inside the `/etc/sysctl.d` directory. The reason they are set initially in the `/etc/sysctl.conf` file is due to Oracle's Universal Installer checking the file `/etc/sysctl.conf` for kernel parameter settings instead of checking the kernel running values.

**NOTE:** Files in `/etc/sysctl.d` are read in order, i.e. `1-example.conf` is read before `2-example.conf`, hence it is possible for two files to have conflicting settings. The last file read, takes precedence.

All nodes within the Oracle RAC cluster, as root user,

1. Edit the `/etc/sysctl.conf` using an editor such as **vi** and remove all Oracle parameter related entries. The `/etc/sysctl.conf` file should be blank and all settings should be set in a `.conf` file within the `/etc/sysctl.d` or set by a tuned profile such as the oracle profile found in `/usr/lib/tuned` from **Section 3.4.4 Optimizing Database Storage using Automatic System Tuning**. This reference environment, uses the custom tuned oracle profile for kernel parameter settings.
2. Save the `/etc/sysctl.conf` file and run the following command to set the changes:

```
sysctl -p /etc/sysctl.conf
```

3. Verify the Oracle tuned profile is still set by:

```
tuned-adm active
Current active profile: oracle
```

**NOTE:** If not using the custom tuned profile oracle, please ensure to have all the oracle kernel parameters set within the `/etc/sysctl.d/98-oracle.conf` file

4. Verify the Oracle kernel parameters are still set by run the following command:

```
sysctl -a | grep <oracle kernel parameter>
```

Example of checking kernel.shmmax:

```
sysctl -a | grep kernel.shmmax
kernel.shmmax = 4398046511104
```



## 6.1.2 Removal of `firewalld` Trusted Source Address

During the installation of the Oracle Grid Infrastructure 12c Release 1, the installation will fail during the **Network Interface Usage** step if any type of firewall is running. Instead of completely disabling the firewall, the source address of our bonded interface, `bond0`, is added to the trusted zone within `firewalld`. However, this should be immediately removed upon completion of the installation of Oracle as `bond0` is a public interface and should not reside in the *trusted* zone. For clarification, interfaces `ens10f2` and `ens10f3` are private interfaces that reside on a private network not accessible to the outside world. Due to this, it is acceptable for these interfaces to be part of the *trusted* zone.

On all nodes within the Oracle RAC cluster as the `root` user, remove the source public address of the other nodes in the cluster. For example, this reference environment consists of a two node Oracle RAC cluster, with node one's public IP set to 10.19.142.51 and node two's public IP set to 10.19.142.52

Within node one of the Oracle RAC cluster, run the following command:

```
firewall-cmd --permanent --zone=trusted --remove-source=10.19.142.52/21
success

systemctl restart firewalld.service
```

Within node two of the Oracle RAC cluster, run the following command:

```
firewall-cmd --permanent --zone=trusted --remove-source=10.19.142.51/21
success

systemctl restart firewalld.service
```

With the removal of the public source addresses, Oracle RAC Database can now run properly with a firewall protecting the public interface.



# 7 Common Tasks when Managing Container Database (CDB) and Pluggable Databases (PDB)

This section describes tasks that are commonly used when dealing with a CDB and PDBs. The tasks covered within this section are as follows:

- Connect to a CDB
- Connect to a PDB
- Managing a CDB
- Managing a PDB
- Location of Data files in a CDB & PDB

## 7.1 Connect to a CDB

Example on how to connect to node one of an Oracle RAC Database cluster.

As the oracle user,

1. Set the environment variable for *ORACLE\_HOME* with the location of your Oracle Database 12c Release 1 (12.1) home. This reference environment sets *ORACLE\_HOME* to */u01/app/oracle/product/12.1.0/dbhome\_1*

```
export ORACLE_HOME=/u01/app/oracle/product/12.1.0/dbhome_1
echo $ORACLE_HOME
/u01/app/oracle/product/12.1.0/dbhome_1
```

**NOTE:** As a precaution, ensure not to include a trailing forward slash (/) when exporting the *ORACLE\_HOME*.

2. Set the Oracle System ID (*ORACLE\_SID*) used to identify the CDB database.

```
export ORACLE_SID=cdb1
echo $ORACLE_SID
cdb1
```





3. Invoke the **sqlplus** binary to log into the Oracle instance as a *sysdba*.

```
$ORACLE_HOME/bin/sqlplus / as sysdba;

SQL*Plus: Release 12.1.0.2.0 Production on Tue Nov 18 10:55:19 2014

Copyright (c) 1982, 2014, Oracle. All rights reserved.

Connected to:
Oracle Database 12c Enterprise Edition Release 12.1.0.2.0 - 64bit
Production
With the Partitioning, Automatic Storage Management, OLAP, Advanced
Analytics
and Real Application Testing options
```

4. Once connected, verify that the instance is connected to the root container, *CDB\$ROOT* with a *CON\_ID* is 1.

```
SQL> show con_name

CON_NAME

CDB$ROOT

SQL> show con_id

CON_ID

1
```

**NOTE:** The *CDB\$ROOT* connection ID is always set to one.

5. To list all the available services and PDBs within the CDB:

```
SQL> select name, con_id from v$active_services;

NAME CON_ID

pdb2 4
pdb1 3
cdbXDB 1
cdb 1
SYS$BACKGROUND 1
SYS$USERS 1

6 rows selected.
```



## 7.2 Connect to a PDB

The syntax to connect to a PDB varies depending on whether or not there is an entry within the *tnsnames.ora* file for the PDB.

Without an entry to the *tnsnames.ora* file, the syntax to connect to a PDB named *pdb* with an instance name of *pdb1* from node one of the Oracle RAC cluster is as follows:

```
$ORACLE_HOME/bin/sqlplus sys/<password>@db-oracle-node1:1521/pdb1 as
sysdba;
SQL*Plus: Release 12.1.0.2.0 Production on Tue Nov 18 10:57:15 2014

Copyright (c) 1982, 2014, Oracle. All rights reserved.

Connected to:
Oracle Database 12c Enterprise Edition Release 12.1.0.2.0 - 64bit
Production
With the Partitioning, Automatic Storage Management, OLAP, Advanced
Analytics
and Real Application Testing options
```

**NOTE:** The value 1521, represents the Oracle Listener port.

With an entry to the *tnsnames.ora* file, the syntax to connect to a PDB instance labeled *pdb1* is as follows:

```
$ORACLE_HOME/bin/sqlplus sys/<password>@PDB1 as sysdba;
```

A snippet of the entry found within the *tnsnames.ora* file is displayed below:

*/u01/app/oracle/product/12.1.0/dbhome\_1/network/admin/tnsnames.ora*

```
PDB1 =
 (DESCRIPTION =
 (ADDRESS = (PROTOCOL = TCP)(HOST = db-oracle-
scan.cloud.lab.eng.bos.redhat.com)(PORT = 1521))
 (CONNECT_DATA =
 (SERVER = DEDICATED)
 (SERVICE_NAME = pdb1)
)
)
```



## 7.3 Managing a CDB

The process of starting and shutting down a CDB database instance is similar to the steps done in previous Oracle database versions for traditional databases. The key difference is to verify that the connection is to the root container prior to shutting down or starting up the Oracle database.

1. Connect to the CDB database as a `sysdba` using **sqlplus**. The steps are the same as shown in **Section 7.1 Connect to a CDB** steps one through three.
2. Once connected, verify the instance is the root container `CDB$ROOT` via the command:

```
SQL> show con_name;
```

```
CON_NAME
```

```

```

```
CDB$ROOT
```

3. To shutdown the Oracle CDB database use the following command:

```
SQL> shutdown immediate;
```

```
Database closed.
```

```
Database dismounted
```

```
ORACLE instance shut down.
```

4. To start the Oracle CDB database use the following command:

```
SQL> startup;
```

The startup command starts the instance, mounts the control files, and then opens the root container.



## 7.4 Managing a PDB

This section focuses on verifying the *OPEN\_MODE* of a PDB, how to open and close a specific PDB, and how to open and close all PDBs within a CDB.

1. To verify the *open\_mode* status of all the PDBs, while logged in as a *SYSDBA* in the CDB, use the following command:

```
SQL> select name, open_mode from v$pdb;
```

| NAME      | OPEN_MODE |
|-----------|-----------|
| PDB\$SEED | READ ONLY |
| PDB1      | MOUNTED   |
| PDB2      | MOUNTED   |

2. When a PDB is closed, the *OPEN\_MODE* is set to *MOUNTED*. To open a PDB and verify the new *OPEN\_MODE* of *READ WRITE*, run the following SQL syntax while logged in as a *SYSDBA* in the CDB:

```
SQL> alter pluggable database pdb1 open;
```

Pluggable database altered.

```
SQL> select name, open_mode from v$pdb;
```

| NAME      | OPEN_MODE  |
|-----------|------------|
| PDB\$SEED | READ ONLY  |
| PDB1      | READ WRITE |
| PDB2      | MOUNTED    |

3. To open all the PDBs connected to a CDB and verify the new *OPEN\_MODE* of *READ WRITE*, run the following SQL syntax while logged in as a *SYSDBA* in the CDB:

```
SQL> alter pluggable database all open;
```

Pluggable database altered.

```
SQL> select name, open_mode from v$pdb;
```

| NAME      | OPEN_MODE  |
|-----------|------------|
| PDB\$SEED | READ ONLY  |
| PDB1      | READ WRITE |
| PDB2      | READ WRITE |



4. To drop a particular PDB i.e. *pdb2*, and its data files, execute the following SQL syntax while logged in as a *SYSDBA* in the CDB:

```
SQL > alter pluggable database pdb2 close immediate;
```

Pluggable database altered.

```
SQL > drop pluggable database pdb2 including datafiles;
```

Pluggable database dropped.

5. To verify if the pluggable database with the name *pdb2* has been dropped:

```
SQL> select name, open_mode from v$pdb;
```

| NAME      | OPEN_MODE  |
|-----------|------------|
| -----     |            |
| PDB\$SEED | READ ONLY  |
| PDB1      | READ WRITE |

## 7.5 Location of Data files of PDBs & CDB

The following section shows how to identify tablespace names, data files associated with the CDB and PDBs, including their temporary files.

1. Connect to the CDB database as a *SYSDBA* using **sqlplus**. The steps are the same as shown in **Section 7.1 Connect to a CDB** steps one through three.
2. To identify the tablespaces associated with the CDB or any of the PDBs installed, use the following syntax where the *con\_id* varies upon the database chosen. The example below uses the *con\_id* of 1 to show the CDB tablespaces.

```
SQL > select tablespace_name, con_id from cdb_tablespaces where con_id = 1;
```

| TABLESPACE_NAME | CON_ID |
|-----------------|--------|
| -----           |        |
| SYSTEM          | 1      |
| SYSAUX          | 1      |
| UNDOTBS1        | 1      |
| TEMP            | 1      |
| USERS           | 1      |

3. To locate the data files from the CDB or PDBs installed, use the following syntax where the *con\_id* varies upon the database chosen. The example below uses the *con\_id* of 1 to show the CDB data file locations.

```
SQL> select file_name, con_id from cdb_data_files where con_id=1;
```

| FILE_NAME                                   | CON_ID |
|---------------------------------------------|--------|
| -----                                       |        |
| +DATADG/CDB/DATAFILE/system.270.836232071   | 1      |
| +DATADG/CDB/DATAFILE/sysaux.273.836232077   | 1      |
| +DATADG/CDB/DATAFILE/undotbs1.262.836232081 | 1      |



| FILE_NAME                                | CON_ID |
|------------------------------------------|--------|
| -----                                    |        |
| +DATADG/CDB/DATAFILE/users.275.836232097 | 1      |

4. To locate the temporary files from the CDB or PDBs installed, use the following syntax where the *con\_id* varies upon the database chosen. The example below uses the *con\_id* of 1 to show the CDB data file locations.

```
SQL> select file_name, con_id from cdb_temp_files where con_id =1 ;
```

| FILE_NAME                               | CON_ID |
|-----------------------------------------|--------|
| -----                                   |        |
| +DATADG/CDB/TEMPFILE/temp.278.836232081 | 1      |



## 8 Conclusion

Red Hat Enterprise Linux 7 provides an excellent foundation for database deployments with demonstrated stability, scalability, and performance. With the support for Oracle RAC Database 12c Release 1 (12.1) on Red Hat Enterprise Linux 7, customers can increasingly look to deploy Oracle RAC Databases in advanced configurations.

The steps and procedures described in this reference architecture should provide system, database, and storage administrators the blueprint required to create a robust and performing solution based on Oracle RAC Databases. Administrators can reference this document to simplify and optimize the deployment process and employ the latest best practices for configuring Red Hat technologies while implementing the following tasks:

- Deploying Oracle Grid Infrastructure 12c Release 1 (12.1.0.2.0)
- Deploying Oracle Database Software 12c Release 1 (12.1.0.2.0)
- Deploying an Oracle RAC Database 12c Release 1 (12.1.0.2.0) with shared SAN disks
- Using Oracle ASM disks with *udev* rules
- Securing the Oracle Database 12c environment with *SELinux*
- Securing the public network using *firewalld*

For any questions or concerns, please email [refarch-feedback@redhat.com](mailto:refarch-feedback@redhat.com) and ensure to visit the Red Hat Reference Architecture page at <http://red.ht/1Nd64ha> to find out about all of our Red Hat solution offerings.



# Appendix A: Revision History

Revision 1.1

Friday March 05, 2015

Roger Lopez

- Within the Conclusion, fixed hyperlink of Reference Architectures to point to

<http://red.ht/1Nd64ha>

- Fixed some missing dependencies within the .tar.gz file that were already correctly reflexed in the scripts within the Appendix sections. Details within the CHANGELOG of the .tar.gz

Revision 1.0

Friday February 20, 2015

Roger Lopez

Initial Release





## Appendix B: Contributors

1. Tom Tracy, content reviewer
2. David Wilson, content reviewer
3. John Boero, content reviewer
4. Yan Fisher, content reviewer
5. Scott McBrien, content reviewer



# Appendix C: DM Multipath Configuration File

```
This is a basic configuration file with some examples, for device mapper
multipath.
For a complete list of the default configuration values, see
/usr/share/doc/device-mapper-multipath-0.4.9/multipath.conf.defaults
For a list of configuration options with descriptions, see
/usr/share/doc/device-mapper-multipath-0.4.9/multipath.conf.annotated
#
REMEMBER: After updating multipath.conf, you must run
#
service multipathd reload
#
for the changes to take effect in multipathd

By default, devices with vendor = "IBM" and product = "S/390.*" are
blacklisted. To enable mulitpathing on these devices, uncomment the
following lines.
#blacklist_exceptions {
device {
vendor "IBM"
product "S/390.*"
}
#}

Here is an example of how to configure some standard options.
##
#
defaults {
 udev_dir /dev
 polling_interval 10
 path_selector "round-robin 0"
 path_grouping_policy multibus
 getuid_callout "/lib/udev/scsi_id --whitelisted
--device=/dev/%n"
 prio alua
 path_checker readsector0
 rr_min_io 100
 max_fds 8192
 rr_weight priorities
 failback immediate
 no_path_retry fail
 user_friendly_names yes
}
##
The wwid line in the following blacklist section is shown as an example
of how to blacklist devices by wwid. The 2 devnode lines are the
compiled in default blacklist. If you want to blacklist entire types
of devices, such as all scsi devices, you should use a devnode line.
However, if you want to blacklist specific devices, you should use
a wwid line. Since there is no guarantee that a specific device will
```



```
not change names on reboot (from /dev/sda to /dev/sdb for example)
devnode lines are not recommended for blacklisting specific devices.
##
blacklist {
 wwid 3600508b1001030353434363646301200
 devnode "^(ram|raw|loop|fd|md|dm-|sr|scd|st)[0-9]*"
 devnode "^hd[a-z]"
}
multipaths {
 multipath {
 wwid 3600c0ff000d7e7a899d8515101000000
 alias db1
 }
 multipath {
 wwid 3600c0ff000dabfe5a7d8515101000000
 alias db2
 }
 multipath {
 wwid 3600c0ff000d7e7a8dbd8515101000000
 alias fra
 }
 multipath {
 wwid 3600c0ff000dabfe5f4d8515101000000
 alias redo
 }
}
#devices {
device {
vendor "COMPAQ "
product "HSV110 (C)COMPAQ"
path_grouping_policy multibus
getuid_callout "/lib/udev/scsi_id --whitelisted"
--device="/dev/%n"
path_checker readsector0
path_selector "round-robin 0"
hardware_handler "0"
failback 15
rr_weight priorities
no_path_retry queue
}
device {
vendor "COMPAQ "
product "MSA1000 "
path_grouping_policy multibus
}
#}
```



## Appendix D: Firewalld Configuration File

*/etc/firewalld/zones/public.xml*

```
<?xml version="1.0" encoding="utf-8"?>
<zone>
 <short>Public</short>
 <description>For use in public areas. You do not trust the other computers
on networks to not harm your computer. Only selected incoming connections
are accepted.</description>
 <service name="dhcpv6-client"/>
 <service name="ssh"/>
 <rule family="ipv4">
 <source address="10.19.142.54"/>
 <port protocol="tcp" port="1521"/>
 <accept/>
 </rule>
 <rule family="ipv4">
 <source address="10.19.142.54"/>
 <port protocol="tcp" port="5500"/>
 <accept/>
 </rule>
</zone>
```

*/etc/firewalld/zones/trusted.xml*

```
<?xml version="1.0" encoding="utf-8"?>
<zone target="ACCEPT">
 <short>Trusted</short>
 <description>All network connections are accepted.</description>
 <interface name="<enter-private-interface-name2>"/>
 <interface name="<enter-private-interface-name2>"/>
</zone>
```



## Appendix E: Huge Pages Script

The following huge pages script is from Tuning Red Hat Enterprise Linux For Oracle & Oracle RAC by Scott Crot, Sr. Consultant, Red Hat<sup>31</sup> and modified to include the values Oracle's soft memlock, hard memlock, and work with kernel 3.10

```
#!/bin/bash
KERN=`uname -r | awk -F. '{ printf("%d.%d\n", $1, $2); }'`
Find out the HugePage size
HPG_SZ=`grep Hugepagesize /proc/meminfo | awk '{print $2}'`
Start from 1 pages to be on the safe side and guarantee 1 free HugePage
NUM_PG=1
Cumulative number of pages required to handle the running shared memory
segments
for SEG_BYTES in `ipcs -m | awk '{print $5}' | grep "[0-9][0-9]*"`
do
MIN_PG=`echo "$SEG_BYTES/($HPG_SZ*1024)" | bc -q`
if [$MIN_PG -gt 0]; then
NUM_PG=`echo "NUM_PG+MIN_PG+1" | bc -q`
fi
done
Finish with results
case $KERN in
'2.4') HUGETLB_POOL=`echo "NUM_PG*HPG_SZ/1024" | bc -q`;
echo "Recommended setting: vm.hugetlb_pool = $HUGETLB_POOL" ;;
'2.6') MEM_LOCK=`echo "NUM_PG*HPG_SZ" | bc -q`;
echo "Recommended setting within the kernel boot command line: hugepages =
$NUM_PG"
echo "Recommended setting within /etc/security/limits.d/99-grid-oracle-
limits.conf: oracle soft memlock $MEM_LOCK"
echo "Recommended setting within /etc/security/limits.d/99-grid-oracle-
limits.conf: oracle hard memlock $MEM_LOCK" ;;
'3.10') MEM_LOCK=`echo "NUM_PG*HPG_SZ" | bc -q`;
echo "Recommended setting within the kernel boot command line: hugepages =
$NUM_PG"
echo "Recommended setting within /etc/security/limits.d/99-grid-oracle-
limits.conf: oracle soft memlock $MEM_LOCK"
echo "Recommended setting within /etc/security/limits.d/99-grid-oracle-
limits.conf: oracle hard memlock $MEM_LOCK" ;;
*) echo "Unrecognized kernel version $KERN. Exiting." ;;
esac
```

31 Tuning Red Hat Enterprise Linux For Oracle & Oracle RAC by Scott Crot, Sr. Consultant, Red Hat, [http://www.redhat.com/promo/summit/2010/presentations/summit/decoding-the-code/fri/scott-945-tuning/summit\\_jbw\\_2010\\_presentation.pdf](http://www.redhat.com/promo/summit/2010/presentations/summit/decoding-the-code/fri/scott-945-tuning/summit_jbw_2010_presentation.pdf)



## Appendix F: Oracle Database Package Requirements Text File

```
binutils
compat-libcap1
compat-libstdc++-33
gcc
gcc-c++
glibc
glibc-devel
ksh
libgcc
libstdc++
libstdc++-devel
libaio
libaio-devel
libXext
libXtst
libX11
libXau
libxcb
libXi
make
sysstat
libXmu
libXt
libXv
libXxf86dgak
libXxf86misc
libXxf86vm
xorg-x11-utils
xorg-x11-xauth
nfs-utils
```



## Appendix G: Kernel Parameters

```
vm.swappiness = 1
vm.dirty_background_ratio = 3
vm.dirty_ratio = 80
vm.dirty_expire_centisecs = 500
vm.dirty_writeback_centisecs = 100
kernel.shmmax = 4398046511104
kernel.shmall = 1073741824
kernel.shmmni = 4096
kernel.sem = 250 32000 100 128
fs.file-max = 6815744
fs.aio-max-nr = 1048576
net.ipv4.ip_local_port_range = 9000 65500
net.core.rmem_default = 262144
net.core.rmem_max = 4194304
net.core.wmem_default = 262144
net.core.wmem_max = 1048576
net.ipv4.conf.<priv-interface-name1>.rp_filter = 2
net.ipv4.conf.<priv-interface-name2>.rp_filter = 2
kernel.panic_on_oops = 1
```



## Appendix H: 98-oracle.conf

```
98-oracle.conf resides in the /etc/sysctl.d directory
Replace <priv-interface-name{1-2}> with the appropriate Ethernet device
name
net.ipv4.conf.<priv-interface-name1>.rp_filter = 2
net.ipv4.conf.<priv-interface-name2>.rp_filter = 2

If not using the custom oracle tuned profile
please uncomment the following kernel parameters.
vm.swappiness = 1
vm.dirty_background_ratio = 3
vm.dirty_ratio = 80
vm.dirty_expire_centisecs = 500
vm.dirty_writeback_centisecs = 100
kernel.shmmax = 4398046511104
kernel.shmall = 1073741824
kernel.shmmni = 4096
kernel.sem = 250 32000 100 128
fs.file-max = 6815744
fs.aio-max-nr = 1048576
net.ipv4.ip_local_port_range = 9000 65500
net.core.rmem_default = 262144
net.core.rmem_max = 4194304
net.core.wmem_default = 262144
net.core.wmem_max = 1048576
kernel.panic_on_oops = 1
```





## Appendix I: Limits Configuration File (99-grid-oracle-limits.conf)

```
oracle soft nproc 16384 #Ora bug 15971421
oracle hard nproc 16384
oracle soft nofile 1024
oracle hard nofile 65536
oracle soft stack 10240
oracle hard stack 32768
oracle soft memlock <value-provided-by-script>
oracle hard memlock <value-provided-by-script>

grid soft nproc 16384 #Ora bug 15971421
grid hard nproc 16384
grid soft nofile 1024
grid hard nofile 65536
grid soft stack 10240
grid hard stack 32768
```



## Appendix J: 99-oracle-asmdevices.rules

```
KERNEL=="dm-*", ENV{DM_UUID}=="<enter-value-according-to-your-
environment>", OWNER="grid", GROUP="asmadmin", MODE="0660"
KERNEL=="dm-*", ENV{DM_UUID}=="<enter-value-according-to-your-
environment>", OWNER="grid", GROUP="asmadmin", MODE="0660"
KERNEL=="dm-*", ENV{DM_UUID}=="<enter-value-according-to-your-
environment>", OWNER="grid", GROUP="asmadmin", MODE="0660"
KERNEL=="dm-*", ENV{DM_UUID}=="<enter-value-according-to-your-
environment>", OWNER="grid", GROUP="asmadmin", MODE="0660"
```



## Appendix K: Sample Kickstart File

```
Red Hat | Oracle Solutions Kickstart Script
```

```
install
url --url=<place-distro-url-here>
lang en_US.UTF-8
keyboard us
network --onboot yes --device em1 --mtu=1500 --bootproto dhcp
rootpw <password-for-system>
Reboot after installation
reboot
authconfig --enablemd5 --enablesshadow
selinux --enforcing
timezone America/New_York
bootloader --location=mbr --driveorder=sda --append="crashkernel=auto rhgb
quiet"
The following is the partition information you requested
Note that any partitions you deleted are not expressed
here so unless you clear all partitions first, this is
not guaranteed to work
clearpart --all

volgroup myvg --pesize=32768 pv.008002
logvol /home --fstype=ext4 --name=home --vgname=myvg --size=8192
logvol / --fstype=ext4 --name=root --vgname=myvg --size=15360
logvol swap --name=swap --vgname=myvg --size=16400
logvol /tmp --fstype=ext4 --name=tmp --vgname=myvg --size=4096
logvol /u01 --fstype=ext4 --name=u01 --vgname=myvg --size=51200
logvol /usr --fstype=ext4 --name=usr --vgname=myvg --size=5120
logvol /var --fstype=ext4 --name=var --vgname=myvg --size=8192

part /boot --fstype=ext4 --size=256
part pv.008002 --grow --size=1000

%packages
@Base
@Core
```



# Appendix L: Custom Oracle Tuned Profile (script.sh & tuned.conf)

/usr/lib/tuned/oracle/script.sh

```
#!/bin/bash

Oracle ktune script

. /usr/lib/tuned/functions

start() {

 CURR_SHMALL=`sysctl -n kernel.shmall`
 CURR_SHMMAX=`sysctl -n kernel.shmmax`
 CURR_SHMMNI=`sysctl -n kernel.shmmni`
 CURR_SEM=`sysctl -n kernel.sem`
 CURR_SWAPPINESS=`sysctl -n vm.swappiness`
 CURR_DIRTY_BCKGRD_RATIO=`sysctl -n vm.dirty_background_ratio`
 CURR_DIRTY_RATIO=`sysctl -n vm.dirty_ratio`
 CURR_DIRTY_EXPIRE_CENTISECS=`sysctl -n vm.dirty_expire_centisecs`
 CURR_FILE_MAX=`sysctl -n fs.file-max`
 CURR_AIO_MAX_NR=`sysctl -n fs.aio-max-nr`
 CURR_LCL_PRT_RNG=`sysctl -n net.ipv4.ip_local_port_range`
 CURR_RMEM=`sysctl -n net.core.rmem_default`
 CURR_RMEM_MAX=`sysctl -n net.core.rmem_max`
 CURR_WMEM=`sysctl -n net.core.wmem_default`
 CURR_WMEM_MAX=`sysctl -n net.core.wmem_max`
 CURR_PANIC_ON_OOPS=`sysctl -n kernel.panic_on_oops`

 save_value kernel.shmall "$CURR_SHMALL"
 save_value kernel.shmmax "$CURR_SHMMAX"
 save_value kernel.shmmni "$CURR_SHMMNI"
 save_value vm.swappiness "$CURR_SWAPPINESS"
 save_value vm.dirty_background_ratio "$CURR_DIRTY_BCKGRD_RATIO"
 save_value vm.dirty_ratio "$CURR_DIRTY_RATIO"
 save_value vm.dirty_expire_centisecs "$CURR_DIRTY_EXPIRE_CENTISECS"
 save_value fs.file-max "$CURR_FILE_MAX"
 save_value fs.aio-max-nr "$CURR_AIO_MAX_NR"
 save_value net.ipv4.ip_local_port_range "$CURR_LCL_PRT_RNG"
 save_value net.core.rmem_default "$CURR_RMEM"
 save_value net.core.rmem_max "$CURR_RMEM_MAX"
 save_value net.core.wmem_default "$CURR_WMEM"
 save_value net.core.wmem_max "$CURR_WMEM_MAX"
 save_value kernel.panic_on_oops "$CURR_PANIC_ON_OOPS"

 return 0
}

stop() {
 SHMALL=`restore_value kernel.shmall`
}
```



```

SHMMAX=`restore_value kernel.shmmax`
SHMMNI=`restore_value kernel.shmmni`
SWAPPINESS=`restore_value vm.swappiness`
DIRTY_BCKGRD_RATIO=`restore_value vm.dirty_background_ratio`
DIRTY_EXPIRE_CENTISECS=`restore_value vm.dirty_expire_centisecs`
FILE_MAX=`restore_value fs.file-max`
LCL_PRT_RNG=`restore_value net.ipv4.ip_local_port_range`
RMEM_DEFAULT=`restore_value net.core.rmem_default`
RMEM_MAX=`restore_value net.core.rmem_max`
WMEM_DEFAULT=`restore_value net.core.wmem_default`
WMEM_MAX=`restore_value net.core.wmem_max`
PANIC_ON_OOPS=`restore_value kernel.panic_on_oops`

["$SHMALL"] && sysctl -w kernel.shmall="$SHMALL"
["$SHMMAX"] && sysctl -w kernel.shmmax="$SHMMAX"
["$SHMMNI"] && sysctl -w kernel.shmmni="$SHMMNI"
["$SWAPPINESS"] && sysctl -w vm.swappiness="$SWAPPINESS"
["$DIRTY_BCKGRD_RATIO"] && sysctl -w
vm.dirty_background_ratio="$DIRTY_BCKGRD_RATIO"
["$DIRTY_EXPIRE_CENTISECS"] && sysctl -w
vm.dirty_expire_centisecs="$DIRTY_EXPIRE_CENTISECS"
["$FILE_MAX"] && sysctl -w fs.file-max="$FILE_MAX"
["$LCL_PRT_RNG"] && sysctl -w
net.ipv4.ip_local_port_range="$LCL_PRT_RNG"
["$RMEM_DEFAULT"] && sysctl -w
net.core.rmem_default="$RMEM_DEFAULT"
["$RMEM_MAX"] && sysctl -w net.core.rmem_max="$RMEM_MAX"
["$WMEM_DEFAULT"] && sysctl -w
net.core.wmem_default="$WMEM_DEFAULT"
["$WMEM_MAX"] && sysctl -w net.core.wmem_max="$WMEM_MAX"
["$PANIC_ON_OOPS"] && sysctl -w
kernel.panic_on_oops="$PANIC_ON_OOPS"

 return 0
}

process $@

```

*/usr/lib/tuned/oracle/tuned.conf*

```

#
tuned configuration
#

[main]
include=throughput-performance

[sysctl]
vm.swappiness = 1
vm.dirty_background_ratio = 3
vm.dirty_ratio = 80
vm.dirty_expire_centisecs = 500

```



```
vm.dirty_writeback_centisecs = 100
kernel.shmmax = 4398046511104
kernel.shmall = 1073741824
kernel.shmmni = 4096
kernel.sem = 250 32000 100 128
fs.file-max = 6815744
fs.aio-max-nr = 1048576
net.ipv4.ip_local_port_range = 9000 65500
net.core.rmem_default = 262144
net.core.rmem_max = 4194304
net.core.wmem_default = 262144
net.core.wmem_max = 1048576
kernel.panic_on_oops = 1

[vm]
transparent_hugepages=never

[script]
script=script.sh
```



## Appendix M: Configuration Files

All configuration files can be downloaded from the Red Hat customer portal<sup>32</sup>. A listing of all the files and a brief description can be seen on the **Table 8.1: Configuration Files**. Some of the configuration files require input with the proper information pertaining to your environment.

Files	Description
req-rpm.txt	The required RPMs to deploy Oracle.
huge_page_settings.sh	Script that provides the proper hugepage values to set.
multipath.conf	Device Mapper Multipath configuration file.
98-oracle.conf	Configuration file for the kernel parameters related to <code>rp_filter</code>
99-grid-oracle-limits.conf	Configuration file to set limits for a user.
99-oracle-asmdevices.rules	<code>udev</code> configuration file for Oracle ASM disks
public.xml	firewalld XML configuration
trusted.xml	firewalld XML configuration
oracle-grid.sh	Shell script used to set user limits
sample-ks.cfg	Sample Kickstart File
oracle	Directory of the custom tuned profile labeled oracle
CHANGELOG	Listing of the latest changes made to the .tar.gz file
script.sh	Oracle tuned profile script
tuned.conf	Configuration file for oracle tuned profile

**Table 8.1: Configuration Files**

<sup>32</sup> <https://access.redhat.com/node/1357883/40/1>



# Appendix N: Troubleshooting ORA-\* Errors

This section focuses on using the command line tool, Automatic Diagnostic Repository Command Interpreter (*ADRCI*), to troubleshoot Oracle database related errors. *ADRCI* was introduced in Oracle Database 11g in order to help users diagnose errors within their Oracle database environments and provide health reports if an issue should arise. The following example shows how one could troubleshoot an Oracle database instance error using the *ADRCI* tool.

**NOTE:** The following steps are intended to produce an ORA-07445 error that can be troubleshooted using the *ADRCI* tool. Do not attempt on a Oracle Database Production environment. The following is for demonstration purposes only and intended only to show how to troubleshoot ORA-\* related errors using the *ADRCI* tool.

Example of running *ADRCI* tool on node one of an Oracle RAC cluster.

1. In order to create an ORA-07445 error, an essential Oracle process will be killed via the following commands:

```
ps -A --format pid,args | grep ora_dbwm | grep -v grep
7811 ora_dbwm_cdb1
kill -SEGV 7811
```

2. Export the ORACLE\_HOME via the command:

```
export ORACLE_HOME=/u01/app/oracle/product/12.1.0/dbhome_1
```

3. Start the *ADRCI* command tool via the command:

```
$ORACLE_HOME/bin/adrci

ADRCI: Release 12.1.0.2.0 - Production on Wed Nov 12 10:44:50 2014

Copyright (c) 1982, 2014, Oracle and/or its affiliates. All rights reserved.

ADR base = "/u01/app/oracle"
adrci>
```

4. At the *ADRCI* prompt, show Oracle Home's available via the command:

```
adrci> show home
ADR Homes:
diag/rdbms/cdb/cdb1
diag/tnslsnr/db-oracle-node1/listener
```

**NOTE:** If more than one Oracle Home is available, one must specify a particular Oracle Database Home. An example on how to set to a particular Oracle Database Home is as follows:

```
adrci> set home diag/rdbms/cdb/cdb1
```

5. At the *ADRCI* prompt, run the following command to see the last 50 entries in the alert log:

```
adrci> show alert -tail -f
[... Output Abbreviated ...]
```





```
Exception [type: SIGSEGV, unknown code] [ADDR:0xD431000076FD]
[PC:0x3199AEB22A, semtimedop()+10] [exception issued by pid: 30461, uid:
54321] [flags: 0x0, count: 1]
Errors in file /u01/app/oracle/diag/rdbms/cdb/cdb/trace/cdb_dbrm_7811.trc
(incident=9673):
ORA-07445: exception encountered: core dump [semtimedop()+10] [SIGSEGV]
[ADDR:0xD431000076FD] [PC:0x3199AEB22A] [unknown code] []
Incident details in:
/u01/app/oracle/diag/rdbms/cdb/cdb/incident/incdir_9673/cdb_dbrm_7811_i96
73.trc
Use ADRCI or Support Workbench to package the incident.
See Note 411.1 at My Oracle Support for error and packaging details.
```

**NOTE:** In this particular case, we are looking for an ORA-07445 within the alert log as seen above. However, this step is just to confirm what is seen in the future *ADRCI* steps. To exit the alert log, execute CTRL+C.

6. Within the *ADRCI*, there are two key terms to be aware of, problem and incident. An incident is a particular time when a problem occurred. For example, it is possible for an Oracle process to crash at different times with the same ORA-07445. The multiple occurrences of the crash are incidents, while the problem is still the ORA-07445 error. In order to view the problem, the following *ADRCI* command needs to be run.

```
adrci> show problem
ADR Home = /u01/app/oracle/diag/rdbms/cdb/cdb1:

PROBLEM_ID PROBLEM_KEY LAST_INCIDENT LASTINC_TIME

1 ORA 7445 [semtimedop] 9673 2014-11-12 10:20:50.273000
-05:00
1 row fetched
```

7. In order to view how many incidents, the following *ADRCI* command must be run. In this example, I only have one incident in which the ORA-07445 problem occurred.

```
adrci> show incident
ADR Home = /u01/app/oracle/diag/rdbms/cdb/cdb:

INCIDENT_ID PROBLEM_KEY CREATE_TIME

9673 ORA 7445 [semtimedop] 2014-11-12 10:20:50.273000 -05:00
1 rows fetched
```

8. In order to view the incident in more detail, run the following command:

```
adrci> show incident -mode detail -p "incident_id=9673"
ADR Home = /u01/app/oracle/diag/rdbms/cdb/cdb:

INCIDENT INFO RECORD 1

INCIDENT_ID 9673
```



```

STATUS ready
CREATE_TIME 2014-11-12 10:20:50.273000 -05:00
PROBLEM_ID 1
CLOSE_TIME <NULL>
FLOOD_CONTROLLED none
ERROR_FACILITY ORA
ERROR_NUMBER 7445
ERROR_ARG1 semtimedop
ERROR_ARG2 SIGSEGV
ERROR_ARG3 ADDR:0xD431000076FD
[... Output Abbreviated ...]
OWNER_ID 1
INCIDENT_FILE
/u01/app/oracle/diag/rdbms/cdb/cdb1/trace/cdb_dbrm_7811.trc
OWNER_ID 1
INCIDENT_FILE
/u01/app/oracle/diag/rdbms/cdb/cdb1/incident/incdir_9673/cdb_dbrm_7811_i9673.trc
1 rows fetched

```

**NOTE:** The two parameters of importance here are the `PROBLEM_ID` and `INCIDENT_FILE`.

9. When looking at the incident in further detail, the following incident file can be examined further via the following command:

```

adrci> show trace
/u01/app/oracle/diag/rdbms/cdb/cdb1/incident/incdir_9673/cdb_dbrm_7811_i9673.trc
Output the results to file: /tmp/utsout_24439_14048_2.ado
/bin/bash: adrci: command not found

```

10. Open the `/tmp/utsout_24439_14048_2.ado` file with an editor such as `vi`.

```

/
u01/app/oracle/diag/rdbms/cdb/cdb1/incident/incdir_9673/cdb_dbrm_7811_i9673.trc

LEVEL PAYLOAD

Dump file
/u01/app/oracle/diag/rdbms/cdb/cdb1/incident/incdir_9673/cdb_dbrm_7811_i9673.trc
Oracle Database 12c Enterprise Edition Release 12.1.0.2.0 - 64bit
Production With the Partitioning, Automatic Storage Management, OLAP,
Advanced Analytics and Real Application Testing options
ORACLE_HOME = /u01/app/oracle/product/12.1.0/dbhome_1
System name: Linux
Node name: db-oracle-node1.cloud.lab.eng.bos.redhat.com
Release: 3.10.0-123.el7.x86_64
Version: #1 SMP Mon May 5 11:19:54 EDT 2014
Machine: x86_64
Instance name: cdb1
Redo thread mounted by this instance: 2
Oracle process number: 9

```



Unix process pid: 7811, image: [oracle@db-oracle-node1](#) (DBRM)

```
*** 2014-11-12 10:20:50.293
*** SESSION ID:(271.1) 2014-11-12 10:20:50.293
*** CLIENT ID:() 2014-11-12 10:20:50.293
*** SERVICE NAME:(SYS$BACKGROUND) 2014-11-12 10:20:50.293
*** MODULE NAME:() 2014-11-12 10:20:50.293
*** ACTION NAME:() 2014-11-12 10:20:50.293
*** CONTAINER ID:(1) 2014-11-12 10:20:50.293
```

Dump continued from file:

/u01/app/oracle/diag/rdbms/cdb/cdb1/trace/cdb\_dbrm\_7811.trc

1> \*\*\*\*\* Error Stack \*\*\*\*\*

ORA-07445: exception encountered: core dump [semtimedop()+10] [SIGSEGV]  
[ADDR:0xD431000076FD] [PC:0x3199AEB22A] [unknown code] []

1< \*\*\*\*\* Error Stack \*\*\*\*\*

1> \*\*\*\*\* Dump for incident 9673 (ORA 7445 [semtimedop]) \*\*\*\*\*

2> \*\*\*\*\* Beginning of Customized Incident Dump(s) \*\*\*\*\*

Dumping swap information

Memory (Avail / Total) = 42169.27M / 64420.08M

Swap (Avail / Total) = 16415.99M / 16415.99M

Exception [type: SIGSEGV, unknown code] [ADDR:0xD431000076FD]

[PC:0x3199AEB22A, semtimedop()+10] [exception issued by pid: 30461, uid: 54321]

[flags: 0x0, count: 1]

Registers:

%rax: 0xffffffffffffffffc %rbx: 0x000000000000a000b %rcx:

0xfffffffffffffffff

[ ... Output Abbreviated ... ]

11. While this concludes how to examine trace files that pertain to a particular ORA error using *ADRCI*; if the issue cannot be solved by the end user, the *ADRCI* provides the *Incident Packaging Service* (IPS) tool to ZIP the necessary trace files based on the problem. It can then be sent to support for further debugging. To create the appropriate ZIP file, use the following commands:

```
adrci> ips create package problem 1 correlate all
Created package 1 based on problem id 1, correlation level all
```

**NOTE:** Problem 1 is the Problem\_ID found in step 6.

```
adrci> ips generate package 1 in "/home/oracle"
Generated package 1 in file
/home/oracle/ORA7445se_20140120153727_COM_1.zip, mode complete
```

**NOTE:** Package 1 is the package ID captured from the ips create output command.

For more information about *ADRCI*, please visit the

[http://docs.oracle.com/cd/E16655\\_01/server.121/e17639/adrci.htm#SUTIL700](http://docs.oracle.com/cd/E16655_01/server.121/e17639/adrci.htm#SUTIL700)



# Appendix O: References

**TECH: Unix Semaphores and Shared Memory Explained [ID 15566.1]<sup>33</sup>**

<https://support.oracle.com/epmos/faces/DocumentDisplay?id=15566.1>

**Oracle Grid Infrastructure, Oracle Documentation**

[http://docs.oracle.com/cd/E16655\\_01/install.121/e17888/concepts.htm#CWLIN506](http://docs.oracle.com/cd/E16655_01/install.121/e17888/concepts.htm#CWLIN506)

**Tuning Red Hat Enterprise Linux For Oracle & Oracle RAC by Scott Crot, Sr. Consultant, Red Hat**

[http://www.redhat.com/promo/summit/2010/presentations/summit/decoding-the-code/fri/scott-945-tuning/summit\\_jbw\\_2010\\_presentation.pdf](http://www.redhat.com/promo/summit/2010/presentations/summit/decoding-the-code/fri/scott-945-tuning/summit_jbw_2010_presentation.pdf)

**USE\_LARGE\_PAGES To Enable HugePages [ID 1392497.1]<sup>33</sup>**

[https://support.oracle.com/epmos/faces/Dashboard?\\_adf.ctrl-state=nvtwimbst\\_252](https://support.oracle.com/epmos/faces/Dashboard?_adf.ctrl-state=nvtwimbst_252)

**Large Pages Information in the Alert Log [ID 1392543.1]<sup>33</sup>**

[https://support.oracle.com/epmos/faces/ui/km/SearchDocDisplay.jspx?\\_adf.ctrl-state=nvtwimbst\\_226](https://support.oracle.com/epmos/faces/ui/km/SearchDocDisplay.jspx?_adf.ctrl-state=nvtwimbst_226)

**Tuning Virtual Memory**

[https://access.redhat.com/site/documentation/en-US/Red\\_Hat\\_Enterprise\\_Linux/6/html/Performance\\_Tuning\\_Guide/s-memory-tunables.html](https://access.redhat.com/site/documentation/en-US/Red_Hat_Enterprise_Linux/6/html/Performance_Tuning_Guide/s-memory-tunables.html)

**Maximum SHMMAX values for Linux x86 and x86-64 [ID 567506.1]<sup>33</sup>**

[https://support.oracle.com/epmos/faces/ui/km/SearchDocDisplay.jspx?\\_adf.ctrl-state=yp0o5bwk6\\_4](https://support.oracle.com/epmos/faces/ui/km/SearchDocDisplay.jspx?_adf.ctrl-state=yp0o5bwk6_4)

**About the Oracle Database Fault Diagnosability Infrastructure**

<http://docs.oracle.com/database/121/ADMIN/diag.htm#ADMIN11257>

**Oracle 12c Database Installation Guide**

<https://docs.oracle.com/database/121/LADBI/toc.htm>

**Pro Oracle Database 11g RAC on Linux – Installation, Administration, Performance by Steve Shaw and Martin Bach**

[http://www.amazon.com/Pro-Oracle-Database-11g-Linux-ebook/dp/B004VJ472I/ref=sr\\_1\\_1?ie=UTF8&qid=1389977723&sr=8-1&keywords=pro+oracle+11g+rac](http://www.amazon.com/Pro-Oracle-Database-11g-Linux-ebook/dp/B004VJ472I/ref=sr_1_1?ie=UTF8&qid=1389977723&sr=8-1&keywords=pro+oracle+11g+rac)

---

<sup>33</sup> My Oracle Support Credentials Required to access article

